

Loom — Manual

Loom is a browser-based, session-based music workstation. It organises sound into **lanes** (instrument tracks), **clips** (patterns of notes), and **scenes** (rows of clips launched together). Eight synthesis engines are available — TB-303, Subtractive, FM, Wavetable, Karplus-Strong, West Coast, Sampler, and Drum Machine — and all synthesis runs live in the browser with no install or sign-in required.

Beyond the engines, Loom has **note velocity & dynamics** (an Ableton-style velocity lane), **per-clip and arrangement-wide loop regions**, **per-lane modulation and FX** with a sidechain mixer, **MIDI import**, ready-made **sampled drum kits**, a **playable arrangement** timeline, improved **session editing workflow** (inline track/scene rename, duplicate track/scene, and capture-playing-to-scene), optional **stem separation + transcription** (split a finished song into 4 lanes via a local helper and optionally derive editable note/drum lanes), **key & scale musical assistance** (a project key/scale with an optional piano-roll scale lock — off by default, genre generators, an examples gallery, and a chord-accompaniment maker), and **editable warp markers** for locking audio loops to tempo.



The Loom interface: transport across the top, the session clip grid in the centre, and per-lane channel strips below.

Table of Contents

#	Chapter
1	Getting Started
2	Transport
3	Sessions, Lanes, Clips & Scenes
4	Engines
5	Editing Clips
6	Modulation & Note FX
7	Mixing & FX
8	MIDI & Samples
9	Saving & Export
10	Performance & Arrangement
11	Developer Guide

[Loom-Manual.pdf](#)

Live demo: <https://ijol.github.io/Loom/>

Getting Started

Open the app

Loom runs entirely in the browser — no installation, no account, no plugins.

- **Live demo:** <https://ijol.github.io/Loom/>
- **Run locally:** clone the repository, then run `npm install` followed by `npm run dev`. The app is available at <http://localhost:5173>.

Make your first sound

When the app opens it loads the *Minimal Techno* demo automatically, so there is already a full arrangement ready to play.

1. Click the ► **Play** button in the transport bar. The first click resumes the browser's `AudioContext` — this is a browser requirement; audio cannot start without a user gesture, so nothing plays until you press Play.
2. Sound starts immediately. Use the **Volume** slider in the transport to adjust the output level.
3. Want to hear a different demo? Select one from the — **load a demo** — dropdown (Minimal Techno, Acid Rain, Cordillera, Neon Drive). The session loads and begins at the first scene.



Loom with the Minimal Techno demo loaded and the session grid visible

The full Loom interface. Transport across the top; coloured clips in the session grid; per-lane channel strips below.

The menu bar

Across the very top sits a desktop-style menu bar. Most of what it offers also has a button somewhere in the interface — the menu is the tidy, discoverable route to the same things, and the place to look when you can't remember where a control lives.

File

Item	Shortcut	What it does
New Session	Ctrl+N	Empties the session and starts fresh
Open...	Ctrl+O	Opens the save manager to load a stored session
Save	Ctrl+S	Opens the save manager to store the current session
Save As...		Same panel, for storing under a new name
Project Options...		Project-wide settings, including the name and the musical key/scale
Open Demo ▶		A submenu of the bundled demos; picking one loads it
Import MIDI...		Opens the MIDI import flow (chapter 8)
Separate into Stems...		Opens stem separation (chapter 8)
Preferences...	Ctrl+,	Not yet available — greyed out, planned for a later release

Edit — **Undo** (Ctrl+Z) and **Redo** (Ctrl+Shift+Z , or Ctrl+Y). Both grey out when there is nothing to undo or redo, so the menu doubles as an indicator of whether your last action was recorded.

View — switches between the **Session** and **Performance** views, with a tick marking the current one, and toggles the **Performance diagnostics (PERF)** overlay (audio load, scheduler lag, FPS, voice counts).

Tools — opens the **MIDI Controller** panel, runs **Capture Scene** (Ctrl+I , snapshots the currently playing clips into a new scene) and **Copy Scenes** → **Performance**.

Help — opens this **Manual** in a new tab, and **About Loom**.

Sessions are stored in your browser, not as files on disk: "Open" and "Save" mean the in-browser save manager, described in [chapter 9](#).

The mental model

Loom organises music in three nested concepts:

- A **lane** is an instrument track. Each lane runs one engine — six melodic ones (TB-303, Subtractive, FM, Wavetable, Karplus-Strong, West Coast), the Sampler, the Drum Machine, or an Audio channel for playing a recording — and has its own mixer strip with EQ, sends, pan, and fader.
- A **clip** is a pattern of notes that lives in a lane. A lane can hold multiple clips — one per scene row.
- A **scene** is a horizontal row across all lanes. Launching a scene fires the clip in each lane at that row simultaneously, so scenes work like song sections or loop variations.



The session clip grid showing four lanes across four scenes with coloured clips

The session grid: lanes run left to right as columns, scenes top to bottom as rows. Click a clip to select it; click the scene launch button on the right to fire the whole row.

Click any clip cell to open its editor in the inspector below the grid. The editor shows a piano roll for melodic lanes and a drum grid for drum lanes.

Next steps

- [Transport](#) — BPM, meter, swing, and the unified REC controls (take / live-WAV / offline-WAV).
- [Sessions, Lanes, Clips & Scenes](#) — how to build and arrange your own session from scratch.

Transport

The header runs across the top of the interface in two rows. The first row (transport & tempo) holds the playback and timing controls you touch while the music plays; the second row (session & I/O) holds the mode toggle, recording/export, and session file management.



Play / Stop

Press ► (`#play` , title "Play — arranca el transporte") to start playback, and ■ (`#stop` , title "Stop — detiene el transporte") to stop it. Play and Stop are now two separate buttons — Play is no longer a play/stop toggle. When you press Play the browser's `AudioContext` resumes automatically if it was suspended (a browser policy requirement on first interaction).

Position and time readout

The readout shows two values side by side:

- **Bar.Beat.Step** (`#transport-position` , e.g. `1.1.1`) — the current song position derived from elapsed time, BPM, and the active meter. The counter resets to `1.1.1` on each new Play.
- **Elapsed** (`#transport-time` , e.g. `00:00:00`) — wall-clock time since the most recent Play press.

Both readouts update via `requestAnimationFrame` while playing and freeze on stop so you can read the final position. Because each lane has its own independent loop clock, the global position is computed from elapsed seconds rather than a single sequencer cursor. The visual playhead is similarly a display-only timer matched to scheduled audio time; it may drift slightly under browser tab throttling, but audio scheduling is unaffected.

Tempo controls

BPM (`#bpm`) — sets the tempo. Range: 40–400 BPM, default 130. (The ceiling is deliberately high: drum & bass written at its true tempo, and the faster end of hardcore/gabber, both live above 240.) The field accepts **floating-point values** (e.g. 128.57 BPM, step 0.01); tempos detected from MIDI imports or stem separation are no longer rounded to an integer. You can type a value directly or use the number-input spinners. BPM changes propagate immediately to the sequencer, all lane engines, delay/LFO sync, and stretch-mode loop buffers. The change takes effect on the next scheduled step; it does not alter a note that is already held.

Each 16th-note step lasts `60 / bpm / 4` seconds.

Meter (`#meter`) — sets the global time signature. The dropdown offers common meters: 4/4, 3/4, 2/4, 5/4, 6/8, 7/8, 9/8, and 12/8. The meter controls how bars map onto the 16th-step grid and how the position readout counts beats. Changing the meter takes effect on the next loop cycle.

Swing (`#swing`) — adds a shuffle feel by delaying the off-beat 16th notes. Range: 0 (straight) to 0.6. A swing of `s` delays each off-beat 16th by `s × one 16th step`, so the off-beat lands `50 × (1 + s)` % of the way to the next on-beat — the familiar swing percentage: **0** = straight (50%), **0.33** = the classic triplet shuffle (66.7%), **0.5** = a heavy shuffle (75%), **0.6** = the hardest limp Loom offers (80%). On-beats never move.

Swing applies to every lane, and to notes that do not sit exactly on a 16th too (32nds, imported MIDI, off-grid drum hits): each note is carried along with the beat it belongs to rather than left behind by it. Note lengths follow their notes, so a TB-303 slide still slides. The value is saved and restored with the session, and WAV exports render with the same swing you hear.

Master volume

Volume (`#volume`) — the master output level, range 0–1 (default 0.5). This is a post-mix gain applied before the output visualiser.

Mode toggle and REC

Session / Performance (`#mode-toggle`) — switches the main view between the Session clip grid and the Performance arrangement view.

† **Copy to Performance** (`#copy-to-performance`) — an icon-only button (tooltip "Copiar las escenas a la timeline de Performance") that copies the current scenes onto the Performance timeline. See [Performance & Arrangement](#).

- **Session** is the default mode: you see the clip grid and can trigger scenes, edit clips, and record automation in real time.
- **Performance** shows the arrangement timeline where recorded takes are displayed as timeline bands and automation curves can be drawn directly. See [Performance & Arrangement](#) for the full workflow.
- **REC** (`#rec` , tooltip "Grabar — el modo se elige al lado") arms recording; the **mode selector** (`#rec-mode`) beside it chooses what gets recorded:
 -  **take** (`data-recmode="take"` , the default) — records knob moves and clip launches into a performance take. See [Modulation & Note FX](#) for how automation lanes work.
 -  **live** (`data-recmode="live"`) — records the real-time audio output to a WAV.

- ⚡ **offline** (`data-recmode="offline"`) — renders the current scene to a WAV offline (fast).

Click **REC** again to disarm. This unified REC group replaced the old standalone WAV-export button.

Output visualiser

The canvas at the far right (`#viz`) shows a real-time waveform of the master output. It updates continuously while the page is open and gives a quick visual check that audio is flowing.

Session management, export, and demos

The remaining controls in the transport row are covered in dedicated chapters:

- **WAV export** — there is no longer a standalone "↓ WAV" button. WAV export now happens through the **REC** group's mode selector (see the Mode toggle and REC section above): 🎙️ **live** records real-time audio to WAV, and ⚡ **offline** renders the scene to WAV offline (fast). See [Saving & Export](#).
 - **New / Save / Load** (`#new-session` , `#save` , `#load`) — session file management. See [Saving & Export](#).
 - — **load a demo** — (`#demo-picker`) — loads a bundled demo arrangement into the session.
 - **MIDI import** — there is no button for this in the transport row. It lives under **File ▶ Import MIDI...**, which opens a modal dialog. See the MIDI Import chapter for details.
 - ☰ (`#stems-open` , an icon-only button tooltipled "Separate a song into stems") — opens the stem-separation dialog, which splits a finished song into four lanes (vocals / drums / bass / other) via a local helper service. Also under **File ▶ Separate into Stems....** Requires the service to be running; see [MIDI & Samples → Stem separation](#).
 - 📸 **Capture** (`#capture-scene`) — snapshots the clips currently playing into a new scene (also `Ctrl+I` , and under **Tools ▶ Capture Scene**).
 - 📐 **XY** (`#xy-open`) — opens the floating XY pad for sweeping two parameters at once. See [Mixing & FX → XY pad](#).
 - **PERF** (`#perf-toggle`) — toggles the performance-diagnostics overlay (audio load, scheduler lag, FPS, voice counts). It only runs while open.
 - ↶ / ↷ (`#undo-btn` , `#redo-btn`) — global undo and redo, greyed out until there is history.
-

Keyboard shortcuts

Key	Action
Space	Pause / resume
R	Arm or disarm clip MIDI recording
Ctrl/Cmd+N	New session
Ctrl/Cmd+O	Open
Ctrl/Cmd+S	Save
Ctrl/Cmd+Z	Undo
Ctrl/Cmd+Shift+Z or Ctrl/Cmd+Y	Redo
Ctrl/Cmd+I	Capture playing clips into a new scene

Two things about **Space** that catch people out:


- It is **pause/resume, not stop**. It remembers the exact position — including the fraction of a bar — and picks up from there, where Stop resets to the beginning.
- It does **nothing from a standing start**. If the transport is idle, Space will not begin playback; use ► or launch a scene. It only pauses something already playing.

Note also that after a pause/resume the position readout restarts its count from **1.1.1** even though the music continues where it left off — the readout times the current run, not the song.

Shortcuts never fire while you are typing in a text field, so they can't steal input from the BPM box or a name field.

Sessions, Lanes, Clips & Scenes

The session view is the main workspace in Loom. It organises everything you hear into a compact grid of lanes and clips that you can launch, edit, and rearrange in any order.



The model

Lane — one instrument track. Each lane runs its own synthesis engine (TB-303, Drums, Subtractive, FM, Wavetable, Karplus, or Sampler) with its own mixer strip and insert chain. Lanes appear as columns in the grid.

Clip — a pattern of notes that lives at one lane × scene cell. Every clip stores a list of note events, an optional name, a colour, and a length in bars. A cell is either filled (it holds a clip) or empty.

Scene — a row in the grid. Launching a scene fires one clip per lane simultaneously, aligning all starts to the same quantize boundary so the whole arrangement snaps together cleanly.

The grid



Lane names run along the top header row; row numbers (1, 2, 3 ...) label the scenes down the left edge. Scenes launch buttons sit in the rightmost column. Each filled cell shows the clip's name (or its row number as a fallback) against the clip's pastel colour. **Click a lane header** to open that lane's instrument editor; once it is open, the ▶ / ▼ **chevron** that appears on that header collapses and reopens the editor.

A newly created instrument lane starts out empty — no placeholder clips are added to its column; you fill cells yourself. Audio and Sampler lanes created from a dropped WAV or loop place their clip in row 1 only. The grid always keeps at least one launchable scene available even when every lane is empty.

Right-clicking grid elements opens a context menu. On a **lane header**: *Rename track*, *Edit instrument*, *Duplicate track*, *Stop track*, and *Delete track* (red). On a **scene cell**: *Rename scene*, *Launch scene*, *Duplicate scene*, *Capture playing -> scene*, *Add scene*, and *Delete scene* (red). On a **filled clip**: *Open editor*, *Play / Stop*, and *Delete clip* (red). On an **empty cell**: *Create clip* — on audio lanes this entry reads *Import audio (WAV)*....

Below the scene rows there is a stop row: each lane has its own ■ stop button, and a global ■ **all** button at the far right stops every lane at once.

Deleting from the grid

A small ✕ cross appears on every lane header (*Delete track*), every filled clip cell (*Delete clip*), and every scene cell (*Delete scene*). Deleting a clip is immediate. Deleting a lane or scene asks you to confirm in an in-app dialog (**OK** / **Cancel**, with the destructive choice in red) only when the target still has content — an empty lane or scene is removed straight away. Every deletion is undoable with **Ctrl+Z**, which restores the lane, scene, or clip along with its audio resources.

Launching clips and scenes

Click the ► **icon** on any filled cell to launch that clip. The icon switches to || once the clip is playing; clicking || stops that lane. Clicking anywhere else on the cell body — outside the play icon — opens the inspector without affecting playback.

From a cold start (transport stopped), launching a clip starts it immediately, and the **Quantize** setting decides the alignment of that first start. **While the session is already playing**, Loom keeps everything phase-aligned instead:

- **Launching a single clip** swaps it in at the end of *its own* current loop, so the new pattern enters cleanly on the downbeat rather than mid-bar.
- **Launching a scene** is an **atomic switch at the end of the governing loop** — the longest loop currently playing (with stray outliers capped, so one odd 1-bar clip doesn't drag a 4-bar groove). Every clip in the scene changes together on that boundary, and any **lane that has no clip in the new scene stops** at the same moment (it pulses a "stopping" state until then). The result is a clean, simultaneous scene change with no lanes left hanging.

Click the ► **scene name** button at the right of a scene row (it reads ► **Scene 1**, or whatever you renamed it to — not a bare number) to launch all clips in that row. The toolbar's ■ **All** button stops every lane at once.

The **Quantize** setting (per clip, falling back to the lane's, then the session's global) governs **cold starts** — the alignment used when you launch from a stopped transport. Once the transport is running, the loop-end sync above takes over so launches stay musically aligned.

Renaming, duplicating, and capturing scenes

Track and scene names support **inline rename**: double-click a lane name in the header or a scene name in the scene cell, type the new label, and press Enter.

You can also use the context menus to duplicate structure quickly:

- **Duplicate track** clones the whole lane with fresh clip ids and keeps its clip content.
- **Duplicate scene** appends a full clone of that scene's lane/clip assignments.

To snapshot the current live state into a new scene, use **Capture scene**:

- Click the ☉ button in the **Scenes** header (tooltip: *New scene from currently playing clips*).

- Or press **Ctrl+I**.
- Or use **Capture playing -> scene** from a scene cell's context menu.

Capture scene appends a new scene row and writes one clip-per-lane mapping from whatever is playing at that moment.

The inspector

Clicking a clip cell body opens the inspector below the grid with controls for that clip.



Control	What it does
Name	Free-text label shown on the cell.
Length (bars)	How many bars the clip plays before looping.
Quantize	Per-clip launch alignment: <i>Default</i> (inherits lane/global setting), <i>Immediate</i> (no wait), or a fixed interval — 1/4, 1/2, 1 bar, 2 bars, or 4 bars.
Copy notes	Copies the clip's notes to an internal clipboard.
Paste ▸ Replace	Replaces the selected clip's notes with the clipboard contents.
Paste ▸ Layer	Merges the clipboard notes into the selected clip rather than replacing them.
Duplicate	Creates a copy of the clip in the next empty cell of the same lane.
↔ Editor	Toggles between the piano-roll editor and the drum-grid editor. The appropriate editor is chosen automatically based on the lane's engine; use this button to switch if needed. See Editing Clips for the full editing reference.
Notes	Randomizes the clip's note content using the current scale and root settings.
Delete	Removes the clip from the grid (also triggered by Delete/Backspace while the inspector is open).

The editor embedded inside the inspector is rebuilt each time a new clip is selected; deep editing of notes is covered in [Editing Clips](#).

Arranging clips

You can move a clip to any other cell by dragging it. Hold **Ctrl** while dragging to copy instead of move. Clips remember their colour across moves and copies. Drag a clip onto a cell that already has a clip to swap or replace it.

Each clip is assigned a colour automatically from a pastel palette. You can distinguish patterns at a glance even before you give them names.

For the mixer strip, insert effects, and per-lane send levels, see [Mixing & FX](#). For the available synthesis engines on each lane, see [Engines](#).


Engines

Every lane in Loom runs exactly one synthesis engine. You choose the engine with the ENGINE selector at the top of the lane's editor panel. Changing the engine replaces the sound source while preserving the lane's clips and modulation routing. Eight engines are available: six melodic synthesisers (TB-303, Subtractive, FM, Wavetable, Karplus-Strong, West Coast), a Sampler, and a Drum Machine.

Each engine exposes a PRESET dropdown (Load / Save As / Delete) and a  **Sound** button that randomises the patch and sets the preset name to "Custom". Presets are JSON assets stored per-engine; they include GM programme tags so that MIDI import can auto-assign the best engine and preset for each track. See [Sessions, Lanes, Clips & Scenes](#) for how to add and configure lanes.

Every engine responds to **note velocity**. Velocity (0–127) scales each note's loudness continuously; accent (velocity ≥ 100) layers additional character on top — brightening the filter envelope and resonance on bass-style engines, and adding brightness on drums. For how to view and edit velocities in the piano-roll or drum-grid, see [Velocity & dynamics](#).

TB-303



Above: TB-303 editor — Wave, Cutoff, Resonance, Env Mod, Decay, Accent, and a per-lane LFO.

The TB-303 is a monophonic, resonant bass synthesiser modelled on the Roland TB-303. It is the natural choice for acid bass lines but works equally well for aggressive leads and any sound that calls for a steep, self-oscillating filter sweep.

TB-303 parameters

Parameter	Description
Wave	Sawtooth or Square oscillator waveform
Cutoff	Filter cutoff frequency (0–100%)
Resonance	Filter Q — see the note on the diode ladder below
Env Mod	How far the filter envelope opens the filter per step
Decay	Filter envelope decay time
Accent	Per-step level: brightens the filter, bumps Q, raises output gain

The filter is a diode ladder. This is the circuit that gives the 303 its voice: its asymmetric clipping adds even harmonics, which is where the squelch comes from. One consequence surprises people used to other synths — a ladder *loses* level as resonance rises, instead of growing a peak on top of the signal. Turning Resonance up thins and quietens the sound rather than making it louder. That is correct behaviour, and the engine compensates internally so that accented steps still punch through.

Slide and accent behaviour

A note's `slide` flag means "slide into the next step". When the scheduler emits step N it checks whether step N-1 carried a slide flag; if so, it ramps the pitch from the previous note and skips the amp re-attack so the gate stays open across the boundary. The outgoing step gets an extended gate (1.5× step length) so the overlap is audible.

Accent is a per-step flag that simultaneously brightens the filter envelope, raises the resonance Q, and boosts the output gain — the classic 303 bassline punctuation technique.

The engine ships with 20+ presets from "BASS Acid Classic" to "LEAD Squelch". See [Editing Clips](#) for how to set slide and accent on individual steps.

Subtractive



Above: Subtractive editor — OSC 1/2, Sub oscillator, Noise, Filter (with built-in envelope), Amp, and POLY controls.

The Subtractive engine is a classic analogue-style polyphonic synthesiser with two oscillators, a sub oscillator, a noise source, a multimode filter, and a full amplitude envelope. It is the most general-purpose engine in Loom and suits pads, leads, basses, and plucks — and with **85 presets** it has by far the largest library of any engine here.

Parameter sections

- **OSC 1 / OSC 2** — waveform, level, detune in cents, **PW**, and **Sync**. Detuning the two oscillators creates the classic "supersaw" chorus effect. See [Oscillator extras](#) below for PW and Sync.
- **SUB / NOISE** — sub oscillator level (one octave below OSC 1) and a noise generator for breath and texture.
- **FILTER** — **Model**, **Type**, Cutoff, Resonance, Env Amount, Drive, Key Track, and a full ADSR filter envelope (toggle with Built-in Env). See [Filter model and type](#).
- **AMP** — Attack/Decay/Sustain/Release amplitude envelope (toggle with Built-in Env).
- **MASTER** — global Tune in semitones, plus **Unison**, **Detune** and **Drift**. See [Unison](#).

- **POLY** — voice count (1–16), poly/mono mode, and legato/retrig behaviour in mono mode.

Filter model and type

Two dropdowns sit above the filter knobs and change its character completely.

Model picks the circuit being emulated:

Model	Character
DIG (default)	A clean state-variable filter. Precise and neutral.
MOG	A four-pole Moog-style ladder. Warmer, and it saturates as you drive it.
303	The diode ladder from the TB-303. Asymmetric clipping adds even harmonics — the acid voice.

Type picks the response: **LP**, **HP**, **BP** or **Notch**.

Important: only **DIG** is a true multimode filter. The two ladders (MOG and 303) are four-pole lowpass topologies and are **lowpass only** — with Model set to MOG or 303, the Type dropdown has no effect. If you want a high-pass or a notch, you need DIG.

A second thing worth knowing about the ladders: they *lose* level as resonance climbs, rather than growing a resonant peak on top. Turning Q up on MOG or 303 thins and quietens the sound. That is faithful to the hardware, and it is why the TB-303 engine compensates with a dedicated accent gain.

Oscillator extras: PW and Sync

PW (Pulse Width) — 0.05 to 0.95, centred at 0.5. On its own it reshapes a square wave from a hollow, nasal thin pulse through the full-bodied square at 0.5 and out the other side. The classic use is **PWM**: assign an LFO to `osc1.pw` and the width breathes back and forth, turning a static square into a wide, shimmering pad. This is the single most rewarding modulation target on the engine.

Sync (hard sync) — select **Sync** as the waveform and the oscillator runs a second, faster oscillator that is force-restarted on every cycle of the first. The restart chops the waveform mid-cycle, producing the aggressive, tearing timbre of a synced lead. Here the **Sync knob (1–8)** sets the ratio between the two, and *that ratio is the timbre* — the pitch you hear still follows the note. Sweeping the Sync knob (or modulating it) gives the classic sync-sweep lead. Presets **LEAD Sync**, **LEAD Sync Sweep** and **BASS Sync** are built on it.

Unison, Detune and Drift

In the MASTER section:

Parameter	Description
Unison	Stack 1–7 copies of the whole voice per note (default 1)
Detune	How far the stack spreads apart, 0–50 cents (default 25)
Drift	Slow, random analogue pitch wander, 0–1 (default 0)

Unison is what makes a lead enormous: seven slightly-detuned copies beating against each other is the supersaw sound. It costs CPU in proportion — seven copies is seven times the synthesis work per note — so raise it for leads, not for dense pads.

Unison is read once when the note is triggered, so unlike most parameters it is deliberately *not* a modulation target; changing it affects the next note you play, not the one currently sounding. **Drift** is the opposite: a touch of it (0.1–0.2) keeps sustained chords from sounding sterile.

For modulation routing see [Modulation & Note FX](#).

FM



Above: FM editor — Algorithm selector, Op 1–4 (Ratio/Detune/Level/ADSR), global Mix and Voices.

The FM engine is a four-operator, DX7-style frequency-modulation synthesiser. Each operator is a sine oscillator with its own ADSR amplitude envelope and level; operators are wired together according to one of four algorithms.

FM parameters

Parameter	Description
Algorithm	1 = serial 4→3→2→1; 2 = three parallel mods → Op 1; 3 = two pairs; 4 = additive
FB (Op4)	Op 4 self-feedback — adds odd harmonics and edge
Op 1–4: Ratio	Frequency ratio relative to the played note (0.1–16×)
Op 1–4: Det	Per-operator detune in cents
Op 1–4: Level	Carrier output level or modulation index (modulators)
Op 1–4: ADSR	Per-operator amplitude envelope
Mix	Final output level (0–1)
Voices	Polyphony cap (1–16; default 6)

FM suits metallic tones, electric pianos, bells, and evolving textures. Small ratio changes yield very different timbres; the preset library covers bells, organs, and electronic basses.

Wavetable



Above: Wavetable editor — Wave A/B selectors, Morph, Detune, Filter, Amp envelope, and Voices.


The Wavetable engine morphs between two pre-computed waveforms — Wave A and Wave B — using the Morph knob or an LFO/ADSR routed to it. Both waves are drawn from a fixed bank of eight anti-aliased tables: Sine, Triangle, Sawtooth, Square, Pulse (25%), Organ, Brass, and Vocal.

Wavetable parameters

Parameter	Description
Wave A / Wave B	Source waveforms to interpolate between
Morph	Crossfade position (0 = full Wave A, 1 = full Wave B)
Detune	Global detune in cents
Cutoff / Res	Resonant low-pass filter
Built-in Env	Toggle the built-in amp ADSR on/off
Attack / Decay / Sustain / Release	Amplitude envelope
Voices	Polyphony cap (1–16; default 8)

Animating Morph with an LFO is the signature technique — sweeping from Sine to Sawtooth or Brass to Vocal while a note sustains produces evolving, living tones. See [Modulation & Note FX](#).

Karplus-Strong



Above: Karplus-Strong editor — String section (Damping, Brightness), Excite section (Excite time, Noise Tone), and Amp controls.

Karplus-Strong is a physical-modelling engine that synthesises plucked-string sounds. Loom renders each note offline (sample-by-sample in JavaScript) into a buffer and plays it back through an amplitude envelope. This gives exact pitch at every frequency, natural high-harmonic roll-off, and no feedback runaway.

Karplus-Strong parameters

Parameter	Description
Damping	T60 decay: 0 = long sustain (~4 s), 1 = muted (~0.12 s)
Brightness	Loop filter: 0 = dark/cello, 1 = open/metallic
Excite	Excitation burst length (pluck sharpness)
Noise Tone	Colour of the excitation noise (dark → bright)
Attack / Release	Amp envelope on buffer playback
Level	Output amplitude
Voices	Polyphony cap (1–16; default 8)

Damping and Brightness are set per-note at the moment of the pluck (baked into the buffer). Level and its envelope are live AudioParams and can be modulated. Karplus suits acoustic bass, guitar, harp, and marimba-style sounds.

West Coast



West Coast editor panel

A Buchla-style "West Coast" voice: complex oscillator → wavefolder → low-pass gate, driven by a built-in AD contour.

The West Coast engine takes a fundamentally different approach to synthesis from the filter-based ("East Coast") engines above. Instead of shaping a harmonically-rich waveform by subtracting frequencies with a filter, it **adds** harmonics by routing an oscillator through a **wavefolder** and then taming the result with a **low-pass gate**. The result is highly percussive and organic — metallic plucks, woody tones, evolving bell-like pads, and abstract textures that are difficult to achieve with conventional subtractive synthesis.

The engine is **polyphonic** — up to 16 voices, 8 by default — with a Poly/Mono switch in the AMP section if you want the classic one-voice-at-a-time behaviour. FM here is native-linear, not through-zero.

The signal chain is: **Complex Oscillator** → **Timbre (Wavefolder)** → **Low-Pass Gate (LPG)** driven by an **AD Contour**.

COMPLEX OSCILLATOR section

Two cross-coupled oscillators produce the raw material. The principal oscillator's frequency is modulated by the modulator oscillator (linear FM), optionally ring-modulated, and a sub-harmonic divider can be added below.

Parameter	Description
Princ Wave	Waveform of the principal oscillator: Sin, Tri, or Saw
Mod Wave	Waveform of the modulator oscillator: Sin or Tri
Ratio	Frequency ratio of the modulator relative to the principal (0.25–16×) — integer ratios produce harmonic tones; non-integers produce inharmonic, bell-like partials
FM Index	Depth of linear FM from the modulator into the principal (0–1) — higher values add more sidebands
Ring/AM	Amount of ring modulation mixed in (0 = off, 1 = full ring mod)
Sub ÷	Sub-harmonic divider: Off, ÷2, ÷3, or ÷4 — adds a sub-oscillator one, two, or three octaves below
Sub Lvl	Output level of the sub-harmonic oscillator (0–1)
Detune	Fine-tune of the principal oscillator in cents (± 50 ¢)

TIMBRE section (wavefolder)

The wavefolder processes the summed oscillator signal through a non-linear waveshaping curve. As the **Fold** amount increases it drives the signal harder into the curve, folding the waveform back on itself and adding a cascade of new harmonics. Accent (velocity ≥ 100) pushes the fold drive harder automatically.

Parameter	Description
Fold	Drive into the fold curve (0 = gentle/clean, 1 = heavy folding/maximum harmonics)
Symmetry	DC bias applied before the folder — shifts the waveform asymmetrically for even-harmonic colouring (–1 to +1)

LOW-PASS GATE section

A low-pass gate combines a resonant filter and a VCA in a single, vactrol-like element so that the contour simultaneously opens the brightness and the volume. The **Mode** selector determines how the contour is routed.

Parameter	Description
Mode	LP — contour sweeps the filter only (VCA stays open); Gate — contour opens the VCA only (filter stays at its base cutoff); Both — contour drives both, the most classic "plonky" Buchla behaviour
Cutoff	Base cutoff frequency (0–1, exponential scaling from ~60 Hz to 18 kHz)
Resonance	Filter Q (0–1) — adds resonant emphasis at the cutoff

CONTOUR section

A single AD (attack–decay) envelope generator, similar to the contour generator in a Buchla 281. It drives both the LPG filter and VCA according to the Mode setting above.


Parameter	Description
Mode	Pluck — decay is gate-independent (fires and fades regardless of note length); Sus — holds at peak until the note ends, then decays
Attack	Rise time (1 ms – 2 s)
Decay	Fall time (5 ms – 4 s) — in Pluck mode this is the T60 envelope; in Sus mode this is the release time after gate-end
Amount	Peak level of the contour (0–1) — scales how far the LPG opens
Cycle	On — re-triggers the AD shape repeatedly while the note is held, turning the contour into a free-running LFO-like modulator

AMP section

Parameter	Description
Level	Master output gain (0–1)
Tune	Global pitch offset in semitones (± 12 st)
Voices	Polyphony, 1–16 (default 8)
Mode	Poly or Mono

The engine ships with **24 presets** — spanning percussive bass plucks (BASS Fold Sub, BASS Growl FM), bells (BELL Metallic, BELL Crystal Ring), pads (PAD Fold Drone, PAD Glass Air), keys (KEYS Fold E-Piano, KEYS Marimba Fold), and abstract textures (FX Cycle Burst, FX Sci-Fi Cycle) — selectable from the lane's preset dropdown. Per-section knob accent colours group the COMPLEX OSCILLATOR, TIMBRE, LOW-PASS GATE, and CONTOUR sections visually.

Sampler



Above: Sampler editor — global Gain/Voices controls; per-pad controls appear once samples are loaded.

The Sampler engine plays back audio samples mapped across the keyboard. Each keymap zone (pad) has its own per-pad parameters read at trigger time, making it possible to tune, filter, and pan individual pads independently.

Global parameters

Parameter	Description
Gain	Master output gain for the lane
Voices	Polyphony cap (1–16; default 8)


Per-pad parameters

Shown in the drum-voice rack once pads are mapped.

Parameter	Description
Tune	Transposition in semitones (–24 to +24)
Cutoff / Res	Per-pad resonant low-pass filter
Attack / Decay	Per-pad amplitude envelope
Level	Per-pad output level
Pan	Stereo position
A / B	Send amounts to Send A (delay by default) and Send B (reverb by default)
Loop / Loop Start / Loop End	Loop mode (one-shot or loop-while-gated) and the loop region (start/end as a fraction of the sample)
Sample Start / End	Trim the played window — start/end as a fraction of the sample; draggable on the waveform in the Selected sample panel
Retrig	Poly (voices overlap) or Mono (re-hit cuts the previous voice)

A Sampler lane with pads mapped to GM drum note numbers automatically enters drumkit mode and shows the drum-grid editor instead of the piano roll. For details on loading samples and building keymaps see [MIDI & Samples](#).

Drums (Drum Machine)



Above: Drums editor — Preset row, master bus knobs (Vol/Pan/Rev/Dly/Lo/Mid/Hi), per-voice rack with Tune/Decay/Rev/Dly controls per voice.

The Drums engine is a fully synthesised **ten-voice** drum machine — no samples required. All ten voices (Kick, Snare, **Rimshot**, Closed Hat, Open Hat, Clap, Cowbell, Tom, Ride, **Crash**) are built from oscillators, noise generators, and simple envelopes. Each voice routes through its own channel strip and then to a shared drum bus.

Voice synthesis rack

The per-voice rack exposes the key parameters for each voice:

Voice	Key parameters
Kick	Tune, Attack (click), Decay, Start/End Freq, Sweep, Wave
Snare	Tune, Tone body, Snap (noise), Body Decay, Noise Decay, Noise Tone
Closed Hat	Tune, Decay
Open Hat	Tune, Decay
Clap	Tone, Decay, Sharp (filter Q)
Tom	Tune, Decay, Sweep, Start/End Freq
Cowbell	Tune, Decay, Detune
Ride	Tune, Decay

Drum preset dropdown — synth kits and sample kits

The preset dropdown for any drum lane lists kits from four groups:

Group	Kits	How it works
GM	KIT Standard, KIT Room, KIT Power, KIT Electronic, KIT TR-808, KIT Jazz, KIT Brush, KIT Orchestra	GM-programme aliases that map to the synth kits below
Synth	TR-909, TR-808, TR-606, CR-78, LinnDrum	100% synthesised DSP — no samples required
Samples	TR-808 (samples), Acoustic (samples), Dirt (samples)	Real one-shot WAVs bundled with Loom
Drum Machines	64 sampled kits from classic boxes (Roland, LinnDrum, Korg, Oberheim, Casio, E-mu...)	Sampled one-shots from the tidal-drum-machines collection — a large library of vintage drum-machine sounds

Synth kits seed every voice's parameters from the kit's characteristic values; you can edit individual voices on top and hit 🎲 Sound to randomise all voices at once.

Sample kits load the matching WAV for each voice (kick, snare, closed hat, open hat, clap, tom, cowbell, ride) from `public/drumkits/` and rebuild the keymap fresh on every session load — you never need to re-import the files manually. Once a sample kit is selected, the lane uses the drum-grid editor and the full per-pad parameter rack, exactly like a Sampler lane in drumkit mode.

The sample WAVs are curated one-shots from the Dirt-Samples collection (used by TidalCycles), classic TR-808 recordings, and — for the **Drum Machines** group — the [tidal-drum-machines](#) library. Full credits are in the repo `README.md` under "Credits — sample sources".

Note: sample kits are loaded by the Sampler engine under the hood. For the Sampler's own Instrument family selector (Drumkit / Melodic / Loop) and per-pad parameters, see [MIDI &](#)

Bus controls

The master bus row gives Vol, Pan, Rev, Dly, Lo, Mid, and Hi (± 18 dB shelves) for the whole drum bus. These are automatable via [Modulation & Note FX](#). Routing to the shared reverb and delay sends is covered in [Mixing & FX](#).

Choke groups

Each drum voice has a **Choke** dropdown in the voice's advanced area of the per-voice rack. Voices assigned to the **same non-zero choke group are mutually exclusive**: triggering one immediately cuts the tails of all other voices in the same group. A voice assigned to group 0 (the default) is never choked by anything.

The default configuration is **closed hat and open hat both in group 1**, which gives the classic hi-hat choke behaviour — hitting the closed hat silences the open hat's ring, just as on a real drum kit. You can reassign any voice to any group number to create other exclusive pairs, for example a conga and cowbell that cannot overlap.

Summary table

Engine	Best for	Standout parameters
TB-303	Acid bass lines, resonant leads	Slide, Accent, Env Mod
Subtractive	Pads, leads, basses, general-purpose	Dual OSC detune, Filter Drive, Key Track, POLY mode
FM	Bells, electric pianos, metallic textures	Algorithm, per-operator Ratio, FB feedback
Wavetable	Evolving tones, digital leads, pads	Morph (A→B crossfade), 8-waveform bank
Karplus-Strong	Plucked strings, guitar, harp	Damping (sustain), Brightness, Excitation
West Coast	Percussive plucks, metallic tones, evolving textures	Wavefolder (Fold), LPG Mode, Contour Cycle
Sampler	Any audio, drum kits with per-pad control	Per-pad Tune/Filter/Envelope, Loop mode
Drums	Synthesised drum machine, percussion	8-voice synth rack, kits-as-presets, bus EQ, Choke groups

Editing Clips

Every clip in Loom holds a sequence of notes. To edit those notes, click the **body** of any filled cell in the session grid — anywhere except the ► play icon or the ✕ delete cross in its corner (clicking ✕ deletes the clip outright, with no confirmation). You can also **right-click** a filled cell and choose **Open editor**. The inspector panel opens below the grid and the editor renders inside it. Closing the inspector does not stop playback — launching and editing are independent.

Melodic lanes (TB-303, Subtractive, FM, Wavetable, Karplus, West Coast, Sampler) open the **piano-roll**. Drum-machine lanes and sampler lanes that have a drum kit loaded open the **drum-grid**. If you want to switch between the two views for a given clip, click the ↔ **Editor** button in the inspector toolbar.

See [Sessions](#), [Lanes](#), [Clips & Scenes](#) for how clips are organised, and [Engines](#) for the controls each engine exposes.

Key, scale & musical assistance



The project key/style bar sits in the transport strip next to Meter; the 🔒 Scale button and the 🎲, Vary, Mirror, Reverse, and Chords controls appear in the clip-editor toolbar.

Loom's musical-assistance layer works from a single shared **project key and style** setting. Everything that generates or reshapes notes — the random generator, the examples gallery, the pattern transforms, the chord harmoniser — reads that setting and stays in key automatically. You can override it on a per-lane basis when one part needs to inhabit a different harmonic world.

Project key & style bar

In the top transport bar, next to the **Meter** control, a button displays the current project tonality and the scale-lock state — for example 🎵 **A minor · Acid / Techno** · 🔒. The 🔒 / 🔓 glyph at the end tells you at a glance whether the scale lock is on, without opening anything. Click the button to open the key/style panel.

The panel has three pickers and a lock toggle:

- **Root** — the tonic note (C through B).
- **Scale** — chosen by feel rather than by musical-theory name. Each option shows a mood label, the musical name in small print, and a one-line usage hint:
 - 🌙 Dark / tense — minor · the classic acid/techno sound
 - 🛡️ Safe (almost anything fits) — pentatonic minor
 - 🌞 Bright / uplifting — major
 - 🌀 Hypnotic / modal — dorian

- 🧐 Unsettling — phrygian
- 🎨 Anything goes — chromatic (no constraint)
- **Style** — the rhythmic and harmonic character used by the generators. There are **20**: Techno, Acid Techno, Dub Techno, House, Deep House, Garage, Trance, Psytrance, EDM, Breakbeat / Big Beat, Drum & Bass, Jungle, Dubstep, Electro, Synthwave, IDM, Glitch, Downtempo, Lo-fi and Ambient. Style affects which rhythmic patterns the chord maker lays down, sets the default feel of generated basslines and melodies, and selects which library of ready-made patterns the **Pattern** dropdown offers.
- **Scale lock** — a checkbox that turns the piano-roll's note snapping on or off **globally**. It is **off by default**, so you can play freely; tick it when you want every placed note pulled into the key. This is the same switch as the piano-roll's 🔒 button — wherever you toggle it, the other reflects it.

Every other musical-assistance feature — the piano-roll scale highlight, the generator, the examples, the transforms, and the chord maker — respects the active root and scale. Closing the panel without choosing commits no change.

Per-lane override

The lane inspector shows a line such as **Key: inherits A minor · Override**. Click **Override** to assign a different root and scale to that lane alone, so one part can be in a different key while the rest of the session stays in the project key.

Scale highlight & lock (piano-roll)

Inside the piano-roll, in-scale pitch rows are highlighted and the keyboard strip on the left colours in-scale keys — the tonic row and key are the brightest. This gives you an instant visual map of which notes belong to the current key.

A 🔒 **Scale** button in the piano-roll toolbar controls the scale lock, mirroring the **Scale lock** checkbox in the project key/style panel — both write the same global setting. The lock is **off by default**: a fresh session never constrains what you play, and a saved session always loads unlocked, so the lock can never surprise you by being on. Switch it to 🔒 whenever you want a safety net. When locked:

- Drawing a note with the pencil snaps the pitch to the nearest in-scale degree.
- Dragging an existing note vertically skips over out-of-scale rows.
- Computer-keyboard note input plays and records only in-scale pitches.
- Paste lands notes on in-scale pitches, transposing each pasted note by the minimum interval needed to put it back in key.

With the lock on it is impossible to place a wrong note; toggle it back to 🔓 (the default) for chromatic freedom — accidentals, blue notes, deliberate dissonance. The arrow-key nudge (semitone up/down) intentionally ignores the lock in both states, so you can always make precise micro-adjustments with the keyboard.

Generate (🎲)

The 🎲 (dice) button in the clip-editor toolbar fills the current clip with a freshly generated musical part in the project key, scale, and style:

- **TB-303-style lanes** — a characteristic bassline (slide and accent placement influenced by the active style).
- **Drum lanes** — a beat pattern matched to the style's rhythmic vocabulary.
- **Other melodic lanes** — a melody or counterpoint line.

The result is not flat random: it is constrained to the scale, shaped by the style, and proportioned to the clip length. Press 🎲 again for another variation; change the global **Style** in the key/style panel for a different character while keeping the same root and scale.

Pattern library

Two dropdowns in the toolbar work as a pair:

- **Style** picks a library — one of the 20 styles listed above. Note that changing it here also changes the **project** style, so the 🎲 generators follow suit.
- **Pattern** lists that style's ready-made patterns, roughly twenty per style. The list is a native dropdown, so you can **type to search** it instead of scrolling.

The list is filtered by editor type — melodic editors offer basslines and riffs, drum editors offer beats — and your own saved examples appear in their own group at the end.

Selecting a pattern:

1. Loads it into the current clip.
2. Transposes it from its stored key into the **current project key**.
3. Repeats or trims it to exactly fill the clip length — short riffs tile; long riffs are cut at the clip boundary.

Two toolbar buttons sit alongside:

- **+ Example** — saves the current clip's notes as your own example. It is stored in this browser (IndexedDB), persists across reloads, and is marked with a ★ in the list.
 - **↓ JSON** — exports the current clip as a portable `.json` file that you can share or import into another browser.
-

Pattern transforms

Three buttons in the toolbar reshape the current pattern — all results stay within the project key and scale.

Button	Name	What it does
Vary	Musical variation	Nudges some pitches to neighbouring scale degrees and lightly adjusts rhythm and velocity; the overall character and contour are preserved.
Mirror	Melodic inversion	Flips the melodic contour so that rising lines fall and falling lines rise, mirrored around the first note of the clip.
Reverse	Retrograde	Plays the pattern backwards in time — the last note becomes the first.


All three are undoable with Ctrl+Z / Cmd+Z.

Chords

The **Chords** button harmonises the current melodic clip. For each bar it picks the most characteristic **diatonic triad** implied by the bar's melody notes, then writes an accompaniment using a **rhythm pattern drawn from the global style** — for example offset stabs in House, sustained pads in Lo-fi, driving eighth notes in Synthwave.

When you click Chords a small dialog asks **which lane** to write the chord part to: pick an existing melodic lane or let Loom create a new lane labelled "Chords". The original clip is not modified.

Piano-roll



The piano-roll is a two-axis canvas: time runs left-to-right, pitch runs bottom-to-top. A vertical keyboard on the left names the rows; a time ruler at the top marks bars and beats.

Zoom and pan

Scrub **vertically on the time ruler** to zoom the time axis; scrub **horizontally** on the ruler to pan. Scrub the **keyboard strip** vertically to zoom the pitch axis. The native scroll bars pan both axes. Zoom state is saved per clip and restored when you reopen it.

A plain **click on the time ruler** (without dragging) **seeks** the song to that point: the global playback position jumps there and every lane re-anchors in phase, so while the transport runs you can jump around the timeline from any clip editor. Dragging the ruler still zooms/pans as above — only a click-without-drag seeks.

The **drum-grid** and the **audio-clip editor** have the same horizontal zoom and scroll: their fixed parts (the voice-name labels on the drum grid, the waveform header on an audio clip) stay pinned on the left while the timeline content scrolls and zooms beside them. The loop brace tracks the zoom in every editor.

Follow (auto-scroll)

A session-global **Follow** toggle keeps the playhead in view: while it's on, every clip editor auto-scrolls horizontally to follow playback. It is **on by default** — turn it off when you want to inspect one region while the transport runs elsewhere. The setting is shared across all editors, but it is a working mode rather than part of the music: it is **not saved with the session** and returns to on when you reload.

Draw mode (pencil)

The default tool is **Draw**. Click an empty area of the grid to place a note at the snap resolution (default: 16th notes). Drag right while placing to extend its duration. Click and drag an existing note's right edge to resize it. Click and drag the body of a note to move it. Alt-click or right-click a note to delete it.

Select mode

Click the **Select** tool button in the toolbar to switch. In Select mode, click a note to select it. **Drag on the grid background** to draw a marquee rectangle; every note whose body intersects the rectangle is selected. Click empty space to deselect all.

With notes selected you can:

- **Move as a group** — drag any selected note; the whole selection moves together and is clamped to the clip boundaries.
- **Delete** — press Delete or Backspace.
- **Nudge** — use the arrow keys to shift the selection one snap unit left/right or one semitone up/down.
- **Cut / Copy / Paste** — Ctrl+X / Ctrl+C / Ctrl+V (Cmd on Mac). Paste anchors the earliest clipboard note at the mouse cursor position, so move the pointer where you want the paste to land before pressing Ctrl+V.

The tool choice and clipboard contents persist across clip re-opens and across clips, so you can copy from one clip and paste into another.

Computer-keyboard note input

The  **Keys** toggle in the clip-editor toolbar turns your computer keyboard into a live instrument. It is **off by default** — switch it on and the keys below play the active lane in real time:

Key row	Notes
a s d f g h j k (home row) — white keys	C D E F G A B C
w e t y u (upper row) — black keys	C# D# F# G# A#
z / x	Shift input octave down / up

This is **live playing, not step entry** — the keys sound the lane exactly as a MIDI keyboard would, and nothing is written into the clip on its own. To capture what you play, arm **• Rec** in the inspector and play along with the loop; see [MIDI live-record](#). (Earlier versions typed notes into the clip at a cursor; that behaviour was replaced by live playing.)

Drum-grid



The drum-grid is a canvas editor where rows correspond to drum voices (kick, snare, hi-hat, etc.) and columns correspond to time positions. Each hit is placed at a precise tick within the clip.

Grid resolution

A resolution selector at the top of the editor sets the snap and the column width. Available resolutions are:

Value	Description
1/4	Quarter notes
1/8	Eighth notes
1/8T	Eighth-note triplets
1/16	Sixteenth notes (default)
1/16T	Sixteenth-note triplets
1/32	Thirty-second notes
free	No snap — place hits at any tick

You can mix resolutions across clips; each clip stores its own `gridResolution` setting.

Euclidean fill (H / S / R)

Every voice row has three small fields between its label and the grid: **H**, **S** and **R**. They generate that voice's pattern by spreading hits as evenly as possible around a cycle — the Euclidean rhythm idea, which produces a surprising number of the world's real drum patterns.

Field	Meaning
H — hits	How many onsets to place, spread as evenly as possible. 0 leaves the voice alone , so you can Euclid one voice without touching the rest.
S — steps	The length of the cycle. Shorter than the clip and it simply repeats; set it to a length that does <i>not</i> divide the clip evenly and the pattern phases , landing differently each time round.
R — rotate	Shifts the cycle so the hits land elsewhere against the beat. Negative values rotate the other way.

Start with **H=4, S=16** on a kick for four-on-the-floor, then try **H=3, S=8** on a hat for the tresillo that underpins most Latin and house grooves. The classic trick is an off-divisor: **S=7** against a 16-step clip never repeats the same way twice within the bar.

The fill is generated into the clip as ordinary notes, so you can edit them by hand afterwards, and the whole operation is undoable.

Placing and removing hits

Click an empty cell to place a hit. Click an existing hit to remove it. In **free** mode, click anywhere in the row — the hit lands at the exact tick under the pointer with no snapping.

Selection and group operations

Drag on the canvas background to draw a marquee rectangle. Hits whose row and time position intersect the rectangle are selected. With a selection active:

- **Move** — drag horizontally to shift selected hits in time; the group is clamped to the clip length.
- **Move rows** — drag vertically to reassign hits to different drum voices; the relative row offsets are preserved and clamped to the available voice list.
- **Delete** — press Delete or Backspace.
- **Cut / Copy / Paste** — Ctrl+X / Ctrl+C / Ctrl+V. Paste anchors the earliest hit at the click position (tick × row), preserving relative offsets within the group.

Selection, clipboard, and group-move all operate on **row indices**, not MIDI numbers, so patterns copy cleanly even between kits that map voices to different MIDI notes.

Playhead

A vertical playhead line moves across the canvas in real time while the clip is playing, driven by the sequencer's look-ahead clock. It updates on every redraw tick and resets to the left edge when the clip stops.

Clip tempo (×2 / ÷2)

Next to the **Length (bars)** field in the inspector toolbar sit two buttons, ***2** and **/2**, that time-scale the open clip's notes in a single, undoable step — the perceived tempo of the pattern doubles or halves while the clip stays musically useful.

- ***2** (*Double tempo — compress notes & repeat to fill, length unchanged*) — every note's start and duration is halved (so the pattern plays **twice as fast**), and the clip's **length in bars stays the same**: the now-shorter pattern is **tiled** to fill it, so you hear a clean double-time version rather than a half-empty clip.
- **/2** (*Halve tempo — stretch notes & grow the clip length*) — the inverse: notes and durations double (so the pattern plays **half as fast**) and the clip's **length grows** to fit the stretched material.

The loop region and any clip automation scale with the notes. The buttons appear for **note and drum clips**; a pure audio channel has no notes to scale, so they are hidden there. The whole rescale is one entry in the global undo history, so a single Ctrl+Z restores the original timing.

Loop regions

A **loop brace** sits above every clip editor — piano-roll and drum-grid — as a narrow strip spanning the full clip length. It lets you mark an A–B sub-region and repeat just that portion while the clip plays.

Setting the loop region

The brace strip has two drag handles (left = A, right = B) and a **Loop** toggle button. To use it:

1. Click **Loop** to enable the loop region. The region highlights between the A and B handles; if no region was set before, it defaults to the full clip length.
2. Drag the **left handle** to move the start point (A). Drag the **right handle** to move the end point (B). Both handles snap to 16th-note grid positions.
3. Drag the **interior** of the region (between the handles) to **slide the whole window** along the timeline — its length is preserved, it snaps to the grid, and it stops at both ends of the clip. Handy for auditioning the same-length loop over different bars.
4. While the clip is playing, the scheduler repeats only the A–B sub-region — the rest of the clip is skipped.

Click **Loop** again to disable it. The clip reverts to playing its full length; the A and B positions are remembered so you can re-enable the same region later.

What the loop brace affects

The loop region works the same way for all clip types:

- **Note clips (piano-roll)** — only notes whose start falls within the A–B range are triggered; the period of repetition equals the duration of that range.
- **Drum clips (drum-grid)** — same: only hits inside the sub-region fire.

- **Sliced note clips (Sampler)** — a clip whose notes trigger bank slices behaves like any note clip: only the hits inside the A–B sub-region fire, still tempo-locked and pitch-preserving.

The brace lives on the note/drum editor. A pure audio channel (the waveform-only audio-clip editor) does not show a brace; to gain per-region control over an audio loop, bring it in through the Sampler's **Loop** family, which slices it into a note clip with a piano-roll (see [MIDI & Samples — Sampler](#)).

The loop region is **per-clip** and saved with the session. Two clips in the same lane or scene can each have their own independent A–B region, or none at all.

Sharing a loop across the scene

By default each clip's loop is its own. The brace strip also carries a **Global** button (tooltip: *"Share this loop across every clip in the scene"*). Turn it on — it lights amber — to push this clip's A–B region onto **every clip in the scene** so they all loop the same bars together, which is handy for auditioning one section across the whole arrangement. Toggle it off to return each clip to its own independent region.

All loop-region edits (moving handles, sliding the region, toggling, and the Global share) are part of the global undo history (Ctrl+Z / Cmd+Z).

For how to use an arrangement-wide A–B loop brace that repeats a section across all lanes at once, see [Performance & Arrangement](#).

Waveform header

Any clip that references an audio buffer shows a **waveform header** above its editor — a peak view of the buffer with a bar/beat ruler, slice markers (orange), and a live playhead. It appears in two situations:

- **Audio clips** open the dedicated **audio-clip editor** (toolbar + waveform header, no note grid).
- **Sliced note clips** (and any normal clip that carries a display-only waveform reference) keep the header **above** the piano-roll or drum-grid, so you can edit notes against the source audio.

To bring a loop into a session as a tempo-locked audio channel, see [MIDI & Samples — Audio channel](#) (the + **Audio** control, the waveform header, the **Warp** tempo-lock). To chop a loop into individually editable note slices, load it through the Sampler's **Loop** family instead.

Velocity & dynamics



Piano-roll with velocity lane

Above: piano-roll editor. The strip beneath the note grid is the velocity lane — one vertical bar per note, height proportional to velocity.

Every note in Loom carries a **velocity** value from 0 to 127. Velocity is set when you draw a note (default: **90**), adjusted in the velocity lane, and captured automatically from MIDI import. It affects the sound in two complementary ways:

- **Loudness** — velocity scales the note's output gain continuously via a smooth curve (`velToGain`). A velocity of 1 is near-silent; 127 is the loudest a note can be. Notes at the default of 90 sit just above the mid-point of the range, so there is clear headroom in both directions.
- **Accent character** — notes with velocity ≥ 100 are accented. On top of the continuous gain, accent adds character to the sound: on bass-style engines (TB-303, Subtractive) it brightens the filter envelope and raises the resonance Q; on drums it increases brightness. This is the same accent model that the 303 bassline and drum sequencer have always used, now unified into the velocity scale.

Reading velocity visually

Each note's **fill colour** shifts along a **blue** → **yellow** ramp as its velocity increases. Low-velocity notes are deep blue; high-velocity notes are warm yellow. The transition is weighted so the blue half of the range covers roughly velocities 0–64 and the yellow half covers 64–127. Accented notes (≥ 100) are additionally outlined with a **white border** — colour alone does not distinguish accent from non-accent.

The velocity lane

Below the note grid (piano-roll) or the drum-voice rows (drum-grid) is the **velocity lane**: a row of vertical bars, one per note, anchored at the note's start position. Bar height is proportional to velocity. A **dashed horizontal line** across the lane marks the accent threshold (velocity 100) so you can see at a glance which notes are accented. The lane scrolls horizontally in sync with the grid.

Editing velocities

You interact with the velocity lane by dragging the bars:

- **Set a single note** — drag a bar up or down. The velocity updates live; the note colour and audible gain change as soon as you release.
- **Adjust a group** — if you have notes selected (marquee selection in the main grid), dragging any bar that belongs to the selection applies the **same delta** to all selected notes. Notes that would go out of range are clamped to 1–127.
- **Paint a ramp** — drag horizontally across multiple bars. Each bar you pass over is set to the velocity corresponding to the current vertical position of the pointer, writing a smooth velocity ramp across the passage in a single gesture.

When several notes share the same start position (a chord), their bars are fanned a few pixels apart in the lane so each one remains individually grabbable.

All velocity edits are undoable (Ctrl+Z / Cmd+Z) in the same undo history as note placement and movement.

Modulation & Note FX

Every lane has two signal-processing layers that sit *before* its engine's audio output and *around* its note stream: **Modulators** shape sounds over time by continuously driving engine parameters, and **Note FX** transform the note events before they reach the engine. Both are per-lane and saved with the session.

The FM engine editor showing the MODULATORS and NOTE FX sections. The MODULATORS section (LFO1 + ADSR1) and the NOTE FX row are visible at the bottom of the panel.

Modulators

Open any lane's engine editor and scroll to the **MODULATORS** section. Press **+ LFO** or **+ ADSR** to add a modulator; you can add as many as you need. Each modulator appears as a card with its own controls, an **ON / OFF** toggle, and a **×** remove button.

The section is not empty to begin with. Every engine arrives with its own modulators already wired, and on some engines they are load-bearing rather than decorative:

Engine	Ships with
Subtractive	Four — two ADSRs (which are the amp and filter envelopes) plus two free LFOs
West Coast	Two ADSRs (wavfolder + low-pass gate) plus two LFOs
Wavetable	An ADSR on filter cutoff
FM, Karplus	One LFO and one ADSR
TB-303	One LFO only — by design; its envelopes are part of the engine

On Subtractive in particular, deleting the two ADSRs removes the amplitude and filter envelopes: they are not extras sitting on top of the sound, they *are* the sound's shape.

A preset can carry its own modulators. Loading one replaces the lane's modulator set with the preset's — because for some patches the modulation *is* the patch. A wobble bass is an LFO on the filter cutoff; without it there is no wobble. Six of the Subtractive presets ship one today: **BASS Wobble LFO**, **BASS Neuro**, **PAD Shimmer**, **PAD Cosmic**, **PAD Ethereal** and **LEAD Sync Sweep**. Load one and look at the MODULATORS section to see how it was built — they are the quickest way to learn what this section can do.

LFO

An LFO generates a periodic waveform that you route to one or more target parameters.

Control	Description
WAVE	Waveform shape: Sine, Tri, Sqr, or Saw
RATE	(<i>FREE mode</i>) the free-running rate, set with a log-scaled knob and shown in bpm (LFO cycles per minute), from ultra-slow sweeps (under 1 bpm) up to audio-rate wobble. The slow half of the knob travel covers the slow rates, so gentle sweeps are easy to dial in.
BARS / FEEL	(<i>SYNC mode</i>) BARS is a free numeric input for the cycle length in bars-per-cycle (e.g. 0.25 = a quarter-bar cycle, 4 = one cycle every four bars; any value from 1/16-bar up to 64 bars). FEEL offsets it — Str (straight), Trip (triplet), Dot (dotted). This replaced the old fixed RATIO dropdown, so you can sync to any cycle length, not just preset divisions.
FREE / SYNC	Toggles between the free-running bpm RATE knob (FREE) and the tempo-locked BARS + FEEL controls (SYNC).
POLARITY	-1..+1 (bipolar, default) oscillates symmetrically around the param's centre. 0..1 (unipolar) only pushes the parameter upward.
TRIG	⚠ Currently has no effect — see the note below. Intended: Free runs the phase continuously; Note resets phase on every note-on.
SCOPE	⚠ Currently has no effect — see the note below. Intended: Shared runs one LFO for the lane; PerVoice gives each played note its own.

TRIG and SCOPE are not connected yet. Since synthesis moved into the AudioWorklet, the only LFO settings that reach the sound engine are the rate, the depth and the destination — the trigger mode and the scope are not sent. Every LFO therefore free-runs and is **shared** across all voices of the lane, whatever these two dropdowns say. They are left in place for the work that will wire them up; until then, do not expect a chord's notes to wobble independently, or an LFO peak to land on each note-on.

The ADSR modulator is unaffected: it is genuinely **per-voice** and works as described below.

ADSR

An ADSR produces a classic Attack–Decay–Sustain–Release envelope that fires once on each note trigger and follows the note's gate duration. Its default scope is **PerVoice**, which is almost always what you want: every note gets its own independent envelope.

Control	Range	Default
A (Attack)	1 ms – 2 s	10 ms
D (Decay)	1 ms – 4 s	300 ms
S (Sustain)	0 – 100 %	70 %
R (Release)	1 ms – 8 s	300 ms

You can switch an ADSR to **Shared** scope if you want a single envelope that re-triggers on each note but whose shape is not tied to any individual voice. In practice PerVoice is the natural choice for most envelopes.

Destinations and depth

Below each modulator card's controls is a destination list. To route the modulator:

1. Select a target from the dropdown — it lists all automatable parameters for the lane's engine, plus any lane inserts, master inserts, and master sends (reverb, delay).
2. Press + **Destination**. A new row appears showing the target name and a **DEPTH** knob (range -1 to +1, default 0.5).

The depth knob scales the modulator's effect. A value of 1.0 means the full output range of the modulator sweeps the full range of the destination parameter (in the parameter's native units). -1.0 inverts the modulation direction. Multiple destinations can share the same modulator at different depth values.

Remove a destination with its ✕ button. Remove the whole modulator with the ✕ on the card header.

Note FX

The **NOTE FX** section sits directly below MODULATORS. Note FX processors intercept the lane's note stream before it reaches the engine, transforming which notes are played and when. Add a processor with + **Arp** or + **Chord**. Each processor shows an **ON / OFF** button and a ✕ remove button.

Note FX are per-lane and persist with the lane's engine state. Loading a demo resets them to the demo's configuration.

Arpeggiator

The Arp takes each held note and generates a rapid sequence of notes from it according to a scale and direction pattern.

Control	Options / Range	Description
PATTERN	up, down, updown, random, cosmic	Direction of the arpeggio. cosmic adds occasional octave jumps and random steps for an unpredictable feel.
SCALE	major, minor, pentMinor, phrygian, chromatic	Scale from which arp notes are drawn, starting from the held root note.
RATE	free, 1/4, 1/8, 1/8t, 1/16, 1/16t, 1/32	Step rate in musical divisions (BPM-synced), or free Hz.
OCT	1–4	Number of octaves the arp spans. Higher values cycle the scale across more octaves before repeating.
GATE	0.05–1.0	Fraction of the step interval during which each arp note is held. Lower values create a more staccato feel.
FREE Hz	0.5–32 Hz	Rate used when RATE is set to <i>free</i> .

The Arp fires as many notes as fit inside the original note's gate duration, so longer notes produce longer runs.

Chord

The Chord processor replaces each incoming note with a chord built on that note as the root.

Control	Options	Description
CHORD	maj, min, maj7, min7, sus2, sus4, dim	Chord voicing to generate
OCT	–2 to +2	Octave offset applied to all notes in the chord

All notes in the chord share the original note's timing and gate length.

Automation

Loom records parameter automation in two ways:

Real-time knob recording (Performance view): in the session/I-O header row, make sure the REC mode selector beside **REC** is set to  **take** (the default), then press **REC** to arm recording and press Play. While recording in take mode, every knob you move *and* every clip launch is captured as automation in the current take. Automation is written at a sub-step resolution derived from the BPM. Press **REC** again to disarm. (The other two REC modes —  **live** and  **offline** — export WAV audio instead of recording automation; see [Transport](#) for the unified REC group.)

Per-clip envelopes (Session view): each clip can carry automation lanes independent of the Performance take system. Open the inspector for a clip and scroll below the note editor to the automation section. Select a parameter from the dropdown and click the add button to create an

envelope lane for that clip. The envelope draws a curve over the clip's length (in bars) and plays back each time the clip loops. Envelopes are stored alongside the clip's notes in the session file.

When you move a clip to a lane running a different engine, envelopes whose parameter ID no longer exists in the new engine are disabled automatically (they remain in the clip but do not play back until re-enabled or deleted).

See also: [Engines](#) for the parameters you can modulate, [Mixing & FX](#) for lane inserts and master FX whose parameters also appear as modulation destinations, and [Transport](#) for the REC button and Performance view.

Mixing & FX

Every lane in Loom has its own signal path from the synthesis engine through to the master output. This chapter explains how that path is structured, what controls are available per lane, and how the shared Master FX panel ties everything together.

Signal flow overview

```
graph LR
    Lane[Lane engine] --> Insert[lane insert chain]
    Insert --> Strip[channel strip]
    Strip --> EQ[EQ]
    EQ --> Comp[comp]
    Comp --> Level[level]
    Level --> Pan[pan]
    Pan --> Mute[mute]
    Mute --> Duck[duck]
    Duck --> Master[master bus]
    Master --> MasterInsert[master insert chain]
    MasterInsert --> MasterComp[master compressor]
    MasterComp --> Output[output]
    Strip --> SendA[Send A gain]
    Strip --> SendB[Send B gain]
    SendA --> Returns[Send A / Send B returns]
    SendB --> Returns
    Returns --> Master
```

In short: the engine's audio passes through any per-lane inserts first, then the channel strip where EQ, sends, pan, and level are applied. The processed signal joins the master bus, which runs through the master insert chain and the master compressor before reaching the speaker. The two **Send A / Send B** buses are parallel return channels: each lane feeds them a post-fader amount, and each send runs its own insert chain (seeded **A = Delay**, **B = Reverb**) before returning to the master.

Per-lane channel strip

Each lane owns a `ChannelStrip`. Its controls are visible below the session grid in the lane's row.



Loom session view — channel strips appear below the clip grid, one row per lane

Level (fader)

The vertical slider sets the lane's output gain as a linear multiplier (0–1 = silence to unity; the percentage label reflects the current value). This is the last gain stage before output, applied after EQ and the per-lane compressor.

Pan

The **PAN** knob positions the lane in the stereo field. Centre (0) is the default; turning it left or right continuously shifts the image. The pan value is automatable and can be modulated — see [Modulation & Note FX](#).

Mute and Solo

M silences the lane by zeroing its mute gain node. **S** solos the lane: all other lanes are muted in the UI while solo is active. Both controls affect what the sidechain tap feeds downstream — a muted lane's tap still carries pre-mute signal so sidechain routing remains stable.

3-band EQ

The three EQ knobs apply before the per-lane compressor:

Knob	Filter type	Centre frequency	Notes
LO	Low-shelf	200 Hz	Boost or cut lows; default ± 0 dB
MID	Peaking	1 000 Hz	Q = 1; adds presence or scoops the midrange
HI	High-shelf	4 500 Hz	Boost or cut highs and air

All three bands are $\pm$ dB adjustments. EQ gain AudioParams are exposed to modulation so you can automate filter sweeps from the modulation panel.

Send A and Send B

The two send knobs — **A** and **B** — control how much of this lane's post-duck signal is fed into the two shared send buses. They replaced the old fixed REV/DLY knobs. Send A and Send B are general-purpose return channels, seeded **A = Delay** and **B = Reverb**, but you can change what effect lives in each — see [Send A / Send B return modules](#). The knobs are independent wet levels: 0 = dry only, higher values mix more of the lane into that send's effect. Sends are post-fader, so they follow the lane's level and sidechain ducking. (Saved knob ids are `mix.<lane>.sendA` / `mix.<lane>.sendB` ; old saves with `...rev` / `...dly` amounts migrate automatically — reverb→B, delay→A.)

Per-lane inserts

Every lane also has a private insert chain that sits *before* the channel strip — the engine's audio passes through it first. If you open the lane's inspector and add FX to its insert list, those effects process the lane signal exclusively and do not affect any other lane.

Inserts vs sends: an insert is a serial in-line processor that the signal passes *through*; a send is a parallel path that taps a copy of the signal into a shared return. Use inserts for tone-shaping a single lane; use sends when several lanes should share one effect (a common reverb space, a tempo-synced delay). Loom no longer privileges any effect — reverb and delay are ordinary inserts too, and they just happen to be the default residents of the Send A / Send B return chains.

The same picker everywhere. Every insert rack — per lane (including audio lanes), on each send return, and on the master — draws from one unified effect picker: **Filter (multifilter)**, **Distortion**, **Reverb**, **Delay**, **Compressor**, **Limit**er, **Tremolo**, **Chorus**, **Flanger**, **Phaser**, and **Bitcrusher**. See [Master FX panel](#) below for each effect's parameters. Any insert's parameters are modulation and Performance-automation destinations, wherever the insert sits.

Master FX panel

The master bus has its own strip at the foot of the **scenes column** of the mixer row — a full column laid out like a lane strip, so it lines up with them: a **MASTER** label, an **EQ** section (HI / MID / LO), an **FX** button (in the lane's SEND slot — the master has no sends), a **PAN** knob, a **Mute** button (no Solo — meaningless on the master), and a fader that mirrors the master **Volume** plus a VU meter. The master EQ, pan and mute shape the whole mix; they are saved with the session and the EQ/pan are undoable like any knob. Click the **FX** button to open the **Master FX panel** below the grid (click again to close it). *(The panel's content — SENDS, MASTER COMP, INSERTS — was previously a separate "Master FX" tab; the controls are identical, only the way you open them changed.)*



SENDS — Send A and Send B

The SENDS section shows the two send buses as **return modules**. Each module is a *simple return* — a **return level**, a **mute**, and an **insert rack** — with no EQ or pan of its own. Whatever effects sit in a send's rack process everything the lanes send into it; the per-lane **A / B** knobs set how much each lane contributes. By default **Send A holds a Delay** and **Send B holds a Reverb**, but you can add, remove, reorder, or replace inserts in either rack from the same picker used everywhere else — so a send can carry a whole chain (say a filter into a delay), not just one effect.

The default reverb and delay expose the parameters below (and behave like any other insert — bypass per slot, modulatable params, etc.).

REVERB parameters:

Param	Range	Description
Wet	0–1.5	Wet output level
PreD	0–0.5 s	Pre-delay before the reverb tail starts
Size	0.05–8 s	Impulse response length (room size)
Decay	0.1–10	Tail decay shape (higher = longer tail)

The reverb is a convolution reverb with a procedurally generated impulse response. Size and Decay rebuild the impulse in real time when adjusted.

DELAY parameters:

Param	Range	Description
Time	0.01–2 s	Delay time. A SYNC toggle on the delay insert locks the time to the project tempo and re-locks whenever the BPM changes
Fbk	0–0.95	Feedback amount
Wet	0–1.5	Wet output level
Damp	200–12 000 Hz	Low-pass filter on the feedback loop; lower values darken repeats

MASTER COMP

The master compressor sits at the tail of the master chain, after all inserts. It uses the same `CompBlock` as the per-lane strip compressor, so the parameters are identical:

Param	Range	Default	Description
Bypass	on/off	on	Pass-through when on
Threshold	–100 to 0 dB	–24 dB	Level above which compression starts
Ratio	1–20	4	Compression ratio
Attack	0–1 s	0.003 s	Gain reduction onset time
Release	0–1 s	0.25 s	Gain recovery time
Knee	0–40 dB	30 dB	Transition softness around the threshold
Makeup	~0–4 (linear)	1	Post-compression gain, up to about +12 dB

The master compressor is bypassed by default. Enable it for glue and loudness control on the final mix, or to tame transient peaks before export. See [Saving & Export](#) for how the master bus feeds the offline render.

INSERTS — the master rack

Below MASTER COMP, the INSERTS section holds the master insert chain. Add a slot from the picker and pick its type. The **same eleven plugin types** are available in every rack — per lane, per send, and here on the master:

Filter (multifilter)

- Type: LP / HP / BP / Notch
- Freq: 20–20 000 Hz (exponential)
- Q: 0.1–24

Distortion (Dist)

- Drive: 0–1 — waveshaper saturation amount (4x oversampled)
- Mix: 0–1 — dry/wet blend

Reverb — same parameters as the Send B reverb above (Wet, PreD, Size, Decay). Use as an insert to reverb the full master rather than via a send.

Delay — same parameters as the Send A delay (Time + SYNC, Fbk, Wet, Damp). Use as an insert for a master-bus slapback or stutter.

Compressor — the same `CompBlock` dynamics compressor as the channel-strip and master compressors, now insertable anywhere. Params: Bypass, Threshold, Ratio, Attack, Release, Knee, Makeup (see [MASTER COMP](#) for ranges).

Limitter — a brickwall limiter (ratio 20:1, hard knee, near-zero attack) for catching peaks. Params: **Ceiling** (the level it will not exceed) and **Release**.

Tremolo (Trem) — an LFO opening and closing the volume. The oldest modulation effect there is: slow and shallow it breathes, fast and deep it chops the sound into a rhythm of its own. Lovely on pads, electric-piano parts and sustained chords.

- Rate: 0.1–12 Hz — how fast it pulses
- Depth: 0–1 — how far it closes between pulses (at 1 it cuts to silence)

Chorus — a delayed copy of the signal, detuned by a slow LFO and mixed back in. The two copies drift in and out of tune with each other, which the ear reads as several players at once. Thickens thin sounds and widens single-note leads without changing their pitch.

- Rate: 0.05–8 Hz — speed of the detuning wobble
- Depth: 0–1 — how far it detunes
- Mix: 0–1 — dry/wet blend (start around 0.5)

Flanger — the same idea as the chorus but with a much shorter delay and its output fed back into itself. Instead of thickening, the copies cancel each other at a comb of frequencies that sweeps up and down — the classic jet-plane whoosh. Feedback sharpens the effect from a gentle sweep to a metallic scream.

- Rate: 0.05–8 Hz — speed of the sweep
- Depth: 0–1 — how far the sweep travels
- Fbk: 0–1 — feedback; the higher, the more resonant and metallic
- Mix: 0–1 — dry/wet blend

Phaser — four all-pass filters whose corner frequencies an LFO sweeps together. They shift phase without changing volume, so mixing them back with the dry signal creates notches that slide through the sound. Similar in spirit to the flanger but smoother and more liquid — a staple on electric pianos, funk guitar and pads.

- Rate: 0.05–8 Hz — sweep speed
- Depth: 0–1 — how far the notches travel
- Fbk: 0–1 — feedback; deepens the notches
- Mix: 0–1 — dry/wet blend

Bitcrusher (Crush) — digital degradation. Reducing the bit depth quantises the waveform to a coarse staircase, adding the gritty, harmonic distortion of early samplers and game consoles; the tone control then dulls the result the way a low sample rate would. At extreme settings it destroys the sound completely, which is often the point.

- Bits: 1–16 — amplitude resolution; 16 is nearly clean, 1 is a square-wave wreck
- Tone: 200–20 000 Hz — a lowpass for the lo-fi dullness
- Mix: 0–1 — dry/wet blend

Slots in the chain are ordered in series: the output of each slot feeds the input of the next. Each slot has a bypass toggle so you can A/B it without removing it. Individual slots can be removed; adding the same type multiple times is allowed.

The master insert chain processes the full mixed signal, after the EQ/pan of the master strip and before the master compressor — distinct from the Send A/B returns, which receive per-lane amounts and sum back into the master independently.

XY pad

The  **XY** button in the session bar opens a floating pad — a Kaoss-style controller for playing two parameters at once with a single gesture.

It stays out of the way rather than taking over the screen: the panel floats above the interface and is **not modal**, so the knobs it drives stay visible and you can watch them move while you sweep.

Assign the axes. Below the square surface are two dropdowns, **X** and **Y**. Each lists every automatable parameter in the session, grouped by lane — the same set of destinations an LFO or ADSR can target: engine parameters, channel-strip and insert controls, anything with a knob. Pick one per axis; leave an axis on — *none* — to make the pad one-dimensional.

Play it. Drag anywhere on the square. Horizontal position sets the X parameter, vertical the Y, each mapped across that parameter's full range — bottom-left is both at minimum, top-right both at maximum. The dot follows your finger, and the assigned knobs turn with it, because the pad drives them through exactly the same path your mouse would.

The classic pairing is a filter's **cutoff on X and resonance on Y**: one diagonal sweep and you have the filter gesture that defines most electronic music. But any two continuous parameters work — delay time against feedback, distortion drive against mix, two engine parameters from different lanes at once.

Two things to know:

- Assignments are **not saved** with the session — the pad starts empty on each reload.
 - Pad moves are **not captured** by Performance take recording; the automation recorder only follows knobs you turn directly.
-

Sidechain compression

Loom includes a sidechain ducking system. Any lane's channel strip can be ducked by the signal level of another lane (the *source*). The ducker follows the source lane's envelope via a full-wave rectifier and two smoothing filters, then reduces the target lane's gain proportionally:

$$\text{duckGain} \approx 1 - \text{depth} \times \text{env}(\text{source})$$

Sidechain parameters (set per lane in the lane inspector):

Param	Range	Default	Description
Source	lane selector	—	Which lane's post-mute tap drives the duck
Depth	0–1	0.6	How deep the gain dips at full envelope
Attack	s	0.005 s	How quickly the ducker opens (gain rises)
Release	s	0.25 s	How quickly the ducker closes (gain falls)
Threshold	dB	–40 dB	Source envelope must exceed this to duck at all

A typical use case is kick-drum ducking: set a bass or pad lane's sidechain source to the kick lane. Every kick hit momentarily ducks the bass, creating a pumping effect common in electronic music. Because the tap is taken from the source lane post-mute (but pre-duck), muting the source stops the ducking without feedback loops.

The sidechain bus is separate from the compressor block available on each channel strip. The per-lane compressor (`CompBlock`) is a standard dynamics compressor in the signal path; the sidechain ducker is a parallel envelope-follower that modulates gain. Both can be active simultaneously.

For engine-level sound design that feeds the channel strips, see [Engines](#). For LFO and ADSR modulation of EQ, sends, pan, and other AudioParams, see [Modulation & Note FX](#).

MIDI & Samples

This chapter covers two ways to bring external material into a Loom session: importing a Standard MIDI File to populate lanes and clips automatically, and loading audio samples into the Sampler engine to build melodic instruments or drum kits.

MIDI Import



The MIDI Import panel lives in the toolbar at the top of the screen, next to the demo picker. Click **MIDI Import** to expand it.

Loading a file

Click the file picker and choose a `.mid` or `.midi` file. Loom reads the file immediately — no server, no upload. The parser extracts every track's name, General MIDI programme number, and all note-on/note-off pairs (converted to start tick, duration, MIDI note, velocity, and channel). Only the first tempo event in the file is used; if no tempo is present, the session BPM stays as it is.

Empty tracks (those with no note events) are silently skipped.

The track list

After parsing, a row appears for each non-empty track showing:

- A **checkbox** to include or exclude that track from the import.
- The **track index**, track name (cleaned of control characters), note count, pitch range, and programme number.
- A **preset dropdown**. The GM programme number is looked up against every engine's preset catalogue, and the matching presets are offered at the top of the list (e.g. programme 33 "Acoustic Bass" will surface TB-303 and Subtractive bass presets). Below a divider, every preset from every engine is available so you can override the suggestion freely.
- A ► **audition button** that plays a short three-note arpeggio through the currently selected preset without touching the session. Use this to compare options before committing.

Importing

Click **Import MIDI**. A confirmation dialog asks whether to **Add** or **Replace**:


- **OK (Add)** appends the imported lanes and a new scene to the current session. The new clips are placed on the new scene row so they line up correctly with the scene's launch button.
- **Cancel (Replace)** clears the session and seeds it with only the imported content.

Loom creates one lane per selected track. The lane's name is taken from the matched preset (e.g. "TB Bass"), while the clip inside it keeps the original track name. If the file contained a tempo, the session BPM is updated to match. The import launches the new scene immediately.

The conversion scales MIDI ticks to Loom's internal grid (based on quarter notes divided into four 16th steps), so the notes land on the correct beats regardless of the file's tick resolution.

See [Engines](#) for what each engine sounds like, and [Sessions, Lanes, Clips & Scenes](#) for how to rearrange lanes and scenes after import.

Sampler



The Sampler is a polyphonic playback engine that maps audio across the keyboard and plays it back at the correct pitch per note. A Sampler lane can be three things — a **melodic instrument** (samples spread across ranges of keys), a **drumkit** (one sample per key), or a **loop** (a sliced loop played as a sequence of notes) — and all three share **one inspector layout**.

Presets are the instruments

The Sampler's bundled instruments — drumkits, melodic instruments, and loops — *are* its presets: pick one straight from the standard **PRESET** dropdown at the top of the inspector (grouped **Drumkit / Melodic / Loop**). Selecting one downloads + decodes its WAVs into IndexedDB (so it persists across reloads) and loads it. There is no separate picker. Loom ships **68 ready-made sample kits**, melodic instruments, and loops — all working on the live deploy without a manual import. Most of the kits are grouped under **Drum Machines**: one-shot sets modelled on classic hardware (TR-808, LinnDrum, MPC60, DMX and dozens more), alongside a few hand-curated kits (**TR-808, Acoustic, Dirt**) and a 31-pad **GM Percussion** kit.

The channel layout

Whatever you load, the inspector has the same shape:

- A **keyboard map** across the top — coloured markers at each sound's key (drumkit / loop) or coloured range bands (melodic) — with **connector lines** down to the strips.
- A row of **channel strips**, one per sound. Each strip carries its **name** (the GM voice for a drumkit — KICK, SNARE...; the note for a melodic zone; **Slice N** for a loop slice), its trigger **key**, a ► **play** button to audition the sample and a ✕ to delete that sound (a row at the top of the strip, above the knobs), per-sound knobs (TUNE / CUTOFF / DECAY / LEVEL / REV / DLY plus an ► **adv** block), and **M / S** mute-solo. A + tile at the end of the row adds a sound.

- A **Selected sample** editor below — click any strip and it shows that sample's waveform on a canvas with a **- / + zoom**, the filename, the key, and a **one-shot / loop badge** (click to toggle). Drag the handles over the waveform to set the **trim** (amber start/end lines — the dimmed regions are silenced) and, when loop is on, the **loop region** (green start/end lines that cycle while a key is held). For a melodic zone, the **root / lo / hi** keys are editable here too.

Every sound is keyed by its own note, so each pad / zone / slice has its **own** parameters and mute-solo — nothing is shared between sounds.

Building your own

Import samples... (multi-select) decodes each file, stores it in IndexedDB (so it persists across reloads), and adds it as a new melodic zone spanning the full keyboard with the root at middle C (MIDI 60); narrow each zone's root and low/high boundary in the sample editor. On a drumkit, the **+** tile clones the last pad onto the next free key.

Keymap and repitch

A sound's **keymap entry** has a root note and a low/high key range. When a note falls in the range, the sample plays at a rate of $2^{((\text{midi} - \text{rootNote}) / 12)}$ — equal-temperament repitching — so a melodic instrument's zones span the keyboard chromatically while a drumkit pad or a loop slice sits on a single key.

Per-sound parameters

Every sound — a drumkit pad, a melodic zone, or a loop slice — has its own set of parameters, all read at trigger time:

Parameter	Range	Default	Description
TUNE	−24 to +24 st	0	Pitch offset in semitones, applied on top of keymap repitch
CUTOFF	0–1	1	Lowpass filter cutoff (0 ≈ 60 Hz, 1 = fully open)
RES	0–1	0	Filter resonance
ATTACK	0.001–2 s	0.005 s	Amplitude envelope attack time
DECAY	0.005–4 s	0.08 s	Release tail after the gate closes
LEVEL	0–1.5	1	Pad output level
PAN	−1 to +1	0	Stereo pan position
REV	0–1	0	Send level to the lane's reverb insert
DLY	0–1	0	Send level to the lane's delay insert
LOOP	Off / On	Off	When On, the sample loops while the gate is held
LSTART	0–1	0	Loop start point as a fraction of sample duration
LEND	0–1	1	Loop end point as a fraction of sample duration
START	0–1	0	Trim in — playback start as a fraction of sample duration
END	0–1	1	Trim out — playback end as a fraction of sample duration
RETRIG	Poly / Mono	Poly	Mono cuts the previous hit on re-trigger; Poly layers them

These live in each channel strip — drumkit, melodic, and loop all share the same per-channel rack (there is no eight-pad limit and no separate knob row).

Drumkits

A drumkit is one sample per key, edited on the **drum-grid** (the same grid as the Drum Machine engine). A kit holds **any number of sounds** — not just eight — and you grow or shrink it with the **+** tile and each strip's **X**. Loom ships **68** ready-made sample kits — the **Drum Machines** group (classic-hardware one-shot sets), three hand-curated kits (**TR-808**, **Acoustic**, **Dirt**) and a 31-pad **GM Percussion**; pick one from the PRESET dropdown and the lane is ready to play. The kit is rebuilt fresh from its manifest on each session load, so the WAVs never need re-importing.

Loops



A Sampler Loop instrument: the whole-loop overview above the per-slice strips


A loop is sliced into segments mapped to **consecutive notes**, so playing the keys in order replays the groove. Above the strips the **whole loop** is drawn as one continuous, colour-coded waveform with a cut line at every slice. Selecting a loop also **drops a note clip** onto the lane — one note per slice, each

placed at its exact position so the notes form a continuous **staircase** — opened in the piano-roll; hit Play and it replays the loop, now as discrete notes you can move, mute, repitch or re-order. The clip keeps the loop's waveform as its header.

See [Editing Clips](#) for drawing patterns in the drum grid, and [Engines](#) for the Drum Machine engine (which lists every kit — synth and sample — in a unified preset table).

Audio channel

The **audio channel** is the first-class way to bring a finished loop into a Loom session: drop a WAV and it plays **tempo-locked to the project without changing pitch**, with its waveform shown as a header above the clip editor. It stays a pure audio loop; to chop a loop into individually editable note slices, load it through the Sampler's **Loop** family instead (see [Sampler](#)).




Creating an audio channel

There are two ways to add one:

- **+ Audio button** — at the end of the lane tab row, next to the engine picker, sits a **+ Audio** button. Click it and pick a WAV. Loom decodes the file, stores it in IndexedDB (so it survives a reload), estimates its original tempo, and creates a **new audio lane** holding the loop as an audio clip. The clip opens automatically in the inspector.
- **Drag onto an audio-lane cell** — once an audio lane exists, you can drag another WAV directly onto one of its grid cells to place a second audio clip there.

Each new audio lane gets a launch button on its scene row, so it is immediately playable alongside the rest of the session.

The audio-clip editor



An audio clip has no note grid. Clicking it opens the **audio-clip editor** — a **⚙ Warp ON / OFF** toggle (tempo-locking, see below) above a **waveform header**.

The waveform header shows a peak view of the buffer with a bar/beat ruler, any detected slice markers (orange), and a live playhead while the clip plays. This same header also appears **above** the normal piano-roll or drum-grid for any clip that references a buffer, so you always see the audio you are editing against.

When the transcription helper is configured, the audio-clip header also shows a **TRANSCRIBE** row with a **melodic/drums** toggle and a **Transcribe loop** button. This sends the clip's effective loop region to the local `/transcribe` pipeline and creates a fresh lane with an editable note/drum clip from that region.

Tempo-lock (Warp)

With **Warp ON** (the default), the audio channel plays in time with the project BPM using a pitch-preserving WSOLA time-stretch:

- **At the loop's native BPM** the stretch ratio is ≈ 1 , so playback is essentially identical to the source file.
- **At any other project BPM** the buffer is time-stretched to fit — faster or slower — **without changing pitch**. The stretched buffer is cached and re-rendered automatically whenever you change the project tempo, so the loop stays locked as you experiment.

With **Warp OFF** the loop plays at its natural speed with no tempo sync — useful when you want the audio exactly as recorded.

First-play note: on the very first loop iteration after import (before the stretch cache is warm) playback briefly falls back to a varispeed render — a slight pitch shift that self-heals from the next iteration. At the loop's native tempo the ratio is ≈ 1 , so even that first pass is near-identical.

Slicing a loop into notes

The audio channel itself is a pure WAV loop. To chop a loop into individually editable hits, load it through the Sampler's **Loop** family (see [Sampler](#)) rather than the audio channel. Loom detects slice points (from embedded **Acid / cue / AIFF** markers when present, or by onset detection plus a tempo estimate), stores one short sample per slice in IndexedDB, and creates a **note clip** that triggers the slices in order on a piano-roll — so the groove plays back identically, now as discrete, editable notes. Move, mute, repitch, or re-order the hits in the piano-roll, and tweak each slice's tune/cutoff/decay/level/pan in its channel strip (see [Per-sound parameters](#)); the clip keeps the original waveform as its header. (*Earlier builds did this from a **< Slice** → **pads** button on the audio channel; that moved to the Sampler's Loop family.*)

Editable warp markers

The audio-clip editor: the Warp toggle and the waveform header

Amber warp markers appear as vertical lines in the waveform header; drag one to re-anchor that beat to the grid.

An audio clip can carry **warp markers** — sparse reference points that pin specific audio beats to exact positions on the session grid. With markers in place, the WSOLA time-stretcher locks each inter-marker region independently, so the groove tracks the grid precisely rather than relying on a single global stretch ratio.

Auto-seeding on stem import. When you separate a track into stems (see [Stem separation](#)), Loom analyses the drums stem (the most rhythmically reliable) and seeds quarter-bar warp markers automatically across all stems of that import. The drums stem becomes the **warp reference** for the group; its clip is the one with a fully editable marker overlay in the waveform header.

Editing markers. Open the drums-stem clip in the inspector. In the waveform header you will see amber vertical lines at each marker. You can:


- **Drag** a marker left or right to realign that beat to the correct grid position.
- **Double-click** an empty area of the header to add a new marker at that audio position.
- **Double-click** an existing marker to delete it.

Edits to the reference clip's markers propagate automatically to the other stems that share the same warp group — move one marker and all four audio clips update together.

Warp pill. An amber **Warp** pill in the clip-editor toolbar toggles warping on and off for the clip. Warping is pitch-preserving (WSOLA) and re-renders automatically whenever the project tempo changes, so stems stay locked to the grid as you experiment with BPM.

Stem separation (optional, local service)

Stem separation lets you drop a finished song into Loom and get it back as four separate Sampler lanes — **Vocals**, **Drums**, **Bass**, and **Other** — so you can mute, solo, and remix each part inside the existing session.



How it works

Click **≡ Stems...** in the session bar (the second header row, alongside Save / Load / MIDI). A dialog titled "Separate into stems" opens and immediately checks whether the local helper service is reachable:

- **Service found** — the hint line reads "4 tracks (Vocals / Drums / Bass / Other) via the local service." and the **Separate** button becomes active once you pick a file.
- **Service not found** — the hint reads "Can't find the stems service at localhost:8765. Is it running?" and Separate stays disabled. Start the service (see below), then re-open the dialog.

To separate a track: pick an audio file with the file picker, then click **Separate**. The dialog shows a progress bar:

1. **"Uploading..."** — the file is being uploaded to the local service.
2. **"Separating... m:ss"** — the service is running Demucs; the counter shows elapsed time. The bar may be indeterminate if the model does not report fine-grained progress.
3. On success the dialog closes automatically and four new **audio lanes** appear in the session — one per stem (**Vocals, Drums, Bass, Other**). Each lane holds a full-length audio clip downbeat-anchored to bar 1 and tempo-warped to the session BPM. Hitting Play reconstructs the original mix; mute or solo any lane to isolate parts.

The session BPM is updated to the **detected floating-point tempo** of the imported track on import, so the warp ratio starts at exactly 1.0 and the grooves lock to the grid from the first bar.

The entire lane-creation is a **single undo step**, so you can undo all four lanes at once.

Transcribe to notes (opt-in). A **"Transcribe to notes"** checkbox in the Separate dialog (off by default) tells the service to additionally analyse each stem and create a note clip alongside the audio:

- Pitched stems (Vocals, Bass, Other) → melodic note clips, with pitches estimated via pYIN onset detection.
- The Drums stem → a drum-machine pattern, with hits mapped to the Drum Machine engine.

The note clips are placed on a second layer within each stem lane and are immediately editable in the piano-roll or drum-grid. They are a convenience starting point, not a perfect transcription — pitch detection works best on monophonic or sparse material.

Cancelar aborts a running job and frees the temporary files on the service. **Cerrar** closes the dialog (only available when no job is running).

Opt-in nature

The feature is entirely opt-in. If you never start the service, nothing else in Loom changes — the ≡ Stems... button is the only touch point, and it degrades gracefully to a clear "service not found" message.

Stems land in IndexedDB as ordinary sample assets: they survive browser reloads just like any other sample you import.

Setting up the local service

The separation runs on your machine via a small Python service in `tools/stem-service/`. It requires **Python 3.10+** and **ffmpeg** on your PATH.

```
cd tools/stem-service
python -m venv .venv
# macOS / Linux:
. .venv/bin/activate
# Windows:
.venv\Scripts\activate

pip install -r requirements.txt
uvicorn app:app --port 8765
```

The first time you separate a track the service downloads the Demucs `htdemucs` model automatically (several hundred MB). Subsequent runs skip the download. Separation takes roughly **1–2 minutes per song** on CPU; a GPU-enabled PyTorch build is much faster.

Codespaces and custom service URL

If you want to run the service in a GitHub Codespace (a Linux VM with Python), start it there with the same commands above, forward port 8765, and paste the resulting HTTPS URL into the browser console:

```
localStorage.loomStemServiceUrl = 'https://your-codespace-url-8765.preview.app.github.dev'
```

The same override works for any non-default host. CORS for `localhost:5173` (dev), `localhost:4173` (preview), and the GitHub Pages origin is already configured in the service.

Note: Chrome's *Private Network Access* policy may add a preflight request when the Pages version of Loom calls `http://localhost:8765`. The lowest-friction setup is running Loom locally (`npm run dev`) alongside the service.

For full notes on CORS, the HTTP contract, and the Codespaces workflow, see [tools/stem-service/README.md](#).

Saving & Export

Loom keeps everything in the browser. There is no account, no cloud, and no server upload — your sessions live in your browser's `localStorage` and on your own filesystem when you export them. This chapter covers how to save and restore sessions, undo your work, and render a scene to a WAV file.

Sessions in the browser


Every parameter, every lane, every clip, and every scene is part of the current session. Changes take effect immediately; nothing is auto-committed to disk. When you close the tab, the browser preserves the last-saved state via an autosave entry in `localStorage`, so reopening Loom typically drops you back where you left off.

Three buttons on the session bar (the second header row) drive session management:

Button	What it does
 New	Discards the current session and starts a blank, empty one (tooltip "Nueva sesión vacía").
Save	Opens the Save Manager with the name field ready to type.
Load	Opens the Save Manager to browse and restore a saved session.

See [Transport](#) for the full two-row header layout (transport & tempo on top, session & I/O below).

Save Manager



Clicking **Save** or **Load** opens the Save Manager modal. It is the single place where all session persistence happens.

Saving a session

Type a name in the text field at the top and click **Save current** (or press Enter). Loom writes the full session state — BPM, time signature, every lane's engine + inserts + clips + scenes, the mixer state, and the arrangement take if one exists — into `localStorage` under a unique key, and also updates the autosave slot. The entry appears in the list immediately.

The save format is versioned (`schemaVersion: 3`). When you load an older file Loom migrates it automatically; saves with an unrecognised schema version are rejected with a warning rather than loading broken state.

The saved-session list

Each row in the list shows the session name, date/time, and size in KB. Per entry you can:

- **Load** — restores the session and closes the modal.
- ⬇ (download) — exports the entry as a `.json` file to your filesystem without closing the modal.
- ✎ (rename) — prompts for a new name in place.
- 🗑 (delete) — confirms and removes the entry.

The topmost row is **Auto-save (latest)**, which always reflects the state at the time of the last named save.

Load from file...

Imports a `.json` file you previously downloaded (or received from someone else). Loom validates the schema before applying it; an invalid file shows an alert.

Clear all saves

Removes every named entry from `localStorage`. The autosave slot is preserved. A confirmation dialog appears before anything is deleted.

Storage readout

The footer shows the total size of all named saves, so you can keep an eye on `localStorage` usage.

Undo / redo

Loom keeps a global undo history that covers **every** session mutation: adding or removing lanes, clips and scenes, editing notes in the piano roll or drum grid, moving or tempo-scaling clips, changing engine parameters via knobs and faders, and inline renames. You can step backwards and forwards from the header buttons or from the keyboard.

Header buttons. A pair of ↶ **Undo** / ↷ **Redo** buttons sits in the transport bar, just to the right of the Play / Stop controls. Each reflects the current state of the history — it is disabled (greyed out) when there is nothing to undo or redo — so you can see at a glance whether a step is available. Click ↶ to undo the last change, ↷ to redo it.

Keyboard shortcuts.

Action	Shortcut
Undo	Ctrl+Z / Cmd+Z
Redo	Ctrl+Shift+Z / Cmd+Shift+Z or Ctrl+Y

Capture is automatic — Loom snapshots the session after each interaction and coalesces a continuous gesture (a knob or fader drag, a note drag, a marquee move) into a **single** undo step, so one drag undoes in one click. The shortcuts are inactive while a text input has focus, so typing a save name — or an inline rename — never accidentally undoes your work. Loading a session, loading a demo, or starting a new session clears the undo history.

WAV export

The REC group on the session bar — REC button plus the take / live / offline mode selector

WAV export is part of the unified **REC** group on the session bar (second header row), not a separate button. The **REC** button (tooltip "Grabar — el modo se elige al lado") records using whichever mode is selected in the adjacent **mode selector** (`#rec-mode`):

-  **take** (default) — records knob moves and clip launches into a performance take (not a WAV).
-  **live** — records the live master output in real time to a stereo 16-bit WAV file.
-  **offline** — renders the current scene to a WAV file offline (faster than real time).

To export audio, pick  **live** or  **offline** and press **REC**. The two WAV backends behave as described below.

live (real-time WAV)

The real-time backend is the ground-truth render, and it works like a tape machine: **arm** → **play** → **stop**.

Pressing **REC** in this mode does not start anything — it only **arms** the recorder, and the button reads **ARMED**. Capture begins on the next downbeat, whichever way you start the transport (the  button, a clip launch, a scene launch), and the button changes to **Recording...**. It then records **open-ended** until you press Stop.

That is the important difference from the offline render: it is not one pass of one scene. It captures **the whole performance, including scene changes** — launch scene A, let it run, switch to B, drop a clip, then stop, and all of it is in the file. It taps the live master output after every insert, the master compression and the master FX, so what you hear is exactly what lands in the file, including any random variation from a voice that uses it.

A 2-second tail is appended after you stop so reverb and delay repeats are not cut off.

offline (fast WAV render)

The offline backend rebuilds the full audio graph — lanes, inserts, master bus — inside an `OfflineAudioContext`, applies every lane's current sound state, batch-schedules all note events, and renders faster than real time without touching the live session. It shares the same encoder and download step as the real-time path, so the output format is identical.

It loops seamlessly. The offline renderer deliberately renders **two cycles of the scene and gives you the second one**. The first cycle is thrown away because it starts from silence — no reverb tail from the previous bar, no delay repeats in flight. The second cycle inherits all of that, so the file loops without a seam. This is the reason to prefer the offline render for loop material.

For the same reason it appends **no** FX tail: the render is exactly the musical, bar-aligned length. A trailing tail would round up to an extra bar and the loop would drift.

Where it differs from the live sound. The divergence is structural, not random: per-pad FX sends and per-voice drum-strip sends/EQ are dropped offline, and sample-mode drum kits render through the sampler path. Both are approximations of the live per-voice mix, so a kit that leans on per-voice sends will not sound identical. Everything else — presets, clip automation, the worklet engines — is applied exactly as live.

What gets exported

- **One pass of the scene (offline only).** The offline export captures one full iteration of the longest clip across all sounding lanes; shorter clips loop to fill that window. The **live** take has no such window — it runs until you stop it.
- **A scene must be playing (offline only).** Press **• REC** in ⚡ offline mode with nothing playing and Loom shows a brief notice and does nothing; launch a scene first. In ⌚ live mode arming with nothing playing is normal — that is the point, since it starts capturing from the downbeat of whatever you launch next.
- **Switching mode disarms.** Changing between take / live / offline clears whatever the previous mode had armed.

Where the recording goes

When a render finishes — live or offline — Loom asks what to do with it:

- **Download a WAV**, named `loom-take-<timestamp>.wav`; or
- **Insert it back into the session** as a new audio channel, tempo-locked to the project BPM. This is how you bounce a busy scene down to a single audio lane and keep building on top of it.

Cancelling the dialog discards the take and writes nothing.

Live build and GitHub Pages

The public instance of Loom is deployed automatically to <https://ijol.github.io/Loom/> — every push to `main` triggers a GitHub Actions workflow that runs `vite build --base=/Loom/` and deploys the result to GitHub Pages. The standard `npm run build` (base `/`) is for local development or self-hosting on any other path.

Performance & Arrangement

Loom has two main views: **Session** and **Performance**. The Session view is the clip grid you work in most of the time — lanes, clips, and scenes launched on demand. The Performance view is a linear **arrangement timeline**: a fixed song laid out from left to right with a playhead that moves forward through it, starts at the beginning, and stops at the end.

Switch between the two views using the **Session / Performance** toggle in the transport bar (`#mode-toggle`). Switching stops playback. See [Transport](#) for the full transport layout.



Performance view — timeline with clip bands and automation curves

Three ways to fill the arrangement

The arrangement starts empty. There are three ways to give it content.

1. Copy to Performance

The fastest route from a working session to a playable song is the ↕ **Copy-to-Performance** button (`#copy-to-performance`) in the session bar of the header — now an icon-only button with the tooltip "Copiar las escenas a la timeline de Performance". Clicking it calls `arrangementFromSession`, which walks your scenes in order and lays them out as a linear song:

- Each scene becomes one section.
- The section length equals the longest effective clip in that scene (measured in bars). If a clip has a loop sub-region enabled, its sub-region length is used instead of the full clip length.
- Every lane that has a clip assigned for that scene gets a timeline band covering the section.

After the layout is computed, Loom switches you to Performance automatically. **This one is not undoable** — it overwrites whatever arrangement was there, so copy before you start editing bands by hand, not after.

2. Record a take live

You can record the arrangement in real time while you play.

Recording is driven by the unified **REC** control in the session bar, which has three modes selectable beside the ● **REC** button: 🎙 **take** (the default — captures knob moves + clip launches into a performance take, described below), 🎧 **live** (records real-time audio to a WAV file), and ⚡ **offline** (renders the current scene to WAV offline, faster than real time). The steps below assume 🎙 **take** is selected.

1. Make sure the REC mode selector (`#rec-mode` , next to the REC button) is set to 📱 **take** — its default. (The other two modes, 🎧 **live** and ⚡ **offline**, record audio to WAV instead; see [Saving & Export](#).) Then click ● **REC** (`#rec`) in the session bar to arm recording.
2. Stay in Session view and press Play. Recording begins.
3. Launch clips and scenes as you would for a performance. Move any knobs whose automation you want captured.
4. Press Stop. Loom finalises the take: any still-open clip events are clamped to the stop time, durations are computed, and the recorded content appears in the Performance view as timeline bands and automation curves.

If you arm REC and then switch to Performance mode before pressing Play, the arm is cleared automatically (a toast notification appears) because Performance mode drives playback from the arrangement directly rather than from the live session.

3. MIDI import

When you import a Standard MIDI File via the **MIDI Import** panel (see [MIDI & Samples](#)), Loom calls the same `arrangementFromSession` logic after building the session. Because an imported MIDI file produces a single scene whose clips span the full song, the arrangement comes out as one long section per lane — the complete track laid out linearly from bar 1.

The timeline

Once the arrangement has content, the Performance view shows:

- **Toolbar** — Length (bars), Zoom slider, automation brush (Line / Flat), Loop A–B toggle, and a readout showing total bars and BPM.
- **Ruler** — a bar-numbered ruler across the top. When the A–B loop is active, a loop brace with two drag handles sits on the ruler.
- **Clip bands** — one row per lane. Each recorded or copied clip event is shown as a coloured block, labelled with the clip name, spanning the bars it occupies.
- **Automation curves** — below each clip band, any recorded or drawn automation curves appear. Each curve corresponds to one parameter (identified by its ID). You can draw into curves directly using the Line or Flat brush.
- **Master automation** — any automation curves routed to global (master) parameters appear in a separate section at the bottom.
- **Playhead** — a vertical line (`#perf-playhead`) that moves in real time via `requestAnimationFrame` while the arrangement is playing.

Bands are fully editable by hand — you do not have to re-record to change the layout:

- **Move** — drag the body of a band. It snaps to the beat, and later bands on that lane ripple forward so the lane stays ordered and non-overlapping.
- **Resize** — drag the handle at either end.

- **Delete** — click the **x** on the band.

All three are undoable (but see the note on Performance's separate undo stack below).

A–B loop brace

The **Loop A–B** button in the Performance toolbar toggles the arrangement-wide loop brace. When active:

- Two drag handles on the ruler mark the loop start (A) and loop end (B). Drag either handle to set the window; handles snap to whole bars.
- The playhead wraps within [A, B). When it reaches B, all lanes are stopped and playback re-anchors to A instantly. Any clip that was already active spanning A is relaunched at the wrap point so there is no gap.
- When Loop A–B is off, Loom plays in **song mode**: the arrangement runs from the beginning to `durationSec` and stops — every lane is halted and the transport stops.

This loop brace operates on the arrangement timeline as a whole. It is distinct from the **per-clip loop brace** available inside the piano-roll and drum-grid editors, which loops a sub-region within a single clip. See [Editing Clips](#) for the per-clip loop.

Song playback

Press Play while in Performance mode to start playback from the beginning of the arrangement. The arrangement's own play state is used — it is independent of the live Session runtime. The playhead advances, clips launch at their scheduled times, and automation curves are applied continuously.

At the end of the arrangement (when Loop A–B is off), all lanes stop and `onArrangementEnd` fires, which stops the transport. You can press Play again to restart from the top.

Knob automation written into the arrangement's curves is applied every lookahead tick alongside clip launches. Automation values are normalised (0–1) internally and mapped to each parameter's min–max range at playback time.

Automation lanes

Automation curves can be added by hand — you do not have to record a take first. Set a non-zero **Length** in the toolbar (or record/copy content so the arrangement has a duration), then use the automation header that appears just below the ruler.

Adding a lane

The header row contains a grouped parameter dropdown and a **+ Automation** button. The dropdown lists every automatable parameter in the project, organised by prefix (lane ID or `master`). Each entry shows the parameter ID and its label — for example `lane-1.fx.reverb.wet — WET`. Select the parameter you want, then click **+ Automation**. A new lane appears below the clip band for that lane (or in the Master section for global parameters). The curve starts flat at the parameter's current value.

Drawing the shape

Two brush buttons in the toolbar control how you paint:

- **Line** — click and drag to draw a ramp between the start and end points of the gesture. Use this for smooth fades, filter sweeps, or any gradual change.
- **Flat** — paints a constant value across the drag range. Use this for step-style automation or to hold a value steady across a section.

The active brush is highlighted. Each lane also exposes **On / Off** and **Smooth / Stepped** toggles in its header. **On / Off** mutes the curve without deleting it. **Stepped** switches interpolation from smooth linear to staircase, snapping the value at each sub-step boundary — useful for parameter jumps that should be instantaneous.

Removing a lane

Click the **×** button on the right side of the lane header to remove the curve entirely. The action is undoable with **Ctrl+Z / Cmd+Z**.

How automation curves play back

Curves added manually behave identically to curves captured by recording. During arrangement playback they are applied continuously alongside clip launches — knob values are updated every lookahead tick, mapped from the normalised 0–1 curve to the parameter's min–max range. Curves generated by a recorded take (see [Transport — REC](#)) and manually drawn curves live in the same list and can coexist on the same lane.

For modulator-driven per-lane automation (LFO / ADSR) that runs in Session view rather than the arrangement timeline, see [Modulation & Note FX](#).

Lane mute, solo & VU meters

Each lane header — visible in both the session lane strip and the mixer — carries three live controls:

- **M (Mute)** — silences the lane's audio output. The lane continues to schedule notes; unmuting restores it at full level with no re-trigger.
- **S (Solo)** — solos this lane, temporarily muting all other lanes. Multiple lanes can be soloed simultaneously; un-soloing the last one restores normal playback.

- **VU meter** — a vertical level meter in the lane header shows the lane's live output RMS. It gives an at-a-glance read of which lanes are active and how loud each one is during a take or a performance.

Mute and solo states are saved with the session.

Persistence

The arrangement is saved as part of the session file (schema version 3). Recorded takes — clip bands, automation curves, loop brace position — survive Save / Load and browser restarts.

Performance has its own, separate undo stack. This surprises people, so it is worth stating plainly: while you are in Performance mode, `Ctrl+Z` / `Ctrl+Y` undo **timeline edits only** — moving, resizing or deleting bands, length and zoom changes, adding or removing curves, drawing automation. Those keystrokes never reach session undo, even when the timeline stack is empty and nothing happens. Conversely, session edits are invisible to Performance's undo. Switch back to Session view to undo session changes.

Two things are outside undo entirely: the **raw recording** (a take is finalised on Stop and does not enter the stack) and **Copy to Performance** (it overwrites the arrangement outright).

See [Saving & Export](#) for how sessions are saved and loaded.

Developer Guide

This chapter is for contributors who want to extend Loom or understand how its internals fit together. Read it alongside `CLAUDE.md` at the repo root, which is the canonical, always-current architecture reference — this chapter expands on it in prose.

The spine

Three structures hold everything together:

1. **A plugin registry** — engines, FX, and modulators are all plugins. They are discovered at build time via a Vite `import.meta.glob` scan, so adding a new one means dropping a file in the right directory, not editing core wiring.
2. **SessionState** — the pure data model: lanes contain clips, scenes reference which clip each lane plays. No audio side-effects live here.
3. **LaneResourceMap** — owns the live Web Audio nodes for each lane. One entry per lane, holding a `ChannelStrip`, a `SynthEngine` instance, and an `InsertChain`. The lane allocator in `src/app/lane-allocator.ts` is the sole path for creating and swapping these resources; nothing else should construct them directly.

The plugin registry

`src/app/plugin-bootstrap.ts` calls `import.meta.glob` at build time over two trees:

```
src/engines/*.ts      - synth engines
src/plugins/fx/*.ts   - FX inserts
src/plugins/modulators/ - LFO / ADSR modulators
```

Every module in those trees is eagerly imported. Any exported value that satisfies the `PluginFactory` shape (`{ kind, manifest, create }`) is collected and registered via `registerPlugin`. The engine registry (`src/engines/registry.ts`) supports both a **singleton** pattern (`registerEngine`) for shared instances and a **factory** pattern (`registerEngineFactory` / `createEngineInstance`) for per-lane instances that need independent state.

`listEngines()` reads from the singleton map and is the source of metadata (name, type, polyphony, parameter specs) used to populate the lane engine selector. Engines declare their parameters as `EngineParamSpec[]`; voices expose continuous `AudioParam`s via `getAudioParams()`, and engines expose shared ones via `getSharedAudioParams()`. That is the entire surface the modulation and automation systems need — nothing else.

SessionState data model

`src/session/session.ts` defines three levels:

- **SessionLane** — has an `engineId`, a list of `SessionClip | null` slots, and an `engineState` bag that persists knob values, modulator configs, note-FX, sampler keymap, pad params, and kit mode.
- **SessionClip** — holds `notes: NoteEvent[]` (the unified note list for both melodic and drum clips), optional `clipEnvelope[]` for per-clip automation, and an optional `sample` field for loop/song audio clips. Clips also carry `loopEnabled` / `loopStartTick` / `loopEndTick` for sub-region looping, and a `gridResolution` hint for the drum editor.
- **SessionScene** — a `clipPerLane` map from lane id to clip slot index (or null for a stopped lane).

Notes carry a `velocity` field (0–127). The `velToColor` function in `src/core/velocity-color.ts` maps velocity to a blue-to-yellow ramp used by both the piano roll and the drum grid.

Saves are written as `schemaVersion: 3` (`SavedStateV3` in `src/save/`). Older saves are normalised by `session-migration.ts` at load time before anything else touches the data.

LaneResourceMap and the audio graph

The master audio path assembled in `src/app/audio-graph.ts` runs:

```
master GainNode → insert chain → MasterCompressor → AnalyserNode
                                     → SidechainBus
```

Each lane's `LaneResources` consists of a `ChannelStrip` (level, EQ, send levels), a `SynthEngine`, and an `InsertChain` of per-lane FX. `LaneResourceMap.replaceEngine` hot-swaps only the engine while keeping the strip and inserts in place — the channel-level resources survive an engine swap.

`ensureLaneResource` in the lane allocator is the only place that constructs a `LaneResources`. Call it once per lane before accessing anything in the map. Test code that needs a lane wired up must call it explicitly as setup.

The scheduler

The `Sequencer` class (`src/core/sequencer.ts`) fires every 25 ms (the poll interval) and looks **200 ms** ahead. On each tick it calls `sessionTick(now, lookaheadSec)` with `lookaheadSec = 0.2`, and the session host fans that out to `tickLane` for each playing lane.

`tickLane` (`src/core/lane-scheduler.ts`) implements the Chris Wilson two-clocks pattern: for every `NoteEvent` whose absolute schedule time falls in the window `[now, now + lookaheadSec)`, it calls `ctx.onTrigger`. Schedule times are derived by converting clip-tick positions to seconds using the

current BPM and projecting onto the absolute timeline from the loop-start anchor. Step duration for a 16th note is `60 / bpm / 4` seconds.

Two important consequences for contributors:

- `bpm` and `length` are mutable at runtime; the next scheduled step picks up the new values immediately.
- Engine params are read at trigger time, not when a note is held. Live knob tweaks apply to the **next** trigger.

When `clip.loopEnabled` is set, `effectiveClipLoop` (`src/core/clip-loop.ts`) constrains the iteration window to `[loopStartTick, loopEndTick)`. The brace UI in `src/core/clip-loop-brace.ts` is the editing surface for that region.

How-to recipes

Add a synth engine

Since synthesis moved into the AudioWorklet, an engine is **two halves in two places**: a metadata descriptor on the main thread and a per-sample renderer inside the worklet bundle. There are **four** steps, and skipping any of the last three gives you an engine that appears in the lane selector and then plays **silence** — the failure is quiet, so do all four.

1. **Metadata descriptor** — `src/engines/<id>.ts`. Build it with `createDescriptorEngine(...)` and register it with `registerEngineFactory(id, ...)`
 - `registerEngine(...)` at module scope. Declare params as `EngineParamSpec[]`. This half carries no DSP: it describes the parameters and the UI.
2. **Renderer** — `src/audio-dsp/<id>-renderer.ts`, a pure per-sample voice renderer that calls `registerRenderer(id, ctor)` at module scope. This is the half that makes sound, and it is plain TypeScript — unit-test it directly, no `AudioContext` required.
3. **Side-effect import in the worklet** — add `import '../audio-dsp/<id>-renderer';` to `src/audio-worklet/loom-processor.ts`. The worklet is a separate bundle; without this import the renderer never registers *inside it* and `createRenderer` throws `no renderer registered for engine '<id>'`.
4. **Route the lane** — add the id to `WORKLET_ENGINE_IDS` in `src/app/lane-allocator.ts`. The allocator only sends listed ids down the worklet path; an unlisted id falls through, the lane gets no engine, and nothing sounds.

Note that `create()` and `getAudioParams()` are **never called** on the registered engine any more — the descriptor's `create()` deliberately throws, and modulation no longer runs on Web Audio `AudioParam`s. It runs sample-accurately inside the worklet (`src/audio-dsp/modulation-runtime.ts`).

See [Engines](#) for the full engine catalogue.

Add an FX insert or modulator

1. Create `src/plugins/fx/<name>.ts` or `src/plugins/modulators/<name>.ts`.
2. Export a `PluginFactory` with `kind: 'fx'` or `kind: 'modulator'` respectively, and call `registerPlugin` at module scope.
3. FX inserts mount per-lane and on master automatically; modulators appear in the modulation panel. See [Modulation and Note FX](#) for the user-facing side.

Add a preset

Open `public/presets/<engine>.json` and append an entry. The `gm` field is optional (an integer GM program number for MIDI-import matching). JSON is the source of truth; `preset-loader.ts` validates and `preset-apply.ts` applies it at runtime by calling `engine.applyPreset`. Each engine's JSON keys are its own vocabulary — do not use a generic `setBaseValue` loop.

Add a synth drum kit

Append an object to the `KITS` array in `src/core/drums.ts`. Kits are parameter bags over shared DSP primitives. To add a new drum *voice* (not just a new kit): extend the `DrumVoice` union, add it to `DRUM_LANES`, add an entry to every kit, implement a `play<Voice>()` method, and add a `trigger()` case.

Add a sampled drum kit

1. Create a subdirectory `public/drumkits/<id>/` containing WAV files for each voice (e.g. `kick.wav`, `snare.wav`, `closedHat.wav`).
2. Add a manifest file `public/drumkits/<id>.json` with `id`, `name`, and a `samples` array. Each entry needs `voice`, `note` (GM MIDI note number), and `file` (path relative to `public/drumkits/`).
3. Register the kit in `public/drumkits/index.json` by appending `{ "id": "<id>", "name": "<display name>" }`.

The existing kits (`tr808`, `acoustic`, `dirt`) follow this layout exactly and are the reference.

Source layout tour

```

src/
  core/      DSP primitives + pure logic (synth, drums, sequencer,
             lane-scheduler, lane-resources, fx, meter, notes,
             history, knob, pianoroll, ...)
             velocity-color.ts / velocity-gain.ts / velocity-lane-editing.ts
             - note-velocity colour ramp, gain curve, lane editing helpers
             clip-loop.ts / clip-loop-brace.ts
             - clip sub-region resolver + drag-brace UI primitive
  engines/   SynthEngine abstraction, registry, one file per engine
             (tb303, subtractive, fm, wavetable, karplus, sampler,
             drums-engine) + engine-selector UI
  session/   SessionState model + all session UI
             (session-host, session-ui, session-inspector,
             clip-editors/, session-migration)
  modulation/ LFO/ADSR voices, ModulationHost, ModulatorScope,
             connection binder
  plugins/   Plugin SPI + registry
             fx/      - multifilter, distortion, reverb, delay, InsertChain
             modulators/ - lfo, adsr
  presets/   Preset loader + apply logic
             (JSON assets live in public/presets/)
  midi/      SMF parser, MIDI-to-session transform, GM lookup, import UI
  samples/   Sample types, IndexedDB store, buffer cache, keymap,
             import metadata
  stems/     Stem-separation client + config + lane-plan builder
             (talks to the local Python service in tools/stem-service/)
  performance/ Arrangement / record model:
             arrangement-from-session, arrangement-ops,
             arrangement-runtime (records clip-launches + knob automation;
             surfaced via the REC group's take mode – see performance-feature)
  audio-dsp/ THE SYNTHESIS KERNEL. Pure per-sample DSP that the worklet
             runs: one <id>-renderer.ts per engine (self-registering into
             renderer-registry.ts), voice-manager, scheduler-queue,
             modulation-runtime, plus primitives (osc, filter, ladder,
             sync-osc, unison, adsr). No AudioContext – unit-test directly
  audio-worklet/ The processors + typed node wrappers: loom-processor/loom-node
             (melodic), drums-*, sampler-*. A processor is referenced ONLY
             via ?worker&url and its registered name – never imported on
             the main thread (see processor-name.ts)
  export/     Offline scene/WAV render + the live take recorder
  patterns/   The pattern library (styles x patterns) + its picker UI
  perf/       Performance diagnostics (the PERF HUD)
  polysynth/  LEGACY/vestigial. Not the Subtractive voice host any more –
             Subtractive is a descriptor engine rendered in the worklet.
             Only the preset store + the extra-poly path still touch it
  app/        Boot factories: audio-graph, lane-allocator, trigger-dispatch,
             knob-mounting, mute-solo, bpm-broadcast, engine-swap,
             plugin-bootstrap, lane-host-wiring
  save/       SaveManager (schemaVersion: 3), auto-history (AutoHistory:

```

	snapshot-diff undo/redo + gesture coalescing, wired to the transport-bar ⏮/⏭ buttons), history-wiring (withUndo / attachKnobUndo + the undo keyboard – LIVE and load-bearing: withUndo wraps mutation sites across the app)
notefx/	Note-FX plugin category (arpeggiator, chord spread) – per-lane
automation/	Clip envelope recording + read-back helpers
control/	Live MIDI controller subsystem: APC Key 25 profile, live keyboard, LED mediator, profile registry
demo/	Baked MIDI demos + demo picker
styles/	SCSS
public/	
presets/	Engine preset JSONs (20+ per engine, GM-tagged)
drumkits/	Sampled drum kits: index.json + <id>.json manifests + WAVs
tools/	
stem-service/	Local Python service (FastAPI + audio-separator / Demucs) exposing an HTTP job queue for stem separation. Run: <code>uvicorn app:app --port 8765</code> Tests: <code>python -m pytest test_app.py</code> (not part of npm test)

Testing

Loom has four test layers, one per risk class.

Pure logic (`src/**/*.test.ts` , excluding `.dsp` and `.wiring` suffixes) — schemas, scales, migrations, session/arrangement logic, modulation math. These run fast and have no audio dependencies.

Scheduling with a fake clock — `src/core/lane-scheduler.test.ts` and `src/session/session-runtime.test.ts` drive the look-ahead scheduler through a mock `AudioContext` clock. The fake clock advances in controlled steps so timing edge-cases are deterministic.

Real DSP (`*.dsp.test.ts`) — every engine and drum kit is rendered through `OfflineAudioContext` via the `node-web-audio-api` package, globalised in `test/setup.ts` . Use the shared battery `runStandardEngineBattery` from `test/dsp-battery.ts` . Each render writes a WAV to `test/output/` (gitignored). Compare against the committed reference in `test/golden/` with `npm run test:wav-diff` ; promote with `npm run test:wav-bless` .

Modulation wiring (`*.wiring.test.ts`) — LFO/ADSR voices connected through a depth-gain bridge into a target `AudioParam` , verified end-to-end.

Assertion rule: always write relative assertions (`a > b` , `a > b * 2`). Never hard-code absolute magnitudes — they are a brittleness smell. If you must write one, justify it in a comment.

Colour-free output: every `npm test` script runs under `cross-env NO_COLOR=1` . When invoking Vitest directly, prefix with `NO_COLOR=1` . Do not add `--reporter=...` — the scripts already configure the right reporter.

Key commands:

Command	What it runs
<code>npm run dev</code>	Vite dev server with hot reload at http://localhost:5173
<code>npm run build</code>	<code>tsc</code> typecheck + Vite bundle to <code>dist/</code>
<code>npm test</code>	Full suite: unit + e2e (always build first)
<code>npm run test:unit</code>	Vitest only, no browser
<code>npm run test:fast</code>	Unit tests excluding DSP renders (inner-loop TDD)
<code>npm run test:dsp</code>	DSP renders only (slow; needs <code>node-web-audio-api</code>)
<code>npm run test:e2e</code>	Playwright against <code>vite preview</code> on port 4173

e2e gotcha: `test:e2e` and `npm test` serve `dist/` with no build step. Playwright boots `vite preview` over the last production bundle. If you changed `src/` without rebuilding, the newest features are absent from the bundle and tests fail with "element not found" — which looks like a regression. Always run `npm run build` before `npm run test:e2e`.

Vitest runs test files serially (`fileParallelism: false`) because `node-web-audio-api` 's `OfflineAudioContext` is not safe under parallel forks. The teardown occasionally exits non-zero with `ERR_IPC_CHANNEL_CLOSED` after all tests pass — that is a tinypool shutdown race, not a test failure; re-run to confirm green.

For the definitive, always-up-to-date architecture reference, read `CLAUDE.md` at the repo root. Implemented design docs are intentionally removed from the tree once shipped (they drift); recover rationale from git history when you need it. Unfinished work lives in `docs/superpowers/REMAINING-WORK.md`.