


Loom — Manual

Loom is a browser-based, session-based music workstation. It organises sound into **lanes** (instrument tracks), **clips** (patterns of notes), and **scenes** (rows of clips launched together). Nine synthesis engines are available — TB-303, Subtractive, FM, Wavetable, Karplus-Strong, West Coast, Sampler, Drum Machine, and Layers, which holds up to four of the melodic ones in a single lane — and all synthesis runs live in the browser with no install or sign-in required.

Beyond the engines, Loom has **note velocity & dynamics** (an Ableton-style velocity lane), **per-clip and arrangement-wide loop regions**, **per-lane modulation and FX** with a sidechain mixer, **MIDI import**, ready-made **sampled drum kits**, a **playable arrangement** timeline, improved **session editing workflow** (inline track/scene rename, duplicate track/scene, and capture-playing-to-scene), optional **stem separation + transcription** (split a finished song into 4 lanes via a local helper and optionally derive editable note/drum lanes), **key & scale musical assistance** (a project key/scale with an optional piano-roll scale lock — off by default, genre generators, an examples gallery, and a chord-accompaniment maker), and **editable warp markers** for locking audio loops to tempo.

WEAVE, the third view beside Session and Performance, is where a session stops being a fixed arrangement: it crossfades between loops and keeps crossfading on its own, with a per-track list of the loops it may draw from and a **note generator** that reads a track's own material rather than playing it back.



The Loom interface: transport across the top, the session clip grid in the centre, and per-lane channel strips below.

Table of Contents

#	Chapter
1	Getting Started
2	Transport
3	Sessions, Lanes, Clips & Scenes
4	Engines
5	Editing Clips
6	Modulation & Note FX
7	Mixing & FX
8	MIDI & Samples
9	Saving & Export
10	Performance & Arrangement
11	WEAVE
12	Developer Guide

[Loom-Manual.pdf](#)

Live demo: <https://ijol.github.io/Loom/>

Getting Started

Open the app

Loom runs entirely in the browser — no installation, no account, no plugins.

- **Live demo:** <https://ijol.github.io/Loom/>
- **Run locally:** clone the repository, then run `npm install` followed by `npm run dev`. The app is available at <http://localhost:5173>.

Make your first sound

When the app opens it loads the *Minimal Techno* demo automatically, so there is already a full arrangement ready to play.

1. Click the ► **Play** button in the transport bar. The first click resumes the browser's AudioContext — this is a browser requirement; audio cannot start without a user gesture, so nothing plays until you press Play.
2. Sound starts immediately. Use the **Volume** slider in the transport to adjust the output level.
3. Want to hear a different demo? There are **37** in the — **load a demo** — dropdown (also under **File** ► **Open Demo**). The first four — Minimal Techno, Acid Rain, Cordillera, Neon Drive — were built here; most of the rest are ports of tunes from the [Strudel](#) library, from Giant Steps to Amensister, and one, Supersaw Mask, came across from a live-coding patch. Picking one loads the session and begins at the first scene.



The full Loom interface. Transport across the top; coloured clips in the session grid; per-lane channel strips below.

The menu bar

Across the very top sits a desktop-style menu bar. Most of what it offers also has a button somewhere in the interface — the menu is the tidy, discoverable route to the same things, and the place to look when you can't remember where a control lives.

FILE EDIT VIEW TOOLS HELP

NEW SESSION CTRL+N

OPEN... CTRL+O

SAVE CTRL+S

SAVE AS...

PROJECT OPTIONS...

OPEN DEMO ▶

IMPORT MIDI...

SEPARATE INTO STEMS...

PREFERENCES... CTRL+ ,

1.1 BPM METER
0:00 **130** 4/4 ▼

⇒ ● REC  

	✕ Sub 1	✕ S
▶	✕ 1 ▶	
▶	✕ 2 ▶	
▶	✕ 3 ▶	✕ 3
▶	✕ 4 ▶	✕ 4

Item	Shortcut	What it does
New Session	Ctrl+N	Empties the session and starts fresh
Open...	Ctrl+O	Opens the save manager to load a stored session
Save	Ctrl+S	Opens the save manager to store the current session
Save As...		Same panel, for storing under a new name
Project Options...		Project-wide settings, including the name and the musical key/scale
Open Demo ▶		A submenu of the bundled demos; picking one loads it
Import MIDI...		Opens the MIDI import flow (chapter 8)
Separate into Stems...		Opens stem separation (chapter 8)
Preferences...	Ctrl+,	Not yet available — greyed out, planned for a later release

Edit — **Undo** (Ctrl+Z) and **Redo** (Ctrl+Shift+Z , or Ctrl+Y). Both grey out when there is nothing to undo or redo, so the menu doubles as an indicator of whether your last action was recorded.

View — switches between the **Session** and **Performance** views, with a tick marking the current one, and toggles the **Performance diagnostics (PERF)** overlay (audio load, scheduler lag, FPS, voice counts).

Tools — opens the **MIDI Controller** panel, runs **Capture Scene** (Ctrl+I , snapshots the currently playing clips into a new scene) and **Copy Scenes** → **Performance**.


Help — opens this **Manual** in a new tab, and **About Loom**.

Sessions are stored in your browser, not as files on disk: "Open" and "Save" mean the in-browser save manager, described in [chapter 9](#).

The mental model

Loom organises music in three nested concepts:

- A **lane** is an instrument track. Each lane runs one engine — six melodic ones (TB-303, Subtractive, FM, Wavetable, Karplus-Strong, West Coast), the Sampler, the Drum Machine, Layers (up to four melodic engines in one lane), or an Audio channel for playing a recording — and has its own mixer strip with EQ, sends, pan, and fader.
- A **clip** is a pattern of notes that lives in a lane. A lane can hold multiple clips — one per scene row.
- A **scene** is a horizontal row across all lanes. Launching a scene fires the clip in each lane at that row simultaneously, so scenes work like song sections or loop variations.



The session grid: lanes run left to right as columns, scenes top to bottom as rows. Click a clip to select it; click the scene launch button on the right to fire the whole row.


Click any clip cell to open its editor in the inspector below the grid. The editor shows a piano roll for melodic lanes and a drum grid for drum lanes.

Next steps

- [Transport](#) — BPM, meter, swing, and the unified REC controls (take / live-WAV / offline-WAV).
- [Sessions, Lanes, Clips & Scenes](#) — how to build and arrange your own session from scratch.

Transport

The header runs across the top of the interface in two rows. The first row (transport & tempo) holds the playback and timing controls you touch while the music plays; the second row (session & I/O) holds the mode toggle, recording/export, and session file management.



Play / Stop

Press ► (`#play` , title "Play — start the transport") to start playback, and ■ (`#stop` , title "Stop — stop the transport") to stop it. Play and Stop are now two separate buttons — Play is no longer a play/stop toggle. When you press Play the browser's `AudioContext` resumes automatically if it was suspended (a browser policy requirement on first interaction).

Position and time readout

The readout shows two values side by side:

- **Bar.Beat.Step** (`#transport-position` , e.g. `1.1.1`) — the current song position derived from elapsed time, BPM, and the active meter. The counter resets to `1.1.1` on each new Play.
- **Elapsed** (`#transport-time` , e.g. `00:00:00`) — wall-clock time since the most recent Play press.

Both readouts update via `requestAnimationFrame` while playing and freeze on stop so you can read the final position. Because each lane has its own independent loop clock, the global position is computed from elapsed seconds rather than a single sequencer cursor. The visual playhead is similarly a display-only timer matched to scheduled audio time; it may drift slightly under browser tab throttling, but audio scheduling is unaffected.

Tempo controls

BPM (`#bpm`) — sets the tempo. Range: 40–400 BPM, default 130. (The ceiling is deliberately high: drum & bass written at its true tempo, and the faster end of hardcore/gabber, both live above 240.) The field accepts **floating-point values** (e.g. 128.57 BPM, step 0.01); tempos detected from MIDI imports or stem separation are no longer rounded to an integer. You can type a value directly or use the number-input spinners. BPM changes propagate immediately to the sequencer, all lane engines, delay/LFO sync, and stretch-mode loop buffers. The change takes effect on the next scheduled step; it does not alter a note that is already held.

Each 16th-note step lasts `60 / bpm / 4` seconds.

Meter (`#meter`) — sets the global time signature. The dropdown offers common meters: 4/4, 3/4, 2/4, 5/4, 6/8, 7/8, 9/8, and 12/8. The meter controls how bars map onto the 16th-step grid and how the position readout counts beats. Changing the meter takes effect on the next loop cycle.

Swing (`#swing`) — adds a shuffle feel by delaying the off-beat 16th notes. Range: 0 (straight) to 0.6. A swing of `s` delays each off-beat 16th by `s × one 16th step`, so the off-beat lands `50 × (1 + s)` % of the way to the next on-beat — the familiar swing percentage: **0** = straight (50%), **0.33** = the classic triplet shuffle (66.7%), **0.5** = a heavy shuffle (75%), **0.6** = the hardest limp Loom offers (80%). On-beats never move.

Swing applies to every lane, and to notes that do not sit exactly on a 16th too (32nds, imported MIDI, off-grid drum hits): each note is carried along with the beat it belongs to rather than left behind by it. Note lengths follow their notes, so a TB-303 slide still slides. The value is saved and restored with the session, and WAV exports render with the same swing you hear.

Master volume

Volume (`#volume`) — the master output level, range 0–1 (default 0.5). This is a post-mix gain applied before the output visualiser.

Mode toggle and REC

Session / Performance (`#mode-toggle`) — switches the main view between the Session clip grid and the Performance arrangement view.

† **Copy to Performance** (`#copy-to-performance`) — an icon-only button (tooltip "Copy the scenes to the Performance timeline") that copies the current scenes onto the Performance timeline. See [Performance & Arrangement](#).

- **Session** is the default mode: you see the clip grid and can trigger scenes, edit clips, and record automation in real time.
- **Performance** shows the arrangement timeline where recorded takes are displayed as timeline bands and automation curves can be drawn directly. See [Performance & Arrangement](#) for the full workflow.
- **REC** (`#rec` , tooltip "Record — pick the mode beside it") arms recording; the **mode selector** (`#rec-mode`) beside it chooses what gets recorded:
 -  **take** (`data-recmode="take"` , the default) — records knob moves and clip launches into a performance take. See [Modulation & Note FX](#) for how automation lanes work.
 -  **live** (`data-recmode="live"`) — records the real-time audio output to a WAV.

- ⚡ **offline** (`data-recmode="offline"`) — renders the current scene to a WAV offline (fast).

Click **REC** again to disarm. This unified REC group replaced the old standalone WAV-export button.

Output visualiser

The canvas at the far right (`#viz`) shows a real-time waveform of the master output. It updates continuously while the page is open and gives a quick visual check that audio is flowing.

Session management, export, and demos

The remaining controls in the transport row are covered in dedicated chapters:

- **WAV export** — there is no longer a standalone "↓ WAV" button. WAV export now happens through the **REC** group's mode selector (see the Mode toggle and REC section above): 🎙️ **live** records real-time audio to WAV, and ⚡ **offline** renders the scene to WAV offline (fast). See [Saving & Export](#).
 - **New / Save / Load** (`#new-session` , `#save` , `#load`) — session file management. See [Saving & Export](#).
 - — **load a demo** — (`#demo-picker`) — loads a bundled demo arrangement into the session.
 - **MIDI import** — there is no button for this in the transport row. It lives under **File ▶ Import MIDI...**, which opens a modal dialog. See the MIDI Import chapter for details.
 - ☰ (`#stems-open` , an icon-only button tooltipled "Separate a song into stems") — opens the stem-separation dialog, which splits a finished song into four lanes (vocals / drums / bass / other) via a local helper service. Also under **File ▶ Separate into Stems....** Requires the service to be running; see [MIDI & Samples → Stem separation](#).
 - 📸 **Capture** (`#capture-scene`) — snapshots the clips currently playing into a new scene (also `Ctrl+I` , and under **Tools ▶ Capture Scene**).
 - 📐 **XY** (`#xy-open`) — opens the floating XY pad for sweeping two parameters at once. See [Mixing & FX → XY pad](#).
 - **PERF** (`#perf-toggle`) — toggles the performance-diagnostics overlay (audio load, scheduler lag, FPS, voice counts). It only runs while open.
 - ↶ / ↷ (`#undo-btn` , `#redo-btn`) — global undo and redo, greyed out until there is history.
-

Keyboard shortcuts

Key	Action
Space	Pause / resume
R	Arm or disarm clip MIDI recording
Ctrl/Cmd+N	New session
Ctrl/Cmd+O	Open
Ctrl/Cmd+S	Save
Ctrl/Cmd+Z	Undo
Ctrl/Cmd+Shift+Z or Ctrl/Cmd+Y	Redo
Ctrl/Cmd+I	Capture playing clips into a new scene

Two things about **Space** that catch people out:


- It is **pause/resume, not stop**. It remembers the exact position — including the fraction of a bar — and picks up from there, where Stop resets to the beginning.
- It does **nothing from a standing start**. If the transport is idle, Space will not begin playback; use ► or launch a scene. It only pauses something already playing.

Note also that after a pause/resume the position readout restarts its count from **1.1.1** even though the music continues where it left off — the readout times the current run, not the song.

Shortcuts never fire while you are typing in a text field, so they can't steal input from the BPM box or a name field.

Sessions, Lanes, Clips & Scenes

The session view is the main workspace in Loom. It organises everything you hear into a compact grid of lanes and clips that you can launch, edit, and rearrange in any order.




The model

Lane — one instrument track. Each lane runs one engine — six melodic ones (TB-303, Subtractive, FM, Wavetable, Karplus-Strong, West Coast), the Sampler, the Drum Machine, Layers, or an Audio channel — with its own mixer strip and insert chain. Lanes appear as columns in the grid.

Clip — a pattern of notes that lives at one lane × scene cell. Every clip stores a list of note events, an optional name, a colour, and a length in bars. A cell is either filled (it holds a clip) or empty.

Scene — a row in the grid. Launching a scene fires one clip per lane simultaneously, aligning all starts to the same quantize boundary so the whole arrangement snaps together cleanly.

The grid



Lane names run along the top header row; row numbers (1, 2, 3 ...) label the scenes down the left edge. Scenes launch buttons sit in the rightmost column. Each filled cell shows the clip's name (or its row number as a fallback) against the clip's pastel colour. **Click a lane header** to open that lane's instrument editor; once it is open, the ► / ▼ **chevron** that appears on that header collapses and reopens the editor.

Folding it away is a standing choice about screen space, not a statement about which lane you are on: once collapsed it **stays** collapsed as you move between lanes, and unfolding shows whichever lane you are now on. (It used to spring open the moment you touched another lane.)

One lane is selected, everywhere

The clip editor and the instrument controls below it always show the **same** lane. There are several ways to change which one that is — clicking a lane header, opening a clip, clicking a mixer column, or selecting the track from a MIDI controller — and all of them move both halves together.

When you switch lanes, the clip editor follows to the **same row** on the new lane: if lane A's clip in row 2 is open and you select lane B, you get lane B's row-2 clip. If that slot is empty the editor closes, and **nothing is created** — an empty cell stays empty until you fill it deliberately. The instrument controls always follow, clip or no clip.

There is one exception, and it is the one you want: opening a clip on another lane does not close the editor it just opened.

A newly created instrument lane starts out empty — no placeholder clips are added to its column; you fill cells yourself. Audio and Sampler lanes created from a dropped WAV or loop place their clip in row 1 only. The grid always keeps at least one launchable scene available even when every lane is empty.

Right-clicking grid elements opens a context menu. On a **lane header**: *Rename track*, *Edit instrument*, *Duplicate track*, *Stop track*, and *Delete track* (red). On a **scene cell**: *Rename scene*, *Launch scene*, *Duplicate scene*, *Capture playing -> scene*, *Add scene*, and *Delete scene* (red). On a **filled clip**: *Open editor*, *Play / Stop*, *Color* (a row of swatches), and *Delete clip* (red). On an **empty cell**: *Create clip* — on audio lanes this entry reads *Import audio (WAV)*....

Below the scene rows there is a stop row: each lane has its own ■ stop button, and a global ■ **all** button at the far right stops every lane at once.

Deleting from the grid

A small ✕ cross appears on every lane header (*Delete track*), every filled clip cell (*Delete clip*), and every scene cell (*Delete scene*). Deleting a clip is immediate. Deleting a lane or scene asks you to confirm in an in-app dialog (**OK** / **Cancel**, with the destructive choice in red) only when the target still has content — an empty lane or scene is removed straight away. Every deletion is undoable with **Ctrl+Z**, which restores the lane, scene, or clip along with its audio resources.

Launching clips and scenes

Click the ► **icon** on any filled cell to launch that clip. The icon switches to || once the clip is playing; clicking || stops that lane. Clicking anywhere else on the cell body — outside the play icon — opens the inspector without affecting playback.

From a cold start (transport stopped), launching a clip starts it immediately, and the **Quantize** setting decides the alignment of that first start. **While the session is already playing**, Loom keeps everything phase-aligned instead:

- **Launching a single clip** swaps it in at the end of *its own* current loop, so the new pattern enters cleanly on the downbeat rather than mid-bar.
- **Launching a scene** is an **atomic switch at the end of the governing loop** — normally the longest loop currently playing, so short clips do not cut a long one off mid-phrase. One stray *very* long clip is ignored rather than allowed to hold everything up: if the longest loop is more than twice the next longest, it is dropped and the comparison repeats, so a lone 16-bar pad does not make a session of 4-bar clips wait sixteen bars to change scene. Every clip in the scene changes together on that boundary, and any **lane that has no clip in the new scene stops** at the same moment (it pulses a "stopping" state until then). The result is a clean, simultaneous scene change with no lanes left hanging.

Click the ► **scene name** button at the right of a scene row (it reads ► **Scene 1**, or whatever you renamed it to — not a bare number) to launch all clips in that row. The toolbar's ■ **All** button stops every lane at once.

The **Quantize** setting (per clip, falling back to the lane's, then the session's global) governs **cold starts** — the alignment used when you launch from a stopped transport. Once the transport is running, the loop-end sync above takes over so launches stay musically aligned.

Renaming, duplicating, and capturing scenes

Track and scene names support **inline rename**: double-click a lane name in the header or a scene name in the scene cell, type the new label, and press Enter.

You can also use the context menus to duplicate structure quickly:

- **Duplicate track** clones the whole lane with fresh clip ids and keeps its clip content.
- **Duplicate scene** appends a full clone of that scene's lane/clip assignments.


To snapshot the current live state into a new scene, use **Capture scene**:

- Click the ☉ button in the **Scenes** header (tooltip: *New scene from currently playing clips*).
- Or press **Ctrl+I**.
- Or use **Capture playing -> scene** from a scene cell's context menu.

Capture scene appends a new scene row and writes one clip-per-lane mapping from whatever is playing at that moment.

The inspector

Clicking a clip cell body opens the inspector below the grid with controls for that clip.



Control	What it does
Name	Free-text label shown on the cell.
Length (bars)	How many bars the clip plays before looping.
Quantize	Per-clip launch alignment: <i>Default</i> (inherits lane/global setting), <i>Immediate</i> (no wait), or a fixed interval — 1/4, 1/2, 1 bar, 2 bars, or 4 bars.
Copy notes	Copies the clip's notes to an internal clipboard.
Paste ▸ Replace	Replaces the selected clip's notes with the clipboard contents.
Paste ▸ Layer	Merges the clipboard notes into the selected clip rather than replacing them.
Duplicate	Creates a copy of the clip in the next empty cell of the same lane.
View as grid / View as piano roll	Toggles between the piano-roll editor and the drum-grid editor — the button always names the view you are <i>not</i> in. The appropriate editor is chosen automatically based on the lane's engine; use this button to switch if needed. It changes only how the clip is edited, never how it sounds. See Editing Clips for the full editing reference.
 Notes	Randomizes the clip's note content using the current scale and root settings.
Delete	Removes the clip from the grid (also triggered by Delete/Backspace while the inspector is open).

The editor embedded inside the inspector is rebuilt each time a new clip is selected; deep editing of notes is covered in [Editing Clips](#).

Arranging clips

You can move a clip to any other cell by dragging it. Hold **Ctrl** while dragging to copy instead of move. Clips remember their colour across moves and copies. Drag a clip onto a cell that already has a clip to swap or replace it.

Each clip is assigned a colour automatically from a pastel palette, so you can distinguish patterns at a glance even before you give them names. To choose a different one, right-click the clip and pick from the **Color** swatches.

For the mixer strip, insert effects, and per-lane send levels, see [Mixing & FX](#). For the available synthesis engines on each lane, see [Engines](#).

Engines

Every lane in Loom runs exactly one synthesis engine. You pick it when you create the lane, from the **+** menu at the top of the clip grid, which lists every engine. Afterwards you can swap a lane's engine from the **ENGINE** selector at the top of its editor panel — that dropdown offers the piano-roll engines only, so a lane can be moved between them freely but not turned into a Drum Machine, which edits on the drum grid. Changing the engine replaces the sound source while preserving the lane's clips and modulation routing.

This chapter covers nine engines: six melodic synthesisers (TB-303, Subtractive, FM, Wavetable, Karplus-Strong, West Coast), a Sampler, a Drum Machine, and **Layers** — a lane that holds up to four of the others at once. A tenth choice, the **Audio channel**, plays a recording rather than synthesising one and is covered in [MIDI & Samples](#).

Each engine exposes a PRESET dropdown with **Load**, **Save As...** and **Delete** beside it, and a  **Sound** button that randomises the patch and sets the preset name to "Custom". (Layers has neither a dice nor factory presets: what a rack sounds like is which instruments you put in it.) See [Sessions, Lanes, Clips & Scenes](#) for how to add and configure lanes.

Factory presets and your own

The dropdown has two groups.

Factory are the presets the engine ships with, carried inside the engine itself and tagged with GM programme numbers so MIDI import can pick a sensible engine and preset for each imported track. There are 23 to 28 of them on each of the melodic engines — and 102 on Subtractive, which is the one with the widest range to cover. Selecting one applies it immediately; there is no need to press Load, which only re-applies whatever is already selected.

User is whatever you have saved. **Save As...** asks for a name and stores the sound exactly as it stands on that lane. Two things about it are worth knowing:

- **A user preset belongs to its engine.** Saving "Bells" on an FM lane files it under FM, and it is offered on FM lanes only — a Wavetable lane can hold its own, different "Bells" without the two ever meeting.
- **It saves the instrument, not the desk.** The lane's volume, pan, sends and EQ are left out on purpose, so recalling a preset changes how the sound is made without moving where it sits in the mix.

Delete removes a user preset; factory presets cannot be deleted, and the button says so if you try.

Changing a lane's engine forgets the preset it was on. Swap FM for Wavetable, or convert a lane to a rack, and the dropdown reads **(custom — no preset)** rather than going on naming a sound that is no longer loaded. Nothing is lost: the parameters are still whatever the swap left them.

The same dropdown is what WEAVE offers on a track row, and it is now the same list — a Sampler track in WEAVE finally sees its bundled instruments and its loops, which it did not before.

Your saved presets live in the browser, not in the session file. That cuts both ways: they follow you from one song to the next, and they are **not** included when you share a save — someone opening your session gets the sound, because the parameter values travel with the lane, but not the entry in their own preset list. It also means clearing your browser's site data takes them with it.

The dice rolls the *instrument*, not the desk: it leaves the lane's mixer column (volume, pan, sends, EQ), the polyphony setting and the master tune exactly where you put them, so a randomised sound stays at the level and in the tuning the rest of the mix expects. Rolls are biased towards each parameter's current value rather than spread flat across its range, so most of them land near a sound that works and the extremes are the tail, not the norm. Hit it repeatedly: it is the fastest way to find out what an engine can do.

Every engine responds to **note velocity**. Velocity (0–127) scales each note's loudness continuously along a curve with a floor, so even the softest note still sounds. Accent (velocity ≥ 100) layers additional character on top, and what it adds is engine-specific: it brightens the filter envelope on the bass-style engines (and on the TB-303 alone also raises the resonance), drives the wavfolder harder on West Coast, and simply hits harder on drums. For how to view and edit velocities in the piano-roll or drum-grid, see [Velocity & dynamics](#).

Turning a knob while a note is sounding

Loom's continuous parameters are **live**. Hold a note — or let a pattern run — and move a cutoff, a resonance, an FM ratio, a wavfolder amount: the sound already playing bends with your hand, the way it does on hardware. You do not have to wait for the next note.

A few parameters deliberately do not work that way, and they are the ones where changing mid-note would produce a click rather than a sweep:

- **Waveform** and **filter model** — the shape itself, not a value on it.
- **Unison size** — how many copies make up a voice.
- **Every envelope time** — attack, decay, release. The envelope is computed from how long the note has been held, so re-reading an attack halfway through would jump the level instead of gliding it.

Those apply from the next note you play. Everything else is immediate. (The Drum Machine is the exception to the whole section: its parameters are read when the hit fires.)

Voices, and what happens when you run out

Every melodic engine has a **Voices** control: the most notes that lane will sound at once. The range is **1 to 64**, the knob lands on whole numbers only, and the default is **10** on every polyphonic engine.


It is worth knowing that this is a real limit, because it did not use to be — the number was read as a mono switch, so anything above 1 meant "no limit at all". A pad with a three-second release, played in 16ths, settles at around 28 notes sounding simultaneously; a Layers rack multiplies that by however many slots it holds. That is the cost this control exists to bound, and the symptom of exceeding it is not a warning but the audio thread missing its deadline while everything on screen still looks right.

When the lane is full, the oldest note goes. It is not cut off: it is faded out over 8 ms and only then let go, so there is no step in the waveform to hear. Notes that are already releasing do not count against the cap, so a fast run cannot steal the same voice twice and thin the part out.

Many voices at once no longer grind the master. A lane's voices used to sum at full level, so eight sustained notes could push the lane far past full scale and park the master limiter in continuous distortion — the classic "any polyphony above minimum turns to noise". Each melodic lane now sums its voices with built-in headroom and saturates softly *inside the lane* when a pileup still exceeds it: one note reaches the master untouched, a stack merely thickens. If a lane sounds saturated, it is that lane's Voices, level or material — the rest of the mix stays clean.

Set it by ear. Lower it on a lane you want to sound monophonic (1 is a genuine mono lane), and lower it on a dense pad if the mix is straining; leave it alone otherwise. The TB-303 has no such control — it is monophonic by declaration, not by a number — and **the Sampler's is inert**: it runs its own voice pool and takes no cap.

TB-303



Above: TB-303 editor — Wave, Cutoff, Resonance, Env, Decay, Accent, and a per-lane LFO.

The TB-303 is a monophonic, resonant bass synthesiser modelled on the Roland TB-303. It is the natural choice for acid bass lines but works equally well for aggressive leads and any sound that calls for a steep, self-oscillating filter sweep.

TB-303 parameters

Parameter	Description
Wave	Sawtooth or Square oscillator waveform
Cutoff	Filter cutoff frequency (0–100%)
Resonance	Filter Q — see the note on the diode ladder below
Env	How far the filter envelope opens the filter per step (the classic "env mod")
Decay	Filter envelope decay time
Accent	Per-step level: brightens the filter, bumps Q, raises output gain

The filter is a diode ladder. This is the circuit that gives the 303 its voice: its asymmetric clipping adds even harmonics, which is where the squelch comes from. One consequence surprises people used to other synths — a ladder *loses* level as resonance rises, instead of growing a peak on top of the signal. Turning Resonance up thins and quietens the sound rather than making it louder. That is correct behaviour, and the engine compensates internally so that accented steps still punch through.


Slide and accent behaviour

A note's `slide` flag means "slide into the next step". When the scheduler emits step N it checks whether step N-1 carried a slide flag; if so, it ramps the pitch from the previous note and skips the amp re-attack so the gate stays open across the boundary. The outgoing step gets an extended gate (1.5× step length) so the overlap is audible.

Accent is a per-step flag that simultaneously brightens the filter envelope, raises the resonance Q, and boosts the output gain — the classic 303 bassline punctuation technique.

The engine ships with 20+ presets from "BASS Acid Classic" to "LEAD Squelch". See [Editing Clips](#) for how to set slide and accent on individual steps.

Subtractive



Above: Subtractive editor — OSC 1/2, Sub oscillator, Noise, Filter (with built-in envelope), Amp, and POLY controls.

The Subtractive engine is a classic analogue-style polyphonic synthesiser with two oscillators, a sub oscillator, a noise source, a multimode filter, and a full amplitude envelope. It is the most general-purpose engine in Loom and suits pads, leads, basses, and plucks — and with **102 presets** it has by far the largest library of any engine here.

Parameter sections

- **OSC 1 / OSC 2** — waveform, level, detune in cents, **PW**, and **Sync**. Detuning the two oscillators creates the classic "supersaw" chorus effect. See [Oscillator extras](#) below for PW and Sync.
- **RING** — level of the ring modulator, OSC 1 × OSC 2. See [Ring modulation](#) below.
- **SUB / NOISE** — sub oscillator level (one octave below OSC 1) and a noise generator for breath and texture.
- **FILTER A** — **Mode**, **Type**, Cutoff, Resonance, Env Amount, Drive, Key Track, and a full ADSR filter envelope (toggle with Built-in Env). See [Mode and Type](#) and [The comb, and its two borrowed knobs](#).
- **FILTER B** — a second filter with its own Mode, Type, Cutoff and Res, plus **Routing**, **Blend** and **Track**, which decide how it combines with Filter A. See [Two filters, and how they are wired](#).
- **AMP** — Attack/Decay/Sustain/Release amplitude envelope (toggle with Built-in Env).
- **MASTER** — global Tune in semitones, plus **Unison**, **Detune** and **Drift**. See [Unison](#).
- **POLY** — the voice count (1–64, default 32). See [Voices](#).

Mode and Type

Two controls. **Mode** picks the circuit; **Type** picks the response you take out of it — and Type only ever offers the responses that circuit can honestly produce.

Mode	Slope	Character	Types it offers
DIG (default)	12 dB/oct	A clean state-variable filter. Precise and neutral, and what most presets are voiced against.	LP, HP, BP, NOTCH
MOG	24 dB/oct	A four-pole Moog-style ladder. Warmer, and it thins as it resonates.	LP, HP, BP
303	24 dB/oct	The diode ladder from the TB-303. Asymmetric clipping adds even harmonics — the acid voice.	LP, HP, BP
COMB	—	A delay summed back on itself: a whole series of peaks instead of one corner. Metallic and hollow.	POS, NEG, FF

Why the ladders have no NOTCH. A ladder's resonance feedback fills a notch's null in, and on the diode model at high resonance the null inverts into a *peak*. A notch that becomes a bump is not a notch, so under MOG or 303 the button is not there — rather than being there and quietly handing you the lowpass, which is what this used to do.

The ladders' HP and BP are the real thing, not the lowpass relabelled: a ladder is four one-pole filters in a feedback loop, and the other responses come out of its stage taps the same way the Oberheim Xpander derives its modes.

A second thing worth knowing about the ladders: they *lose* level as resonance climbs, rather than growing a resonant peak on top. Turning Q up on MOG or 303 thins and quietens the sound. That is faithful to the hardware, and it is why the TB-303 engine compensates with a dedicated accent gain.

The comb, and its two borrowed knobs

Under **COMB** the filter delays the signal and adds it back to itself. The delayed copy reinforces every frequency whose period fits the delay and cancels the ones that fall between, so instead of one corner you get a series of evenly spaced peaks — which is why it sounds like a plucked string or a hollow tube rather than a filter.

Its three types are three different instruments:

Type	What it does	Sounds like
POS	Peaks on every harmonic of the tuning	a plucked string
NEG	Peaks on the ODD harmonics only	a stopped pipe, a clarinet
FF	No feedback: notches instead of peaks	a flanger frozen mid-sweep

POS and NEG differ by a single sign and sound nothing alike — cancelling the even harmonics is what makes a clarinet a clarinet.

Two knobs mean something else while COMB is selected, and it is worth knowing before you reach for them:

- **Cutoff is the comb's TUNING** — the frequency its peaks are spaced by, not a corner frequency. Sweeping it slides the whole series.
- **Resonance is the feedback** — how much comes back round, so how long it rings and how sharp the peaks are. Under **FF** there is no feedback path at all, so it sets how deep the notches cut and cannot ring however far you push it.

Two filters, and how they are wired

FILTER B is a second filter with its own Mode, Type, Cutoff and Res. It is off until **Routing** says otherwise:

Routing	What comes out
Off (default)	Filter A alone. Filter B is not even built.
Series	A feeds B — two filters in a row, steeper and darker.
Parallel	Both filters see the same signal and the results are summed.
Difference	A minus B.

Blend always means the same thing: how much of B is in the result. At 0 all three modes sound exactly like Off, so you can bring the second filter in by hand — or put an LFO on Blend and have the routing itself breathe.

Difference is the one worth explaining. Subtracting one lowpass from another leaves only what sits between their two cutoffs: a band-pass whose two edges you set separately, with its own resonance on each. No single circuit here produces that, and it is why having the same filter in both slots is useful rather than redundant.

A comb added to a filter needs no special setting — that is Filter A = DIG, Filter B = COMB, Routing = Parallel, which IS a sum. Series combs what the filter left, and Difference removes exactly what the comb reinforces.

Track (0–1) decides how filter B moves. Everything that sweeps filter A — its envelope and its key tracking — is expressed as a ratio, and Track is how much of that ratio B follows:

- **0** — B stays exactly where its knob puts it. The classic fixed high-pass sitting under a low-pass that sweeps.
- **1** — B moves by the same ratio as A, so the interval between the two stays constant in octaves. Two formants sweeping as a block.

Oscillator extras: PW and Sync

PW (Pulse Width) — 0.05 to 0.95, centred at 0.5. On its own it reshapes a square wave from a hollow, nasal thin pulse through the full-bodied square at 0.5 and out the other side. The classic use is **PWM**: assign an LFO to `osc1.pw` and the width breathes back and forth, turning a static square into a wide, shimmering pad. This is the single most rewarding modulation target on the engine.

Sync (hard sync) — select **Sync** as the waveform and the oscillator runs a second, faster oscillator that is force-restarted on every cycle of the first. The restart chops the waveform mid-cycle, producing the aggressive, tearing timbre of a synced lead. Here the **Sync knob (1–8)** sets the ratio between the two, and *that ratio is the timbre* — the pitch you hear still follows the note. Sweeping the Sync knob (or modulating it) gives the classic sync-sweep lead. Presets **LEAD Sync**, **LEAD Sync Sweep** and **BASS Sync** are built on it.

Ring modulation

Ring (0–1, default 0) multiplies OSC 1 by OSC 2 and mixes the result in as a source of its own, alongside Sub and Noise — it does not replace either oscillator. Multiplying two tones produces their **sum and difference** frequencies instead of their harmonics, so the result is inharmonic: bells, gongs, metallic clangs and robot voices.

Two things drive the sound:

- **The interval between the oscillators.** A few cents of Osc2 Det gives a slow beat and a tone an octave up; wide detunes turn it clangorous. The interval *is* the timbre, so a note played higher changes the character as well as the pitch — which is exactly why ring mod sounds inharmonic.
- **The oscillator levels, which Ring ignores.** The product is taken before the level knobs, so you can pull OSC 1 and OSC 2 down to zero and hear the ring modulator alone — the classic bell — or leave them up and let the metal sit on top of the raw oscillators.

Ring is continuous, so it is a modulation target like any other: an ADSR on `ring.level1` gives a metallic attack that decays into a clean tone, and a slow LFO breathes the clang in and out.

Unison, Detune and Drift

In the MASTER section:


Parameter	Description
Unison	Stack 1–7 copies of the whole voice per note (default 1)
Detune	How far the stack spreads apart, 0–50 cents (default 25)
Drift	Slow, random analogue pitch wander, 0–1 (default 0)

Unison is what makes a lead enormous: seven slightly-detuned copies beating against each other is the supersaw sound. It costs CPU in proportion — seven copies is seven times the synthesis work per note — so raise it for leads, not for dense pads.

Unison is read once when the note is triggered — it is one of the structural parameters listed in [Turning a knob while a note is sounding](#), so it is deliberately *not* a modulation target and changing it affects the next note you play, not the one currently sounding. **Drift** is the opposite: a touch of it (0.1–0.2) keeps sustained chords from sounding sterile.

For modulation routing see [Modulation & Note FX](#).

FM



Above: FM editor — Algorithm selector, Op 1–4 (Ratio/Detune/Level/ADSR), global Mix and Voices.

The FM engine is a four-operator, DX7-style frequency-modulation synthesiser. Each operator is a sine oscillator with its own ADSR amplitude envelope and level; operators are wired together according to one of four algorithms.

FM parameters

Parameter	Description
Algorithm	1 = serial 4→3→2→1; 2 = three parallel mods → Op 1; 3 = two pairs; 4 = additive
FB (Op4)	Op 4 self-feedback — adds odd harmonics and edge
Op 1–4: Ratio	Frequency ratio relative to the played note (0.1–16×)
Op 1–4: Det	Per-operator detune in cents
Op 1–4: Level	Carrier output level or modulation index (modulators)
Op 1–4: ADSR	Per-operator amplitude envelope
Mix	Final output level (0–1)
Voices	Polyphony cap (1–64; default 24). See Voices

FM suits metallic tones, electric pianos, bells, and evolving textures. Small ratio changes yield very different timbres; the preset library covers bells, organs, and electronic basses.

Wavetable



Above: Wavetable editor — Wave A/B selectors, Morph, Detune, Filter, Amp envelope, and Voices.

The Wavetable engine morphs between two pre-computed waveforms — Wave A and Wave B — using the Morph knob or an LFO/ADSR routed to it. Both waves are drawn from a fixed bank of eight anti-aliased tables: Sine, Triangle, Sawtooth, Square, PWM 25%, Organ, Brass, and Vocal.

Wavetable parameters

Parameter	Description
Wave A / Wave B	Source waveforms to interpolate between
Morph	Crossfade position (0 = full Wave A, 1 = full Wave B)
Detune	Global detune in cents
Cutoff / Res	Resonant low-pass filter
Built-in Env	Toggle the built-in amp ADSR on/off
Attack / Decay / Sustain / Release	Amplitude envelope
Voices	Polyphony cap (1–64; default 32). See Voices

Animating Morph with an LFO is the signature technique — sweeping from Sine to Sawtooth or Brass to Vocal while a note sustains produces evolving, living tones. See [Modulation & Note FX](#).

Karplus-Strong

ENGINE

GUIAR

Karp

PRESET

GTR Nylon Soft Fingerpick

LOAD

SAVE AS...

DELETE

SOUND

STRING

DAMPING

BRIGHTNESS

STRINGS

SPREAD

EXCITE

NOISE TONE

STACK

Unison

8.00t

0.62

0.38

1.00

0.02s

0.20

AMP

BUILT-IN ENV

OFF

ON

ATTACK

RELEASE

LEVEL

POLY

VOICES

0.01s

1.40s

0.79

10.00

MODULATORS

+ ADSR

+ LFO

+ CURVE LFO

+ CURVE ENV

+ DRIFT

+ PER-NOTE

+ S&H

+ STEPS

LFO1

FREE

RATE

240 bpm

ON

0..1

-1..+1

FREE

NOTE

VOICE

Damping

+ DESTINATION

ADSR1

A

D

S

R

10ms

360ms

70%

360ms

ON

Damping

+ DESTINATION

NOTE FX

+ ARP

+ CHORD

+ RANDOM

+ ADD INSERT

FX

COMP

channel compressor

THR

RAT

ATK

REL

KNEE

HKUP

GR

BYP

-24.0

4.0:1

3ms

250ms

+30.0

1.00x

0.0

SIDECHAIN

duck this lane

off

off

Above: Karplus-Strong editor — String section (Damping, Brightness), Excite section (Excite time, Noise Tone), and Amp controls.


Karplus-Strong is a physical-modelling engine that synthesises plucked-string sounds. Loom renders each note offline (sample-by-sample in JavaScript) into a buffer and plays it back through an amplitude envelope. This gives exact pitch at every frequency, natural high-harmonic roll-off, and no feedback runaway.

Karplus-Strong parameters

Parameter	Description
Damping	T60 decay: 0 = long sustain (~4 s), 1 = muted (~0.12 s)
Brightness	Loop filter: 0 = dark/cello, 1 = open/metallic
Excite	Excitation burst length (pluck sharpness)
Noise Tone	Colour of the excitation noise (dark → bright)
Built-in Env	Toggle the built-in amp envelope on/off
Attack / Release	Amp envelope on buffer playback
Level	Output amplitude
Voices	Polyphony cap (1–64; default 32). See Voices

Damping and Brightness are set per-note at the moment of the pluck (baked into the buffer). Level and its envelope stay live and can be modulated. Karplus suits acoustic bass, guitar, harp, and marimba-style sounds.

West Coast



A Buchla-style "West Coast" voice: complex oscillator → wavefolder → low-pass gate, driven by a built-in AD contour.

The West Coast engine takes a fundamentally different approach to synthesis from the filter-based ("East Coast") engines above. Instead of shaping a harmonically-rich waveform by subtracting frequencies with a filter, it **adds** harmonics by routing an oscillator through a **wavefolder** and then taming the result with a **low-pass gate**. The result is highly percussive and organic — metallic plucks, woody tones, evolving bell-like pads, and abstract textures that are difficult to achieve with conventional subtractive synthesis.

The engine is **polyphonic** — up to 64 voices, 32 by default; set Voices to 1 in the AMP section for the classic one-voice-at-a-time behaviour. FM here is native-linear, not through-zero.

The signal chain is: **Complex Oscillator** → **Timbre (Wavefolder)** → **Low-Pass Gate (LPG)** driven by an **AD Contour**.

COMPLEX OSCILLATOR section

Two cross-coupled oscillators produce the raw material. The principal oscillator's frequency is modulated by the modulator oscillator (linear FM), optionally ring-modulated, and a sub-harmonic divider can be added below.

Parameter	Description
Princ Wave	Waveform of the principal oscillator: Sin, Tri, or Saw
Mod Wave	Waveform of the modulator oscillator: Sin or Tri
Ratio	Frequency ratio of the modulator relative to the principal (0.25–16×) — integer ratios produce harmonic tones; non-integers produce inharmonic, bell-like partials
FM Index	Depth of linear FM from the modulator into the principal (0–1) — higher values add more sidebands
Ring/AM	Amount of ring modulation mixed in (0 = off, 1 = full ring mod)
Sub ÷	Sub-harmonic divider: Off, 2, 3, or 4 — it divides the frequency, so 2 is an octave below, 3 is an octave and a fifth below, and 4 is two octaves below. 3 therefore adds a fifth rather than a plain sub
Sub Lvl	Output level of the sub-harmonic oscillator (0–1)
Detune	Fine-tune of the principal oscillator in cents (± 50 ¢)

TIMBRE section (wavefolder)

The wavefolder processes the summed oscillator signal through a non-linear waveshaping curve. As the **Fold** amount increases it drives the signal harder into the curve, folding the waveform back on itself and adding a cascade of new harmonics. Accent (velocity ≥ 100) pushes the fold drive harder automatically.

Parameter	Description
Fold	Drive into the fold curve (0 = gentle/clean, 1 = heavy folding/maximum harmonics)
Symmetry	DC bias applied before the folder — shifts the waveform asymmetrically for even-harmonic colouring (–1 to +1)

LOW-PASS GATE section

A low-pass gate combines a resonant filter and a VCA in a single, vactrol-like element so that the contour simultaneously opens the brightness and the volume. The **Mode** selector determines how the contour is routed.

Parameter	Description
Mode	LP — contour sweeps the filter only (VCA stays open); Gate — contour opens the VCA only (filter stays at its base cutoff); Both — contour drives both, the most classic "plonky" Buchla behaviour
Cutoff	Base cutoff frequency (0–1, exponential scaling from ~60 Hz to 18 kHz)
Resonance	Filter Q (0–1) — adds resonant emphasis at the cutoff

CONTOUR section

A single AD (attack–decay) envelope generator, similar to the contour generator in a Buchla 281. It drives both the LPG filter and VCA according to the Mode setting above.

Parameter	Description
Mode	Pluck — decay is gate-independent (fires and fades regardless of note length); Sus — holds at peak until the note ends, then decays
Attack	Rise time (1 ms – 2 s)
Decay	Fall time (5 ms – 4 s) — in Pluck mode this is the T60 envelope; in Sus mode this is the release time after gate-end
Amount	Peak level of the contour (0–1) — scales how far the LPG opens
Cycle	On — re-triggers the AD shape repeatedly while the note is held, turning the contour into a free-running LFO-like modulator

AMP section

Parameter	Description
Level	Master output gain (0–1)
Tune	Global pitch offset in semitones (± 12 st)
Voices	Polyphony, 1–64 (default 32). See Voices

The engine ships with **24 presets** — spanning percussive bass plucks (BASS Fold Sub, BASS Growl FM), bells (BELL Metallic, BELL Crystal Ring), pads (PAD Fold Drone, PAD Glass Air), keys (KEYS Fold E-Piano, KEYS Marimba Fold), and abstract textures (FX Cycle Burst, FX Sci-Fi Cycle) — selectable from the lane's preset dropdown. Per-section knob accent colours group the COMPLEX OSCILLATOR, TIMBRE, LOW-PASS GATE, and CONTOUR sections visually.

Sampler



Above: Sampler editor — global Gain/Voices controls; per-pad controls appear once samples are loaded.

The Sampler engine plays back audio samples mapped across the keyboard. Each keymap zone (pad) has its own per-pad parameters read at trigger time, making it possible to tune, filter, and pan individual pads independently.

Global parameters

Parameter	Description
Gain	Master output gain for the lane
Voices	1–16, default 8 — but the Sampler does not honour it . It runs in a voice pool of its own with no cap; the control is there and does nothing. Every other engine's Voices knob is a real limit


Per-pad parameters

Shown in the drum-voice rack once pads are mapped.

Parameter	Description
Tune	Transposition in semitones (–24 to +24)
Cutoff / Res	Per-pad resonant low-pass filter
Attack / Decay	Per-pad amplitude envelope
Level	Per-pad output level
Pan	Stereo position
REV / DLY	Send amounts to Send B (a Reverb by default) and Send A (a Delay by default)
Loop / Loop Start / Loop End	Loop mode (one-shot or loop-while-gated) and the loop region (start/end as a fraction of the sample)
Sample Start / End	Trim the played window — start/end as a fraction of the sample; draggable on the waveform in the Selected sample panel
Retrig	Poly (voices overlap) or Mono (re-hit cuts the previous voice)

A Sampler lane with pads mapped to GM drum note numbers automatically enters drumkit mode and shows the drum-grid editor instead of the piano roll. For details on loading samples and building keymaps see [MIDI & Samples](#).

Drums (Drum Machine)



Above: Drums editor — Preset row, master bus knobs (Vol/Pan/A/B/Lo/Mid/Hi), per-voice rack with each voice's own synthesis and mixer controls.

The Drums engine is a fully synthesised **ten-voice** drum machine — no samples required. All ten voices (Kick, Snare, **Rimshot**, Closed Hat, Open Hat, Clap, Cowbell, Tom, Ride, **Crash**) are built from oscillators, noise generators, and simple envelopes. Each voice routes through its own channel strip and then to a shared drum bus.

Voice synthesis rack

The per-voice rack exposes the key parameters for each voice:

Voice	Key parameters
Kick	Tune, Attack (click), Decay, Tone, Snap, Thud, Boom, Body — plus Start/End Freq, Sweep, Wave, Drive and the three shape controls under ► advanced . See Building a kick
Snare	Tune, Tone body, Snap (noise), Body Decay, Noise Decay, Noise Tone
Rimshot	Tune, Decay, Freq
Closed Hat	Tune, Decay, Filter
Open Hat	Tune, Decay, Filter
Clap	Tone, Decay, Sharp (filter Q)
Cowbell	Tune, Decay, Detune
Tom	Tune, Decay, Sweep, End Freq
Ride	Tune, Decay
Crash	Tune, Decay

Each voice also has its own mixer row — **Level, A, B, Pan, Lo, Mid, Hi** — where **A** and **B** are that voice's send amounts to Send A (a Delay by default) and Send B (a Reverb by default).

Building a kick

The kick used to be a pitch sweep with a 15 ms click on top, and the click had to fight its way out from under the body's own envelope — which meant no layer could outlast it. It now carries **four layers with lives of their own**, plus a filter and a saturator across the sum of them. Everything here is **off by default**, so no kit you already have sounds any different.

Control	Range	Default	What it adds
SNAP	0–1	0	A noise transient at the very front — the beater. Unlike the old click it has its own decay, so it can die instantly or ring on past the body
SDEC	2–150 ms	20 ms	How long that transient lasts (<i>advanced</i>)
THUD	0–1	0	A sine an octave above where the sweep lands — the knock, the part you hear on a laptop speaker
BOOM	0–1	0	A sine an octave below it, on a tail one and a half times the kick's own decay — the weight. It grows with the kick rather than sitting at a fixed length
BODY	0–1	0	The resonant shell: noise through a bandpass, the sound of the drum rather than of the beater
BCTR	80–1200 Hz	220 Hz	Where that shell resonates (<i>advanced</i>)
BLN	10–600 ms	120 ms	How long it rings (<i>advanced</i>)
TONE	200–12 000 Hz	12 000 (open)	A 24 dB/oct lowpass across the whole kick. At the top it is an exact bypass
DRIVE	0–1	0	Saturation after the filter, normalised so the peak stays where it was — it changes the shape, not the level (<i>advanced</i>)

TONE, SNAP, THUD, BOOM and BODY sit in the kick's front row; the four marked (*advanced*) are in its collapsed ► **advanced** block, along with Start/End Freq, Sweep and Wave.

A practical order: set the pitch and length first (Tune, Decay, Sweep), then **THUD** until it speaks through the mix, then **BOOM** for weight, then **SNAP** for attack, and only then **BODY** if the drum should sound like a drum rather than a tone. **TONE** and **DRIVE** are the finish, not the start.

Drum preset dropdown — synth kits and sample kits

The preset dropdown for any drum lane lists kits from five groups:

Group	Kits	How it works
GM	KIT Standard, KIT Room, KIT Power, KIT Electronic, KIT TR-808, KIT Jazz, KIT Brush, KIT Orchestra	GM-programme aliases that map to the synth kits below
Synth	TR-909, TR-808, TR-606, CR-78, LinnDrum	100% synthesised DSP — no samples required
Samples	TR-808 (samples), Acoustic (samples), Dirt (samples)	Real one-shot WAVs bundled with Loom
Drum Machines	64 sampled kits from classic boxes (Roland, LinnDrum, Korg, Oberheim, Casio, E-mu...)	Sampled one-shots from the tidal-drum-machines collection — a large library of vintage drum-machine sounds
Percussion	GM Percussion	A 31-pad General-MIDI percussion set (sampled)

Synth kits seed every voice's parameters from the kit's characteristic values; you can edit individual voices on top and hit  Sound to randomise all voices at once.

Sample kits load the matching WAV for each voice (kick, snare, closed hat, open hat, clap, tom, cowbell, ride) from `public/drumkits/` and rebuild the keymap fresh on every session load — you never need to re-import the files manually. Once a sample kit is selected, the lane uses the drum-grid editor and the full per-pad parameter rack, exactly like a Sampler lane in drumkit mode.

The sample WAVs are curated one-shots from the Dirt-Samples collection (used by TidalCycles), classic TR-808 recordings, and — for the **Drum Machines** group — the [tidal-drum-machines](#) library. Full credits are in the repo `README.md` under "Credits — sample sources".

Note: sample kits are loaded by the Sampler engine under the hood. For the Sampler's own PRESET dropdown — two groups now, **Melodic** and **Loops**, each sorted by name — and its per-pad parameters, see [MIDI & Samples](#). Drum kits are no longer offered there: they are this page's shelf, and a Sampler lane already on one keeps playing it.

Bus controls


The master bus row gives Vol, Pan, A, B, Lo, Mid, and Hi (± 18 dB) for the whole drum bus — **A** and **B** being the bus's sends to Send A and Send B. Lo and Hi are shelves; Mid is a peaking band. These are automatable via [Modulation & Note FX](#). Routing to the shared reverb and delay sends is covered in [Mixing & FX](#).


Choke groups

Each drum voice has a **Choke** dropdown in the voice's advanced area of the per-voice rack. Voices assigned to the **same non-zero choke group are mutually exclusive**: triggering one immediately cuts the tails of all other voices in the same group. A voice assigned to group 0 (the default) is never choked by anything.

The default configuration is **closed hat and open hat both in group 1**, which gives the classic hi-hat choke behaviour — hitting the closed hat silences the open hat's ring, just as on a real drum kit. You can reassign any voice to any group number to create other exclusive pairs, for example a conga and cowbell that cannot overlap.

Layers





A **Layers** lane does not synthesise anything itself. It holds **up to four other instruments** and plays them together, from one set of clips, through one mixer channel. A piano with a pad under it, a bass whose top octave is a different synth, a lead that is three detuned engines at once: all of it without spending four lanes and four channel strips on one part.

Three ways in. Pick *Layers* from the + menu when you add a lane, pick it from the **ENGINE** selector on a lane you already have, or — the useful one — right-click a lane header and choose **Convert to layered**. Converting keeps the lane's current sound, patch and modulators and puts it in the first slot, with the second slot holding the same instrument at zero gain: the lane sounds exactly as it did a moment ago, and you have somewhere to build.

The rack. Four numbered tabs, **1 2 3 4**, sit above the ordinary parameter grid. Click one and the panel below shows **that slot's own instrument** — its real page, its real sections, the same knobs you would see on a lane of its own. Under the tabs is the **Instrument** dropdown for the open slot (– empty – clears it) and its **Preset** dropdown.

What a slot may hold is narrower than what a lane may be: the six melodic engines only. The Sampler, the Drum Machine, the Audio channel and Layers itself are not offered — the first three run in machinery of their own that a rack cannot reach into, and a rack inside a rack has no bottom. This is a real restriction rather than an oversight: it exists because a slot naming the Sampler played silence.

Each slot has three controls of its own, in the LAYER section:

Control	Range	Default	What it does
Gain	0–2	1	That slot's level within the rack. It goes above 1 on purpose — a quiet instrument used to sit at the top of its fader with nowhere left to go
Low / High	0–127	0 / 127	The range of notes that reach this slot. Leave them alone and every note does

Overlapping ranges are the point: two slots that both cover the middle both sound there. Set them to meet rather than overlap and you have a split keyboard instead of a stack.

The MIX control sits between the tabs and the slot pickers, and its shape follows how many slots are loaded. With **two or three**, it is a horizontal fader with the loaded slots numbered underneath as stops: it crosses between them at constant power, only ever two at a time, and parking it on a stop plays that instrument alone. With **four**, it becomes a square with a slot in each corner — the same square the WEAVE cloud and the WEAVE sound control use. With fewer than two slots loaded there is nothing to balance, so nothing is drawn.

It is a gesture over the ordinary Gain params rather than a control of its own: it writes them, reads its position back from them, and so a rack balance automates, undoes and saves like any other knob. Empty slots take no share.

Level the slots is a checkbox, on by default. Loom's presets span some 46 dB end to end, so a rack of four sounds picked for their character will usually be badly out of balance before you touch anything. With this on, each slot is corrected toward a common loudness — by measurement, not guesswork, and never by more than ± 12 dB — so the faders start meaning *balance* rather than *compensation*. Turn it off when the difference between two instruments **is** the arrangement.

Saving a rack. *Save As...* on a Layers lane stores which instrument is in each slot as well as all four sets of knobs, so recalling it rebuilds the rack rather than dropping four sets of values onto whatever happened to be loaded. Presets saved before this still load; they simply carry knobs alone.

Which slot plays a note is normally decided by the note's pitch and the slots' ranges. There is a second way, and WEAVE uses it: a note can carry the slot it belongs to, in which case it goes there whatever its pitch. That is what lets one part be played by several instruments — see [WEAVE](#).

Changing which engine sits in a slot rebuilds the lane, which takes a moment and resets that slot's parameters to the new engine's defaults. Moving a gain or a range does not.

Summary table

Engine	Best for	Standout parameters
TB-303	Acid bass lines, resonant leads	Slide, Accent, Env
Subtractive	Pads, leads, basses, general-purpose	Dual OSC detune, Filter Drive, Key Track, Unison
FM	Bells, electric pianos, metallic textures	Algorithm, per-operator Ratio, FB feedback
Wavetable	Evolving tones, digital leads, pads	Morph (A→B crossfade), 8-waveform bank
Karplus-Strong	Plucked strings, guitar, harp	Damping (sustain), Brightness, Excitation
West Coast	Percussive plucks, metallic tones, evolving textures	Wavefolder (Fold), LPG Mode, Contour Cycle
Sampler	Any audio, drum kits with per-pad control	Per-pad Tune/Filter/Envelope, Loop mode
Drums	Synthesised drum machine, percussion	10-voice synth rack, the kick's four layers, kits-as-presets, bus EQ, Choke groups
Layers	One part on up to four instruments at once	Four slots, per-slot Gain and key range, the MIX gesture, Level the slots

Editing Clips

Every clip in Loom holds a sequence of notes. The editor always belongs to the **selected lane** — the same one the instrument controls below it are showing, however you got there ([One lane is selected, everywhere](#)). Its header carries the clip's ► and ● **Rec**, and the lane's **M** and **S**, so you can mute or solo the part you are editing without going back to the mixer.

To edit those notes, click the **body** of any filled cell in the session grid — anywhere except the ► play icon or the ✕ delete cross in its corner (clicking ✕ deletes the clip outright, with no confirmation). You can also **right-click** a filled cell and choose **Open editor**. The inspector panel opens below the grid and the editor renders inside it. Closing the inspector does not stop playback — launching and editing are independent.

Melodic lanes (TB-303, Subtractive, FM, Wavetable, Karplus, West Coast, Sampler) open the **piano-roll**. Drum-machine lanes and sampler lanes that have a drum kit loaded open the **drum-grid**. If you want to switch between the two views for a given clip, click the button in the inspector toolbar that names the *other* view — it reads **View as grid** while you are in the piano roll and **View as piano roll** while you are in the grid. Its tooltip states the important part: "Change how this clip is EDITED; it does not change the sound."

See [Sessions, Lanes, Clips & Scenes](#) for how clips are organised, and [Engines](#) for the controls each engine exposes.

Key, scale & musical assistance

Project Options

×

PROJECT

Name

Minimal Techno

KEY & STYLE

Root

C ▾

Scale

🌙 Dark / tense — minor · the classic acid/techno sound ▾

Style

Acid Techno ▾

Lock notes to key

☐

THIS COMPUTER

Autosave

☐

Overwrites the single recovery copy as you work. Your named saves are never touched.

CLOSE

The project key/style bar sits in the transport strip next to Meter; the 🔒 Scale button and the 🎛️, Vary, Mirror, Reverse, and Chords controls appear in the clip-editor toolbar.

Loom's musical-assistance layer works from a single shared **project key and style** setting. Everything that generates or reshapes notes — the random generator, the examples gallery, the pattern transforms, the chord harmoniser — reads that setting and stays in key automatically. You can override it on a per-lane basis when one part needs to inhabit a different harmonic world.

Project key & style bar

In the top transport bar, next to the **Meter** control, a button displays the current project tonality and the scale-lock state — for example 🎵 **A minor · Acid / Techno** · 🔒. The 🔒 / 🔓 glyph at the end tells you at a glance whether the scale lock is on, without opening anything. Click the button to open the key/style panel.

The panel has three pickers and a lock toggle:

- **Root** — the tonic note (C through B).
- **Scale** — chosen by feel rather than by musical-theory name. Each option shows a mood label, the musical name in small print, and a one-line usage hint:
 - 🌙 Dark / tense — minor · the classic acid/techno sound
 - 🛡️ Safe (almost anything fits) — pentatonic minor · hard to sound wrong; riffs & basslines
 - 🌻 Bright / open — major · pop, most "happy" music
 - 🔥 Mysterious / hypnotic — phrygian · dark acid, EBM
 - 🌀 Groovy / swung — dorian · house & funk

-  Anything goes (no net) — chromatic · any note; no safety net
- **Style** — the rhythmic and harmonic character used by the generators. There are **20**: Techno, Acid Techno, Dub Techno, House, Deep House, Garage, Trance, Psytrance, EDM, Breakbeat / Big Beat, Drum & Bass, Jungle, Dubstep, Electro, Synthwave, IDM, Glitch, Downtempo, Lo-fi and Ambient. Style affects which rhythmic patterns the chord maker lays down, sets the default feel of generated basslines and melodies, and selects which library of ready-made patterns the **Pattern** dropdown offers.
- **Scale lock** — a checkbox that turns the piano-roll's note snapping on or off **globally**. It is **off by default**, so you can play freely; tick it when you want every placed note pulled into the key. This is the same switch as the piano-roll's  button — wherever you toggle it, the other reflects it.

Every other musical-assistance feature — the piano-roll scale highlight, the generator, the examples, the transforms, and the chord maker — respects the active root and scale. Closing the panel without choosing commits no change.

Per-lane override

The lane inspector shows a line such as **Key: inherits A minor · Override**. Click **Override** to assign a different root and scale to that lane alone, so one part can be in a different key while the rest of the session stays in the project key.

Scale highlight & lock (piano-roll)

Inside the piano-roll, in-scale pitch rows are highlighted and the keyboard strip on the left colours in-scale keys — the tonic row and key are the brightest. This gives you an instant visual map of which notes belong to the current key.

A  **Scale** button in the piano-roll toolbar controls the scale lock, mirroring the **Scale lock** checkbox in the project key/style panel — both write the same global setting. The lock is **off by default**: a fresh session never constrains what you play, and a saved session always loads unlocked, so the lock can never surprise you by being on. Switch it to  whenever you want a safety net. When locked:

- Drawing a note with the pencil snaps the pitch to the nearest in-scale degree.
- Dragging an existing note vertically skips over out-of-scale rows.
- Computer-keyboard note input plays and records only in-scale pitches.
- Paste lands notes on in-scale pitches, transposing each pasted note by the minimum interval needed to put it back in key.

With the lock on it is impossible to place a wrong note; toggle it back to  (the default) for chromatic freedom — accidentals, blue notes, deliberate dissonance. The arrow-key nudge (semitone up/down) intentionally ignores the lock in both states, so you can always make precise micro-adjustments with the keyboard.

Generate ()

The 🎲 (dice) button in the clip-editor toolbar fills the current clip with a freshly generated musical part in the project key, scale, and style:

- **TB-303-style lanes** — a characteristic bassline (slide and accent placement influenced by the active style).
- **Drum lanes** — a beat pattern matched to the style's rhythmic vocabulary.
- **Other melodic lanes** — a melody or counterpoint line.

The result is not flat random: it is constrained to the scale, shaped by the style, and proportioned to the clip length. Press 🎲 again for another variation; change the global **Style** in the key/style panel for a different character while keeping the same root and scale.

Pattern library

Two dropdowns in the toolbar work as a pair:

- **Style** picks a library — one of the 20 styles listed above. Note that changing it here also changes the **project** style, so the 🎲 generators follow suit.
- **Pattern** lists that style's ready-made patterns, roughly twenty per style. The list is a native dropdown, so you can **type to search** it instead of scrolling.

The list is filtered by editor type — melodic editors offer basslines and riffs, drum editors offer beats — and your own saved examples appear in their own group at the end.

Selecting a pattern:

1. Loads it into the current clip.
2. Transposes it from its stored key into the **current project key**.
3. Repeats or trims it to exactly fill the clip length — short riffs tile; long riffs are cut at the clip boundary.

Two toolbar buttons sit alongside:

- **+ Example** — saves the current clip's notes as your own example. It is stored in this browser (IndexedDB), persists across reloads, and is marked with a ★ in the list.
 - **↓ JSON** — exports the current clip as a portable `.json` file that you can share or import into another browser.
-

Pattern transforms

Three buttons in the toolbar reshape the current pattern — all results stay within the project key and scale.

Button	Name	What it does
Vary	Musical variation	Nudges some pitches to neighbouring scale degrees and lightly adjusts rhythm and velocity; the overall character and contour are preserved.
Mirror	Melodic inversion	Flips the melodic contour so that rising lines fall and falling lines rise, mirrored around the first note of the clip.
Reverse	Retrograde	Plays the pattern backwards in time — the last note becomes the first.


All three are undoable with Ctrl+Z / Cmd+Z. **Mirror** only appears on melodic clips — inverting a contour is meaningless for a drum pattern, where a row picks a voice rather than a pitch — so a drum clip shows Vary and Reverse alone.

Chords

The **Chords** button harmonises the current melodic clip. For each bar it picks the most characteristic **diatonic triad** implied by the bar's melody notes, then writes an accompaniment using a **rhythm pattern drawn from the global style** — for example offset stabs in House, sustained pads in Lo-fi, driving eighth notes in Synthwave.

When you click Chords a small dialog asks **which lane** to write the chord part to: pick an existing melodic lane or let Loom create a new lane labelled "Chords". The original clip is not modified.

Piano-roll



The piano-roll is a two-axis canvas: time runs left-to-right, pitch runs bottom-to-top. A vertical keyboard on the left names the rows; a time ruler at the top marks bars and beats.

Zoom and pan

Scrub **vertically on the time ruler** to zoom the time axis; scrub **horizontally** on the ruler to pan. Scrub the **keyboard strip** vertically to zoom the pitch axis. The native scroll bars pan both axes. Zoom state is saved per clip and restored when you reopen it.

A plain **click on the time ruler** (without dragging) **seeks** the song to that point: the global playback position jumps there and every lane re-anchors in phase, so while the transport runs you can jump around the timeline from any clip editor. Dragging the ruler still zooms/pans as above — only a click-without-drag seeks.

The **drum-grid** and the **audio-clip editor** have the same horizontal zoom and scroll: their fixed parts (the voice-name labels on the drum grid, the waveform header on an audio clip) stay pinned on the left while the timeline content scrolls and zooms beside them. The loop brace tracks the zoom in every editor.

Follow (auto-scroll)

A session-global **Follow** toggle keeps the playhead in view: while it's on, every clip editor auto-scrolls horizontally to follow playback. It is **on by default** — turn it off when you want to inspect one region while the transport runs elsewhere. The setting is shared across all editors, but it is a working mode rather than part of the music: it is **not saved with the session** and returns to on when you reload.

Draw mode (pencil)

The default tool is **Draw**. Click an empty area of the grid to place a note at the snap resolution (default: 16th notes). Drag right while placing to extend its duration. Click and drag an existing note's right edge to resize it. Click and drag the body of a note to move it. Alt-click or right-click a note to delete it.

Select mode

Click the **Select** tool button in the toolbar to switch, or press **2** (**1** switches back to the pencil). In Select mode, click a note to select it. **Drag on the grid background** to draw a marquee rectangle; every note whose body intersects the rectangle is selected. **Ctrl+A** selects every note in the clip; **Esc** or a click on empty space deselects all.

With notes selected you can:

- **Move as a group** — drag any selected note; the whole selection moves together and is clamped to the clip boundaries.
- **Delete** — press Delete or Backspace.
- **Nudge** — use the arrow keys to shift the selection one snap unit left/right or one semitone up/down.
- **Cut / Copy / Paste** — Ctrl+X / Ctrl+C / Ctrl+V (Cmd on Mac). Paste anchors the earliest clipboard note at the mouse cursor position, so move the pointer where you want the paste to land before pressing Ctrl+V.

The tool choice and clipboard contents persist across clip re-opens and across clips, so you can copy from one clip and paste into another.

Computer-keyboard note input

The  **Keys** toggle in the clip-editor toolbar turns your computer keyboard into a live instrument. It is **off by default** — switch it on and the keys below play the active lane in real time:

Key row	Notes
a s d f g h j k (home row) — white keys	C D E F G A B C
w e t y u (upper row) — black keys	C# D# F# G# A#
z / x	Shift input octave down / up

This is **live playing, not step entry** — the keys sound the lane exactly as a MIDI keyboard would, and nothing is written into the clip on its own. (Earlier versions typed notes into the clip at a cursor; that behaviour was replaced by live playing.)

Recording notes into a clip


To capture what you play, use the **• Rec** button in the inspector, next to the clip's **►** play button. Beside it is a mode selector with two entries:

- **Merge** — the notes you play are added to whatever the clip already holds.
- **Replace** — the clip's existing notes are cleared and replaced by what you play.

Click **• Rec** (it becomes **■ Stop** while recording), then play along with the loop using the computer keyboard (with  **Keys** on) or a connected MIDI keyboard. Click it again to finish.

Pressing **R** starts and stops recording too, so you can arm and disarm without reaching for the mouse. Note that **R always records in Merge mode** whatever the selector says — the mode selector is read by the button, not by the shortcut. Like the other shortcuts, R is ignored while you are typing in a text field.

Drum-grid



The drum-grid is a canvas editor where rows correspond to drum voices (kick, snare, hi-hat, etc.) and columns correspond to time positions. Each hit is placed at a precise tick within the clip.

Grid resolution

A resolution selector at the top of the editor sets the snap and the column width. Available resolutions are:

Value	Description
1/4	Quarter notes
1/8	Eighth notes
1/8T	Eighth-note triplets
1/16	Sixteenth notes (default)
1/16T	Sixteenth-note triplets
1/32	Thirty-second notes
free	No snap — place hits at any tick

You can mix resolutions across clips; each clip stores its own `gridResolution` setting.

Euclidean fill (H / S / R)

Every voice row has three small fields between its label and the grid: **H**, **S** and **R**. They generate that voice's pattern by spreading hits as evenly as possible around a cycle — the Euclidean rhythm idea, which produces a surprising number of the world's real drum patterns.

Field	Meaning
H — hits	How many onsets to place, spread as evenly as possible. 0 leaves the voice alone , so you can Euclid one voice without touching the rest.
S — steps	The length of the cycle. Shorter than the clip and it simply repeats; set it to a length that does <i>not</i> divide the clip evenly and the pattern phases , landing differently each time round.
R — rotate	Shifts the cycle so the hits land elsewhere against the beat. Negative values rotate the other way.

Start with **H=4**, **S=16** on a kick for four-on-the-floor, then try **H=3**, **S=8** on a hat for the tresillo that underpins most Latin and house grooves. The classic trick is an off-divisor: **S=7** against a 16-step clip never repeats the same way twice within the bar.

The fill is generated into the clip as ordinary notes, so you can edit them by hand afterwards, and the whole operation is undoable.

Placing and removing hits

Click an empty cell to place a hit. Click an existing hit to remove it. In **free** mode, click anywhere in the row — the hit lands at the exact tick under the pointer with no snapping.

Selection and group operations

Drag on the canvas background to draw a marquee rectangle. Hits whose row and time position intersect the rectangle are selected. With a selection active:

- **Move** — drag horizontally to shift selected hits in time; the group is clamped to the clip length.
- **Move rows** — drag vertically to reassign hits to different drum voices; the relative row offsets are preserved and clamped to the available voice list.
- **Delete** — press Delete or Backspace.
- **Cut / Copy / Paste** — Ctrl+X / Ctrl+C / Ctrl+V. Paste anchors the earliest hit at the click position (tick × row), preserving relative offsets within the group.

Selection, clipboard, and group-move all operate on **row indices**, not MIDI numbers, so patterns copy cleanly even between kits that map voices to different MIDI notes.

Playhead

A vertical playhead line moves across the canvas in real time while the clip is playing, driven by the sequencer's look-ahead clock. It updates on every redraw tick and resets to the left edge when the clip stops.

Clip tempo (×2 / ÷2)

Next to the **Length (bars)** field in the inspector toolbar sit two buttons, ***2** and **/2**, that time-scale the open clip's notes in a single, undoable step — the perceived tempo of the pattern doubles or halves while the clip stays musically useful.

- ***2** (*Double tempo — compress notes & repeat to fill, length unchanged*) — every note's start and duration is halved (so the pattern plays **twice as fast**), and the clip's **length in bars stays the same**: the now-shorter pattern is **tiled** to fill it, so you hear a clean double-time version rather than a half-empty clip.
- **/2** (*Halve tempo — stretch notes & grow the clip length*) — the inverse: notes and durations double (so the pattern plays **half as fast**) and the clip's **length grows** to fit the stretched material.

The loop region and any clip automation scale with the notes. The buttons appear for **note and drum clips**; a pure audio channel has no notes to scale, so they are hidden there. The whole rescale is one entry in the global undo history, so a single Ctrl+Z restores the original timing.

Loop regions

A **loop brace** sits above every clip editor — piano-roll and drum-grid — as a narrow strip spanning the full clip length. It lets you mark an A–B sub-region and repeat just that portion while the clip plays.

Setting the loop region

The brace strip has two drag handles (left = A, right = B) and a **Loop** toggle button. To use it:

1. Click **Loop** to enable the loop region. The region highlights between the A and B handles; if no region was set before, it defaults to the full clip length.
2. Drag the **left handle** to move the start point (A). Drag the **right handle** to move the end point (B). Both handles snap to 16th-note grid positions.
3. Drag the **interior** of the region (between the handles) to **slide the whole window** along the timeline — its length is preserved, it snaps to the grid, and it stops at both ends of the clip. Handy for auditioning the same-length loop over different bars.
4. While the clip is playing, the scheduler repeats only the A–B sub-region — the rest of the clip is skipped.

Click **Loop** again to disable it. The clip reverts to playing its full length; the A and B positions are remembered so you can re-enable the same region later.

What the loop brace affects

The loop region works the same way for all clip types:

- **Note clips (piano-roll)** — only notes whose start falls within the A–B range are triggered; the period of repetition equals the duration of that range.
- **Drum clips (drum-grid)** — same: only hits inside the sub-region fire.

- **Sliced note clips (Sampler)** — a clip whose notes trigger bank slices behaves like any note clip: only the hits inside the A–B sub-region fire, still tempo-locked and pitch-preserving.

The brace lives on the note/drum editor. A pure audio channel (the waveform-only audio-clip editor) does not show a brace; to gain per-region control over an audio loop, bring it in through the Sampler's **Loop** family, which slices it into a note clip with a piano-roll (see [MIDI & Samples — Sampler](#)).

The loop region is **per-clip** and saved with the session. Two clips in the same lane or scene can each have their own independent A–B region, or none at all.

Sharing a loop across the scene

By default each clip's loop is its own. The brace strip also carries a **Global** button (tooltip: *"Share this loop across every clip in the scene"*). Turn it on — it lights amber — to push this clip's A–B region onto **every clip in the scene** so they all loop the same bars together, which is handy for auditioning one section across the whole arrangement. Toggle it off to return each clip to its own independent region.

The scene's shared loop is the **same region** as the Performance view's **A–B** brace: set it here and the A–B brace picks it up when you switch to Performance, and setting A–B there writes it back to the scene. See [Performance & Arrangement](#).

All loop-region edits (moving handles, sliding the region, toggling, and the Global share) are part of the global undo history (Ctrl+Z / Cmd+Z).

For how to use an arrangement-wide A–B loop brace that repeats a section across all lanes at once, see [Performance & Arrangement](#).


Waveform header

Any clip that references an audio buffer shows a **waveform header** above its editor — a peak view of the buffer with a bar/beat ruler, slice markers (orange), and a live playhead. It appears in two situations:

- **Audio clips** open the dedicated **audio-clip editor** (toolbar + waveform header, no note grid).
- **Sliced note clips** (and any normal clip that carries a display-only waveform reference) keep the header **above** the piano-roll or drum-grid, so you can edit notes against the source audio.

To bring a loop into a session as a tempo-locked audio channel, see [MIDI & Samples — Audio channel](#) (the grid's + ► **Audio channel** entry, the waveform header, the **Warp** tempo-lock). To chop a loop into individually editable note slices, load it through the Sampler's **Loop** family instead.

Velocity & dynamics



Above: piano-roll editor. The strip beneath the note grid is the velocity lane — one vertical bar per note, height proportional to velocity.

Every note in Loom carries a **velocity** value from 0 to 127. Velocity is set when you draw a note (default: **90**), adjusted in the velocity lane, and captured automatically from MIDI import. It affects the sound in two complementary ways:

- **Loudness** — velocity scales the note's output gain continuously along a smooth curve with a **floor**: the softest possible note still sounds, at roughly a fifth of full level (about 13 dB down), rather than vanishing. 127 is the loudest a note can be. Notes at the default of 90 sit well up the range, so there is clear headroom in both directions.
- **Accent character** — notes with velocity ≥ 100 are accented. On top of the continuous gain, accent adds character, and what it adds depends on the engine:
 - **TB-303** — brightens the filter envelope, raises the resonance Q, and shortens the filter decay. It also gets a bigger loudness jump than the other engines, because a diode ladder loses level as Q climbs and would otherwise leave the accent quieter than the note it is meant to punch.
 - **Subtractive** — opens the filter envelope wider. Its resonance is untouched.
 - **West Coast** — drives the wavefolder (and its cutoff sweep) harder.
 - **Drums** — purely a louder hit. There is no filter or tone change.

This is the same accent model that the 303 bassline and drum sequencer have always used, now unified into the velocity scale.

Reading velocity visually

Each note's **fill colour** shifts along a **blue** → **yellow** ramp as its velocity increases. Low-velocity notes are deep blue; high-velocity notes are warm yellow. The transition is weighted so the blue half of the range covers roughly velocities 0–64 and the yellow half covers 64–127. Accented notes (≥ 100) are additionally outlined with a **white border** — colour alone does not distinguish accent from non-accent.

The velocity lane

Below the note grid (piano-roll) or the drum-voice rows (drum-grid) is the **velocity lane**: a row of vertical bars, one per note, anchored at the note's start position. Bar height is proportional to velocity. A **dashed horizontal line** across the lane marks the accent threshold (velocity 100) so you can see at a glance which notes are accented. The lane scrolls horizontally in sync with the grid.

Editing velocities

You interact with the velocity lane by dragging the bars:


- **Set a single note** — drag a bar up or down. The velocity updates live; the note colour and audible gain change as soon as you release.
- **Adjust a group** — if you have notes selected (marquee selection in the main grid), dragging any bar that belongs to the selection applies the **same delta** to all selected notes. Notes that would go out of range are clamped to 1–127.
- **Paint a ramp** — drag horizontally across multiple bars. Each bar you pass over is set to the velocity corresponding to the current vertical position of the pointer, writing a smooth velocity ramp across the passage in a single gesture.

When several notes share the same start position (a chord), their bars are fanned a few pixels apart in the lane so each one remains individually grabbable.

All velocity edits are undoable (Ctrl+Z / Cmd+Z) in the same undo history as note placement and movement.

Modulation & Note FX

Every lane has two signal-processing layers that sit *before* its engine's audio output and *around* its note stream: **Modulators** shape sounds over time by continuously driving engine parameters, and **Note FX** transform the note events before they reach the engine. Both are per-lane and saved with the session.



The FM engine editor. The MODULATORS section (LFO1 + ADSR1) and the NOTE FX row are visible at the bottom of the panel.

Modulators

Open any lane's engine editor and scroll to the **MODULATORS** section. Press + **ADSR**, + **LFO**, + **Per-Note** or + **S&H** to add a modulator; you can add as many as you need. Each modulator appears as a card with its own controls, an **ON / OFF** toggle, and a **x** remove button.

That row of buttons is not a fixed list — it is built from whatever modulators are registered, so a plugin that ships one adds its own button. **S&H** and **Per-Note** are **exactly that**: modulators that arrive as external plugins rather than from Loom's own source, and behave like any other.

The section is not empty to begin with. Every engine arrives with its own modulators already wired, and on some engines they are load-bearing rather than decorative:

Engine	Ships with
Subtractive	Four — two ADSRs (which are the amp and filter envelopes) plus two free LFOs
West Coast	Two ADSRs (wavefolder + low-pass gate) plus two LFOs
Wavetable	An ADSR on filter cutoff, plus a free LFO
FM, Karplus	One LFO and one ADSR
TB-303	One LFO only — by design; its envelopes are part of the engine

On Subtractive in particular, deleting the two ADSRs removes the amplitude and filter envelopes: they are not extras sitting on top of the sound, they *are* the sound's shape.

A preset can carry its own modulators. Loading one replaces the lane's modulator set with the preset's — because for some patches the modulation *is* the patch. A wobble bass is an LFO on the filter cutoff; without it there is no wobble. On Subtractive that is not the exception but the rule: all 103 of its factory presets arrive with their envelopes and LFOs wired, which is why loading one and reading its MODULATORS section is the quickest way to learn what this section can do. One of them, **PLUCK Per-Note Pulse**, is built on the Per-Note modulator described below.

LFO

An LFO generates a periodic waveform that you route to one or more target parameters.

Control	Description
WAVE	Waveform shape: Sine, Tri, Sqr, or Saw
RATE	(<i>FREE mode</i>) the free-running rate, set with a log-scaled knob and shown in bpm (LFO cycles per minute), from ultra-slow sweeps (under 1 bpm) up to audio-rate wobble. The slow half of the knob travel covers the slow rates, so gentle sweeps are easy to dial in.
BARS / FEEL	(<i>SYNC mode</i>) BARS is a free numeric input for the cycle length in bars-per-cycle (e.g. 0.25 = a quarter-bar cycle, 4 = one cycle every four bars; any value from 1/16-bar up to 64 bars). FEEL offsets it — Str (straight), Trip (triplet), Dot (dotted). This replaced the old fixed RATIO dropdown, so you can sync to any cycle length, not just preset divisions.
FREE / SYNC	Toggles between the free-running bpm RATE knob (FREE) and the tempo-locked BARS + FEEL controls (SYNC).
POLARITY	-1..+1 (bipolar, default) oscillates symmetrically around the param's centre. 0..1 (unipolar) only pushes the parameter upward.
RETRIG	One 3-way control with three states: Free — one lane-wide LFO whose phase runs continuously off the clock, like a classic analogue LFO. Note — one lane-wide LFO whose phase restarts on every note-on, so its shape lands the same way on each note. Voice — each played note gets its own LFO, born with the note.

Musical difference between the lane-wide settings and Voice: with Free or Note the whole lane shares one LFO, so a chord breathes as one — all notes rise and fall together, which is what you want for a pad that should feel like a single instrument. **Voice** gives each note its own cycle, so notes played at different moments drift out of step with each other and the chord shimmers instead of pulsing. On fast, staccato playing the difference is dramatic; on a slow sustained chord it is subtle.

Free vs Note matters most with a slow LFO and short notes: with Free, a note might catch the LFO anywhere in its cycle, so consecutive notes sound inconsistent; with Note, every note gets the same sweep from the same starting point. (On **Voice** a retrigger setting would be redundant — the LFO already starts with the note — which is why the three states share one control rather than two.)

ADSR

An ADSR produces a classic Attack–Decay–Sustain–Release envelope that fires once on each note trigger and follows the note's gate duration. Every note always gets its own independent envelope — unlike the LFO, an ADSR has no RETRIG control, because per-note is the only behaviour that makes sense for it. Its card carries the four knobs and nothing else.

Control	Range	Default
A (Attack)	1 ms – 2 s	10 ms
D (Decay)	1 ms – 4 s	300 ms
S (Sustain)	0 – 100 %	70 %
R (Release)	1 ms – 8 s	300 ms

S&H (sample & hold)

S&H latches a new random value at a steady rate and **holds** it until the next one. Where the LFO glides, this jumps: it is the stepped, jittery movement of a classic modular random source. Aim it at a filter cutoff for burbling per-step colour, at pitch for a small amount of analogue drift, or at a send for a randomly appearing effect.

Control	Range	Default
Rate	0.1 – 20 Hz	6 Hz
Bipolar	Unipolar / Bipolar	Bipolar

Bipolar swings both ways around the destination's current value; **Unipolar** only adds. Like the LFO it can run shared across the lane or per voice.

The steps are random but not *unrepeatable*: the value of a given step is derived from its index, so the same passage lands identically every time — the offline render and what you heard live agree exactly.

Per-Note

Every other modulator answers *what time is it*. This one answers **which note is this** — it is a function of the note's position in the sequence of notes played, not of the clock. The value is fixed when the note starts and **never moves while the note sounds**, and the same note always gets the same value.

It is the modulator to reach for when consecutive notes should not be identical: a touch on filter cutoff and the line stops sounding machine-stamped; a touch on pulse width and every note has its own body.

Control	Range	Default
Pattern	0 – 1	0.618...
Skew	0 – 1	0
Polarity	Unipolar / Bipolar	Bipolar

Pattern is the instrument. It says how far the value walks from one note to the next: **0.5** alternates between two values, **0.25** cycles every four notes, **0.75** every four but in a different order — and the golden-ratio default, 0.618..., **never repeats**. That precision matters: any short decimal is a fraction, and every fraction comes back round eventually. **Skew** slides the whole sequence along without changing its shape, which is how two Per-Notes on the same Pattern can be made to disagree. **Polarity** follows the S&H convention — Unipolar only pushes the destination one way from where you set it, Bipolar lets a note land under the knob as well as over it.

There is no depth control on the card, because every destination row already has one.

One caveat. Because it counts notes rather than reading the clock, an offline render that starts partway through does **not** reproduce what you heard from the top: the note that was the fortieth is now the first. This is the opposite of the S&H guarantee above, and it is the price of the value depending on which note it is.

Destinations and depth

Below each modulator card's controls is a destination list. To route the modulator:

1. Select a target from the dropdown — it lists all automatable parameters for the lane's engine, the lane's own mixer column (level, pan, both sends, the three EQ bands), plus any lane inserts, master inserts, and master sends (reverb, delay). Modulating the lane's level or a send is how you get tremolo, auto-pan or a rhythmic send; it swings around wherever the fader sits rather than replacing it — see [Automating and modulating the mixer column](#).
2. Press **+ Destination**. A new row appears showing the target name and a **DEPTH** knob (range –1 to +1, default 0.5).

The depth knob scales the modulator's effect. A value of 1.0 means the full output range of the modulator sweeps the full range of the destination parameter (in the parameter's native units). –1.0 inverts the modulation direction. Multiple destinations can share the same modulator at different depth values.

Remove a destination with its **×** button. Remove the whole modulator with the **×** on the card header.

Note FX

The **NOTE FX** section sits directly below MODULATORS. Note FX processors intercept the lane's note stream before it reaches the engine, transforming which notes are played and when. Add a processor with + **Arp**, + **Chord** or + **Random**. Each processor shows an **ON / OFF** button and a × remove button.


Note FX are per-lane and persist with the lane's engine state. Loading a demo resets them to the demo's configuration.

Arpeggiator

The Arp takes each held note and generates a rapid sequence of notes from it according to a scale and direction pattern.

Control	Options / Range	Description
PATTERN	up, down, updown, random, cosmic, steps	Direction of the arpeggio. cosmic adds occasional octave jumps and random steps for an unpredictable feel; steps is a pattern you write yourself — see below.
SCALE	global , major, minor, pentMinor, phrygian, chromatic	Where the arp's notes come from. global is the default and is not a scale of its own: the run walks the degrees of the <i>session</i> 's key upward from the note you played, so it stays in the song. The named scales are transposed onto the played note instead, which is the older behaviour and can leave the key — in C major, playing E gives E-F-G-A-B on global and E-F#-G#-A-B on major .
RATE	free, 1/4, 1/8, 1/8t, 1/16 (<i>default</i>), 1/16t, 1/32	Step rate in musical divisions (BPM-synced), or free Hz.
OCT	1–4	Number of octaves the arp spans, 1 by default : the run stays in the octave you played. Higher values are opt-in and cycle the scale across more octaves before repeating.
GATE	0.05–1.0 (default 0.7)	Fraction of the step interval during which each arp note is held. Lower values create a more staccato feel.
FREE Hz	0.5–32 Hz (default 8)	Rate used when RATE is set to <i>free</i> .

The Arp fires as many notes as fit inside the original note's gate duration, so longer notes produce longer runs.



PATTERN = steps — the pattern you paint. Choosing *steps* reveals a **STEPS** row of bars beneath the controls. Drag across it to paint, exactly as you would an automation lane; **-** and **+** change the pattern's length, from 1 to 32 steps, and lengthening repeats the last step rather than the whole cycle.

What a bar's height sets is **which note of the run**, not which pitch — the pattern is written in positions, so it survives everything that would otherwise break it. Transpose the part, change the key, change the scale, change OCT: the pattern still means the same thing, because SCALE and OCT decide the pool the positions point into and the positions themselves do not move. Eight rungs are available above the floor.

A bar dragged all the way down is a rest. It keeps its slot — nothing sounds there and nothing shifts up to fill it — which is what lets you write a rhythm and not just an order. Every fourth bar is drawn brighter so you can count without a ruler.

The default is **0 1 2 3**, the plain upward walk written out, so switching to *steps* changes nothing until you edit it. A pattern with every bar at rest plays silence; it does not fall back to the walk.

Chord

The Chord processor expands each incoming note into a chord built on that note as the root. Out of the box the note you played is part of what you hear — the octave shift is a switch you throw deliberately, never something the processor does behind your back.

Control	Options	Description
CHORD	maj, min, maj7, min7, sus2, sus4, dim, free	Chord voicing to generate. free replaces the named voicing with three intervals you dial yourself.
INT 1 / INT 2 / INT 3	–24 to +24 semitones (4 / 7 / 0)	Only with CHORD = <i>free</i> : the three intervals above (or below) the root. A zero switches that voice off rather than doubling the root — the root always sounds, so a two-note voicing is one interval at 0.
IN KEY	off (<i>default</i>) / scale / chord	Whether the notes it produces are pulled into the song's harmony. See below.
OCT SHIFT	ON / OFF (off by default)	Enables the octave shift. While off the OCT slider is hidden and completely bypassed, so the chord stays rooted on the note you played.
OCT	–2 to +2	Only with OCT SHIFT on: transposes the whole chord, root included, that many octaves.

All notes in the chord share the original note's timing and gate length. The chord is also built in the **session's** key now: playing live used to be answered in A minor whatever the song said, so the same lane with the same settings was in two different keys depending on whether the note came from a clip or from your fingers.

IN KEY — staying in the key is not the same as landing on the chord.

- **scale** snaps every note to the session's key and scale. You cannot play a wrong note, but you can still play a passing note the chord underneath does not contain — which is often what you want, and is why this is not the stronger setting. Be aware of one consequence: in a major scale every out-of-scale note sits exactly between two neighbours, and a tie resolves upward, so a D major triad in C major comes out D-G-A rather than D minor.
- **chord** snaps to the pitch classes of the chord sounding at that bar of the progression. It cannot sound wrong against the harmony — and it cannot sound like a melody either, because it is locked to three notes. With no progression to read it falls back to *scale*.

Conforming can land two intervals on the same note; the duplicate is dropped rather than played twice.

Random

The Random processor introduces controlled uncertainty into a part — the humanising pass that keeps a loop from sounding like a loop. It is built around **chance sliders**: each of the four things it can vary has its own 0–1 probability, and **every one of them starts at 0**, so adding the processor changes nothing until you ask it to. Turn one up and it intervenes on that fraction of notes.

Control	Range / Options	Description
CHANCE	0–1 (default 0)	How often the pitch is moved
CHOICES	1–24 (default 6)	How many alternative pitches are in the pool
INTERVAL	1–12 semitones (default 1)	The spacing between those choices
MODE	random / alt	<i>random</i> picks freely; <i>alt</i> walks round-robin through the pool, which is more even and less clumpy
SIGN	add / sub / bi (default bi)	Whether it may move up only, down only, or both ways
SCALE	ON / OFF (on by default)	Snaps every generated pitch to a key and scale, so the randomness stays in the song. Switching it on reveals ROOT and SCALE selects, both of which start on Global — they follow the song's own key until you pick something else
VEL CHANCE / VEL RND	0–1 (0 / 0.3)	How often, and how far, the velocity varies. This is the one to reach for first: a little velocity variation is most of what "played by a human" means
VEL SMOOTH	0–1 (default 0)	How much the velocity variation is smoothed from note to note. At 0 each note is drawn independently, which reads as twitchy; turned up, the loudness drifts instead of jumping
VEL DRIFT	0.05–8 Hz (default 1)	How fast that drift moves. Only does anything while VEL SMOOTH is above 0
DUR CHANCE / DUR RND	0–1 (0 / 0.3)	How often, and how far, the gate length varies (at 1, between half and double)
DROP	0–1 (default 0)	How often a note is silenced outright — the fastest way to thin a dense pattern into something with space in it

The randomness is reseeded each time the transport starts, and stable within one run: a pass you liked stays the way it was until you stop and start again.

Automation

Loom records parameter automation in two ways:

Real-time knob recording (Performance view): in the session/I-O header row, make sure the REC mode selector beside **REC** is set to  **take** (the default), then press **REC** to arm recording and press Play. While recording in take mode, every knob you move *and* every clip launch is captured as

automation in the current take. Automation is written at a sub-step resolution derived from the BPM. Press • **REC** again to disarm. (The other two REC modes — 🕒 **live** and ⚡ **offline** — export WAV audio instead of recording automation; see [Transport](#) for the unified REC group.)

Per-clip envelopes (Session view): each clip can carry automation lanes independent of the Performance take system. Open the inspector for a clip and scroll below the note editor to the automation section. Select a parameter from the dropdown and click the add button to create an envelope lane for that clip. Envelopes are stored alongside the clip's notes in the session file.

The quickest way in is not that dropdown at all: **right-click the knob you want to automate**. The menu offers *Automate in clip "<name>"* (or *Edit automation in clip ...* if a lane already exists) and takes you straight to it. A knob that cannot be automated at all — the master strip's, for instance — simply does not open a menu. One that *is* automatable but not reachable right now does open it, with the entry greyed out and the reason on it ("This track has no clips", "Master and send FX automate on the timeline — switch to Performance").

Drawing into a clip envelope

Each lane sits below the note editor with its own header:

- **Line / Flat** — the two brush buttons at the top of the automation section. **Line** draws a ramp between where you press and where you release; **Flat** paints one constant value across the drag. The choice applies to every lane.
- **On / Off** — mutes the curve without deleting it.
- **Smooth / Stepped** — switches interpolation from smooth to a staircase that holds one value per step.
- **x** — removes the lane.

The LFO generator

Rather than draw a repeating shape by hand, each lane's header carries a small generator that writes one for you:

- **Shape** — Sine, Triangle, Saw up, Saw down, Square or Random.
- **Rate** — a musical division of the **bar**: 4 bars, 2 bars, 1 bar, 1/2, 1/4, 1/8 or 1/16. 1/16 means sixteen cycles per bar, the fastest this menu offers.
- **Depth** — 100 %, 75 %, 50 % or 25 % of the lane's full height.
- **Loop only** — when the clip has a loop region, confines the fill to it instead of the whole lane.
- **LFO** — the button that actually draws it. One press is one undoable step.

On a **Stepped** lane the wave is sampled once per step, so a rate faster than the step grid cannot show a full cycle.

The length an envelope runs on

An envelope spans the clip's **full length in bars** and repeats on that period. That matches the notes only while the whole clip is playing. If you shorten what plays — either with the clip's own loop region or with the scene's Global loop — the notes repeat on the shorter window but the envelope keeps running on the full clip, so the curve gradually slides against them. Set the loop first, then draw.

When you move a clip to a lane running a different engine, envelopes whose parameter ID no longer exists in the new engine are disabled automatically (they remain in the clip but do not play back until re-enabled or deleted).

See also: [Engines](#) for the parameters you can modulate, [Mixing & FX](#) for lane inserts and master FX whose parameters also appear as modulation destinations, and [Transport](#) for the REC button and Performance view.

Mixing & FX

Every lane in Loom has its own signal path from the synthesis engine through to the master output. This chapter explains how that path is structured, what controls are available per lane, and how the shared Master FX panel ties everything together.

Signal flow overview


Lane engine → lane insert chain → channel strip (EQ → comp → level → pan → mute → duck) → master bus
master bus → master insert chain → master compressor → output

Send A (gain) →
Send B (gain) →
Send A / Send B returns →

In short: the engine's audio passes through any per-lane inserts first, then the channel strip where EQ, sends, pan, and level are applied. The processed signal joins the master bus, which runs through the master insert chain and the master compressor before reaching the speaker. The two **Send A / Send B** buses are parallel return channels: each lane feeds them a post-fader amount, and each send runs its own insert chain (seeded **A = Delay**, **B = Reverb**) before returning to the master.

Per-lane channel strip

Each lane owns a `ChannelStrip`. Its controls are visible below the session grid in the lane's row.



Clicking a column selects its lane. The selected column is highlighted, and selecting it brings the clip editor and the instrument controls with it — see [One lane is selected, everywhere](#).

Only the column's *background* does that: the track name and the empty space between controls. A click that lands on a fader, a knob, M, S, an insert slot or the preset dropdown does that control's job and nothing else, so moving a fader never yanks you out of the clip you were editing.

Level (fader)

The vertical slider sets the lane's output gain as a linear multiplier (0–1 = silence to unity; the percentage label reflects the current value). This is the last gain stage before output, applied after EQ and the per-lane compressor.

Pan

The **PAN** knob positions the lane in the stereo field. Centre (0) is the default; turning it left or right continuously shifts the image. The pan value is automatable and can be modulated — see [Modulation & Note FX](#).

Mute and Solo

M silences the lane by zeroing its mute gain node. **S** solos the lane: all other lanes are muted in the UI while solo is active. Both controls affect what the sidechain tap feeds downstream — a muted lane's tap still carries pre-mute signal so sidechain routing remains stable.

The same two buttons also sit in the **clip editor's header**, next to ► and ●, acting on the lane whose clip is open — so you can drop a part out while editing it without going back to the mixer. They are the same two switches, not a second pair: muting from either place lights both, and with no clip open the header pair is disabled.

3-band EQ

The three EQ knobs apply before the per-lane compressor:

Knob	Filter type	Centre frequency	Notes
LO	Low-shelf	200 Hz	Boost or cut lows; default ±0 dB
MID	Peaking	1 000 Hz	Q = 1; adds presence or scoops the midrange
HI	High-shelf	4 500 Hz	Boost or cut highs and air

All three bands are ±dB adjustments. EQ gain AudioParams are exposed to modulation so you can automate filter sweeps from the modulation panel.

Send A and Send B

The two send knobs — **A** and **B** — control how much of this lane's post-duck signal is fed into the two shared send buses. They replaced the old fixed REV/DLY knobs. Send A and Send B are general-purpose return channels, seeded **A = Delay** and **B = Reverb**, but you can change what effect lives in each — see [Send A / Send B return modules](#). The knobs are independent wet levels: 0 = dry only, higher values mix more of the lane into that send's effect. Sends are post-fader, so they follow the lane's level and sidechain ducking. (Old saves with `...rev` / `...dly` amounts migrate automatically — reverb→B, delay→A.)

Automating and modulating the mixer column

Level, pan, both sends and the three EQ bands are ordinary lane parameters, which means each of them shows up in the automation destination picker, in the XY pad's axis lists, in the modulation panel's target dropdown and in the knob's own right-click menu — on **every** lane, whatever engine it runs. This used to be true of drum lanes only, so a drum lane's volume could be automated and a Subtractive lane's could not; the mixer now declares its parameters once for everybody. Their automation ids are `<lane>.bus.level`, `<lane>.bus.pan`, `<lane>.bus.delaySend`, `<lane>.bus.reverbSend` and `<lane>.bus.eq.low` / `.mid` / `.high` — the lane id first, exactly like an engine param. (Sessions saved before this carry the old `mix.<lane>.<param>` form for drum lanes; they load, but the ids you see in the picker today are the `bus.*` ones.)

One difference worth knowing between **automating** a fader and **modulating** it: automation writes the gain directly, while a modulator is summed onto a trim that sits in front of it. That is deliberate — a bipolar LFO summed straight onto a gain would push it below zero on the trough, which does not quieten the lane, it flips its phase and makes it cancel against everything else in the mix. With the trim, full depth swings the lane between silence and double its fader position, and the modulation stays relative to wherever you left the fader.

Per-lane inserts

Every lane also has a private insert chain that sits *before* the channel strip — the engine's audio passes through it first. If you open the lane's inspector and add FX to its insert list, those effects process the lane signal exclusively and do not affect any other lane.


Inserts vs sends: an insert is a serial in-line processor that the signal passes *through*; a send is a parallel path that taps a copy of the signal into a shared return. Use inserts for tone-shaping a single lane; use sends when several lanes should share one effect (a common reverb space, a tempo-synced delay). Loom no longer privileges any effect — reverb and delay are ordinary inserts too, and they just happen to be the default residents of the Send A / Send B return chains.

The same picker everywhere. Every insert rack — per lane (including audio lanes), on each send return, and on the master — draws from one unified effect picker, whose fifteen entries read: **Filter, Dist, Reverb, Delay, Compressor, Limiter, Trem/Gate, Gate, Chorus, Flanger, Phaser, Crush, Auto-Wah, Ring** and **Width**. See [Master FX panel](#) below for each effect's parameters. Any insert's parameters are modulation and Performance-automation destinations, wherever the insert sits.

The list is not fixed in the app: each of those effects is a plugin loaded at start-up, and the picker simply shows whatever loaded. That is why a broken or missing one shows up as an entry that is absent rather than as an app that fails.

Master FX panel

The master bus has its own strip at the foot of the **scenes column** of the mixer row — a full column laid out like a lane strip, so it lines up with them: a **MASTER** label, an **EQ** section (HI / MID / LO), an **FX** button (in the lane's SEND slot — the master has no sends), a **PAN** knob, a **Mute** button (no Solo — meaningless on the master), and a fader that mirrors the master **Volume** plus a VU meter. The master EQ, pan and mute shape the whole mix; they are saved with the session and the EQ/pan are undoable like any knob. Click the **FX** button to open the **Master FX panel** below the grid (click again to close it). *(The panel's content — SENDS, MASTER COMP, INSERTS — was previously a separate "Master FX" tab; the controls are identical, only the way you open them changed.)*



SENDS — Send A and Send B

The SENDS section shows the two send buses as **return modules**. Each module is a *simple return* — a **return level**, a **mute**, and an **insert rack** — with no EQ or pan of its own. Whatever effects sit in a send's rack process everything the lanes send into it; the per-lane **A / B** knobs set how much each lane contributes. By default **Send A holds a Delay** and **Send B holds a Reverb**, but you can add, remove, reorder, or replace inserts in either rack from the same picker used everywhere else — so a send can carry a whole chain (say a filter into a delay), not just one effect.

The default reverb and delay expose the parameters below (and behave like any other insert — bypass per slot, modulatable params, etc.).

REVERB parameters:

Param	Range	Description
Wet	0–1.5	Wet output level
PreD	0–0.5 s	Pre-delay before the reverb tail starts
Size	0.05–8 s	Impulse response length (room size)
Decay	0.1–10	Tail decay shape (higher = longer tail)
Type	ROOM / HALL / PLATE / SPRING	The character of the generated space

The reverb is a convolution reverb with a procedurally generated impulse response. Size, Decay and Type rebuild the impulse in real time when adjusted.

DELAY parameters:

Param	Range	Description
Time	0.01–2 s	Delay time, used when Sync is Free
Sync	Free, 1/4, 1/8, 1/8., 1/8t, 1/16, 1/16t	Pick a musical division and the delay time locks to the project tempo, re-locking on every BPM change. Free uses the Time knob instead
Fbk	0–0.95	Feedback amount
Wet	0–1.5	Wet output level
Damp	200–12 000 Hz	Low-pass filter on the feedback loop; lower values darken repeats
Width	0–1	Stereo width of the repeats (default 1)

MASTER COMP

The master compressor sits at the tail of the master chain, after all inserts. It uses the same `CompBlock` as the per-lane strip compressor, so the parameters are identical:

Param	Range	Default	Description
Bypass	on/off	on	Pass-through when on
Threshold	–100 to 0 dB	–24 dB	Level above which compression starts
Ratio	1–20	4	Compression ratio
Attack	0–1 s	0.003 s	Gain reduction onset time
Release	0–1 s	0.25 s	Gain recovery time
Knee	0–40 dB	30 dB	Transition softness around the threshold
Makeup	~0–4 (linear)	1	Post-compression gain, up to about +12 dB

The master compressor is bypassed by default. Enable it for glue and loudness control on the final mix, or to tame transient peaks before export. See [Saving & Export](#) for how the master bus feeds the offline render.

INSERTS — the master rack

Below MASTER COMP, the INSERTS section holds the master insert chain. Add a slot from the picker and pick its type. The **same fifteen plugin types** are available in every rack — per lane, per send, and here on the master:

Filter (multifilter)




- Type: LP / HP / BP / Notch
- Freq: 20–20 000 Hz (exponential)
- Q: 0.1–24

Distortion (Dist)




- Drive: 0–1 — waveshaper saturation amount (4x oversampled)
- Mix: 0–1 — dry/wet blend


Reverb — same parameters as the Send B reverb above (Wet, PreD, Size, Decay, Type). Use as an insert to reverb the full master rather than via a send.



Delay — same parameters as the Send A delay (Time, Sync, Fbk, Wet, Damp, Width). Use as an insert for a master-bus slapback or stutter.




Compressor — the same `CompBlock` dynamics compressor as the channel-strip and master compressors, now insertable anywhere. Its own ranges are narrower than the master's:



- Thr: -60 to 0 dB (default -24)
- Ratio: 1–20 (default 4)
- Atk: 0.001–1 s (default 0.003)
- Rel: 0.001–1 s (default 0.25)
- Knee: 0–40 dB (default 30)
- Mkup: 0–4 (default 1)


There is no Bypass among them — to A/B an insert, use the slot's own bypass toggle described at the end of this section.

Limiter — a brickwall limiter (ratio 20:1, hard knee, near-zero attack) for catching peaks.




- Ceil: -30 to 0 dB — the level it will not exceed
- Rel: 0.001–0.5 s — how quickly it lets go

Tremolo (Trem/Gate) — an LFO opening and closing the volume. The oldest modulation effect there is: slow and shallow it breathes, fast and deep it chops the sound into a rhythm of its own. Lovely on pads, electric-piano parts and sustained chords. It is also the trance gate: **Square shape + a 1/16 Sync + a little Smth** is that sound. (That is a *rhythmic* gate, driven by the clock. The **Gate** insert below is the other kind — driven by the signal's own level.)




- Rate: 0.1–12 Hz — how fast it pulses, used when Sync is Free
- Depth: 0–1 — how far it closes between pulses (at 1 it cuts to silence)
- Shape: SIN / SQR / TRI / SAW — the wave it opens and closes with
- Sync: Free, 1/4, 1/8, 1/8., 1/8t, 1/16, 1/16t — lock the pulse to the project tempo instead of the Rate knob
- Smth: 0.2–50 ms — smooths the LFO, so a square shape does not click at the edges

Chorus — a delayed copy of the signal, detuned by a slow LFO and mixed back in. The two copies drift in and out of tune with each other, which the ear reads as several players at once. Thickens thin sounds and widens single-note leads without changing their pitch.



- Rate: 0.05–8 Hz — speed of the detuning wobble
- Depth: 0–1 — how far it detunes
- Mix: 0–1 — dry/wet blend (start around 0.5)


Flanger — the same idea as the chorus but with a much shorter delay and its output fed back into itself. Instead of thickening, the copies cancel each other at a comb of frequencies that sweeps up and down — the classic jet-plane whoosh. Feedback sharpens the effect from a gentle sweep to a metallic scream.



- Rate: 0.05–8 Hz — speed of the sweep
- Depth: 0–1 — how far the sweep travels
- Fbk: 0–1 — feedback; the higher, the more resonant and metallic
- Mix: 0–1 — dry/wet blend


Phaser — four all-pass filters whose corner frequencies an LFO sweeps together. They shift phase without changing volume, so mixing them back with the dry signal creates notches that slide through the sound. Similar in spirit to the flanger but smoother and more liquid — a staple on electric pianos,

funk guitar and pads.



- Rate: 0.05–8 Hz — sweep speed
- Depth: 0–1 — how far the notches travel
- Fbk: 0–1 — feedback; deepens the notches
- Mix: 0–1 — dry/wet blend


Bitcrusher (Crush) — digital degradation. Reducing the bit depth quantises the waveform to a coarse staircase, adding the gritty, harmonic distortion of early samplers and game consoles; the tone control then dulls the result the way a low sample rate would. At extreme settings it destroys the sound completely, which is often the point.



- Bits: 1–16 — amplitude resolution; 16 is nearly clean, 1 is a square-wave wreck


- Tone: 200–20 000 Hz — a lowpass for the lo-fi dullness
- Mix: 0–1 — dry/wet blend
- Dith: 0–2 — noise summed in *before* the bit reduction, scaled to the step size. It trades the gritty, signal-locked distortion of low bit depths for an even hiss, which usually sounds better on quiet material

Auto-Wah — a bandpass filter whose frequency follows how loud the signal is, so the sound opens up on every transient and closes as it decays. It is the funk-guitar wah pedal with your playing on the treadle instead of your foot; on a bassline or a clav it produces the "talking" quality no static filter gets.



- Sens: 0–1 — how strongly the level moves the filter (default 0.6)
- Range: 0–4800 ¢ — how far it travels, in cents above Base (default 2400, two octaves)
- Base: 80–2000 Hz — the resting frequency it sits at when nothing is playing (default 300)
- Atk: 1–200 ms / Rel: 10–500 ms — how quickly the follower opens and lets go. Short attack for percussive quack, long release for a slow sag
- Q: 0.5–12 — how narrow the peak is; high Q is where the vowel sound lives
- Mix: 0–1 — dry/wet blend

Gate — the opposite of the compressor: below the threshold it turns the signal *down*. Use it to clean the silence between hits — kill the bleed on a sampled drum loop, cut the tail of a reverb-soaked lane, or chop a sustained pad into stabs.




- Thr: -60 to 0 dB — the level below which it closes (default -30)
- Range: -60 to 0 dB — how far down it closes. At -60 it silences; at -12 it just ducks, which sounds far more natural
- Atk: 0.5–100 ms — how fast it opens. Too slow and it eats the transient
- Hold: 0–500 ms — how long it stays open after the level drops back, which stops it chattering on a wobbling signal
- Rel: 10–1000 ms — how gradually it closes

Ring (ring modulator) — multiplies the signal by a sine wave, which replaces every partial with a pair at the sum and difference frequencies. The result is inharmonic by construction: bell-like and metallic at high Freq, a tremolo-ish shudder at low Freq, and unrecognisably alien in between. A blunt instrument, and the fastest route out of anything that sounds too clean.



- Freq: 20–4000 Hz (exponential) — the carrier. Below ~30 Hz it reads as amplitude wobble rather than pitch
- Mix: 0–1 — dry/wet blend. Keep the dry signal in to stay musical

Width — a stereo widener: it scales the side channel against the mid, and can wobble that scale with an LFO. Below 1 it collapses toward mono, above 1 it exaggerates the sides.




- Width: 0–2 — 0 is mono, 1 is untouched, 2 is double the side signal (default 1)
- Rate: 0.05–8 Hz and Depth: 0–1 — an optional LFO on the width. Depth defaults to 0, so it is a static widener until you ask for movement
- Sync: Free, 1/1, 1/2, 1/2., 1/2t, 1/4, 1/4t — lock that LFO to the project tempo
- **Check it in mono.** Pushing Width above 1 works by exaggerating what the two channels do *not* share, and everything you add that way cancels when the mix is summed.

Slots in the chain are ordered in series: the output of each slot feeds the input of the next. Each slot has a bypass toggle so you can A/B it without removing it. Individual slots can be removed; adding the same type multiple times is allowed.

The master insert chain processes the full mixed signal, after the EQ/pan of the master strip and before the master compressor — distinct from the Send A/B returns, which receive per-lane amounts and sum back into the master independently.

XY pad

The  **XY** button in the session bar opens a floating pad — a Kaoss-style controller for playing two parameters at once with a single gesture.



It stays out of the way rather than taking over the screen: the panel floats above the interface and is **not modal**, so the knobs it drives stay visible and you can watch them move while you sweep.

Assign the axes. Below the square surface are two dropdowns, **X** and **Y**. Each lists every automatable parameter in the session, grouped by lane — the same set of destinations an LFO or ADSR can target: engine parameters, channel-strip and insert controls, anything with a knob. Pick one per axis; leave an axis on — *none* — to make the pad one-dimensional.

Play it. Drag anywhere on the square. Horizontal position sets the X parameter, vertical the Y, each mapped across that parameter's full range — bottom-left is both at minimum, top-right both at maximum. The dot follows your finger, and the assigned knobs turn with it, because the pad drives them through exactly the same path your mouse would.

The classic pairing is a filter's **cutoff on X and resonance on Y**: one diagonal sweep and you have the filter gesture that defines most electronic music. But any two continuous parameters work — delay time against feedback, distortion drive against mix, two engine parameters from different lanes at once.

Two things to know:

- Assignments are **not saved** with the session — the pad starts empty on each reload.
- Pad moves are **not captured** by Performance take recording; the automation recorder only follows knobs you turn directly.

Sidechain compression

Loom includes a sidechain ducking system. Any lane's channel strip can be ducked by the signal level of another lane (the *source*). A small envelope follower rectifies the source, chases it upwards at the **Attack** time constant and falls back at the **Release** one, and turns whatever exceeds the threshold into a gain multiplier for the target lane:

$$\text{duckGain} = 1 - \text{depth} \times (\text{env}(\text{source}) - \text{threshold}) \quad \text{clamped to } [0, 1]$$

The multiplier can never leave `[0, 1]`: at worst the lane goes silent, and when the source stops the follower returns to zero and the lane comes back. That is a guarantee, not an observation — it is why the follower is a one-pole running in its own audio worklet rather than a pair of filter nodes, which is what it was until 2026-07-27. A filter slow enough to smooth a 0.25 s release sits at 0.64 Hz, close enough to the edge of stability that single-precision rounding error accumulated instead of decaying: the "envelope" grew without bound even with the source at exactly zero, the multiplier drifted past zero into negative territory, and the ducked lane first disappeared and then came back **phase-inverted and louder** — a fault that survived pressing stop, which is why relaunching a loop could sound completely different from letting it run.

Sidechain parameters (set per lane in the lane inspector):

Param	Range	Default	Description
Source	lane selector	—	Which lane's post-mute tap drives the duck
Depth	0–1	0.6	How deep the gain dips at full envelope
Attack	s	0.005 s	How fast the duck clamps down when the source hits
Release	s	0.25 s	How fast the lane comes back up once the source falls
Threshold	dB	–40 dB	Source envelope must exceed this to duck at all

A typical use case is kick-drum ducking: set a bass or pad lane's sidechain source to the kick lane. Every kick hit momentarily ducks the bass, creating a pumping effect common in electronic music. Because the tap is taken from the source lane post-mute (but pre-duck), muting the source stops the ducking

without feedback loops.


The sidechain bus is separate from the compressor block available on each channel strip. The per-lane compressor (`CompBlock`) is a standard dynamics compressor in the signal path; the sidechain ducker is a parallel envelope-follower that modulates gain. Both can be active simultaneously.

For engine-level sound design that feeds the channel strips, see [Engines](#). For LFO and ADSR modulation of EQ, sends, pan, and other AudioParams, see [Modulation & Note FX](#).

MIDI & Samples

This chapter covers two ways to bring external material into a Loom session: importing a Standard MIDI File to populate lanes and clips automatically, and loading audio samples into the Sampler engine to build melodic instruments or drum kits.

MIDI Import



MIDI import lives under **File ► Import MIDI...**, which opens a modal dialog titled "Import MIDI". There is no MIDI button in the transport row.

Loading a file

Click the file picker inside the dialog and choose a `.mid` or `.midi` file. Loom reads the file immediately — no server, no upload. The parser extracts every track's name, General MIDI programme number, and all note-on/note-off pairs (converted to start tick, duration, MIDI note, velocity, and channel).

Loom reads **every** tempo change in the file, not just the first. The session BPM is set from the *effective starting* tempo: exporters often cram several tempo events at bar 1 (say seven at 100 followed by the real 128), and Loom collapses such a cluster to its last value rather than believing the literal first event. A file that genuinely changes tempo keeps its whole tempo map, so the transport BPM follows the song. If no tempo is present at all, the session BPM stays as it is.

Empty tracks (those with no note events) are silently skipped.

The track list

After parsing, a row appears for each non-empty track showing:

- A **checkbox** to include or exclude that track from the import.
- The **track index**, track name (cleaned of control characters), note count, pitch range, and programme number.
- A **preset dropdown**. The GM programme number is looked up against every engine's preset catalogue, and the matching presets are offered at the top of the list (e.g. programme 33 "Acoustic Bass" will surface TB-303 and Subtractive bass presets). Below a divider, every preset from every engine is available so you can override the suggestion freely.
- A ► **audition button** that plays a short three-note arpeggio through the currently selected preset without touching the session. Use this to compare options before committing.

Importing

Click **Import MIDI**. A second dialog asks what to do with the tracks, with three named buttons:

- **Add** appends the imported lanes and a new scene to the current session. The new clips are placed on the new scene row so they line up correctly with the scene's launch button.
- **Replace** clears the session and seeds it with only the imported content.
- **Cancel** aborts the import and changes nothing.

Loom creates one lane per selected track. The lane's name is taken from the matched preset (e.g. "TB Bass"), while the clip inside it keeps the original track name. If the file contained a tempo, the session BPM is updated to match. The import launches the new scene immediately.

The conversion scales MIDI ticks to Loom's internal grid (based on quarter notes divided into four 16th steps), so the notes land on the correct beats regardless of the file's tick resolution.

See [Engines](#) for what each engine sounds like, and [Sessions, Lanes, Clips & Scenes](#) for how to rearrange lanes and scenes after import.

Sampler



The Sampler is a polyphonic playback engine that maps audio across the keyboard and plays it back at the correct pitch per note. A Sampler lane can be three things — a **melodic instrument** (samples spread across ranges of keys), a **drumkit** (one sample per key), or a **loop** (a sliced loop played as a sequence of notes) — and all three share **one inspector layout**.

Presets are the instruments

The Sampler's bundled instruments — drumkits, melodic instruments, and loops — *are* its presets: pick one straight from the standard **PRESET** dropdown at the top of the inspector, which has two groups: **Melodic** — the multi-zone instruments, whether their samples are carried inline or fetched by name, merged into one shelf and sorted, because which of the two a sound is stored as is Loom's business and not yours — and **Loops**, kept apart because a chopped amen is not an instrument you play up the keyboard. **Drum kits are not in this dropdown**: they are the Drum Machine page's shelf, and a Sampler lane already loaded with one goes on playing it, named, from wherever you left it. Selecting one downloads + decodes its WAVs into IndexedDB (so it persists across reloads) and loads it. There is no separate picker. Loom ships **68 ready-made sample kits**, plus a small set of melodic instruments and loops — all working on the live deploy without a manual import. Most of the kits are grouped under **Drum Machines**: one-shot sets modelled on classic hardware (TR-808, LinnDrum, MPC60, DMX and dozens more), alongside a few hand-curated kits (**TR-808**, **Acoustic**, **Dirt**) and a 31-pad **GM Percussion** kit.

The channel layout

Whatever you load, the inspector has the same shape:

- A **keyboard map** across the top — coloured markers at each sound's key (drumkit / loop) or coloured range bands (melodic) — with **connector lines** down to the strips.
- A row of **channel strips**, one per sound. Each strip carries its **name** (the GM voice for a drumkit — KICK, SNARE...; the note for a melodic zone; **Slice N** for a loop slice), its trigger **key**, a ► **play** button to audition the sample and a ✕ to delete that sound (a row at the top of the strip, above the knobs), per-sound knobs (TUNE / CUTOFF / DECAY / LEVEL / REV / DLY plus an ► **adv** block), and **M / S** mute-solo. A + tile at the end of the row adds a sound.
- A **Selected sample** editor below — click any strip and it shows that sample's waveform on a canvas with a -/+ **zoom**, the filename, the key, and a **one-shot / loop badge** (click to toggle). Drag the handles over the waveform to set the **trim** (amber start/end lines — the dimmed regions are silenced) and, when loop is on, the **loop region** (green start/end lines that cycle while a key is held). For a melodic zone, the **root / lo / hi** keys are editable here too.

Every sound is keyed by its own note, so each pad / zone / slice has its **own** parameters and mute-solo — nothing is shared between sounds.

Building your own

Import samples... (multi-select) decodes each file, stores it in IndexedDB (so it persists across reloads), and adds it as a new melodic zone spanning the full keyboard with the root at middle C (MIDI 60); narrow each zone's root and low/high boundary in the sample editor. On a drumkit, the + tile clones the last pad onto the next free key.

Keymap and repitch

A sound's **keymap entry** has a root note and a low/high key range. When a note falls in the range, the sample plays at a rate of $2^{((\text{midi} - \text{rootNote}) / 12)}$ — equal-temperament repitching — so a melodic instrument's zones span the keyboard chromatically while a drumkit pad or a loop slice sits on a single key.

Per-sound parameters

Every sound — a drumkit pad, a melodic zone, or a loop slice — has its own set of parameters:

Parameter	Range	Default	Description
TUNE	−24 to +24 st	0	Pitch offset in semitones, applied on top of keymap repitch
CUTOFF	0–1	1	Lowpass filter cutoff (0 ≈ 60 Hz, 1 = fully open)
RES	0–1	0	Filter resonance
ATTACK	0.001–2 s	0.005 s	Amplitude envelope attack time
DECAY	0.005–4 s	0.08 s	Release tail after the gate closes
LEVEL	0–1.5	1	Pad output level
PAN	−1 to +1	0	Stereo pan position
REV	0–1	0	How much of this sound goes to Send B (a Reverb by default)
DLY	0–1	0	How much of this sound goes to Send A (a Delay by default)
LOOP	Off / On	Off	When On, the sample loops while the gate is held
LSTART	0–1	0	Loop start point as a fraction of sample duration
LEND	0–1	1	Loop end point as a fraction of sample duration
START	0–1	0	Trim in — playback start as a fraction of sample duration
END	0–1	1	Trim out — playback end as a fraction of sample duration
RETRIG	Poly / Mono	Poly	Mono cuts the previous hit on re-trigger; Poly layers them
CHOKE	— / 1 / 2 / 3 / 4	—	Sounds sharing the same non-zero group cut each other

CUTOFF, RES, LEVEL, PAN, REV and DLY are live: turn one while a pad is ringing and you hear it on the sound already playing. The rest are read when the hit fires — ATTACK and DECAY because they are envelope times, TUNE, the trim/loop points and RETRIG/CHOKE because they decide how the voice is built.


REV and DLY feed the two shared session send buses, not an insert on this lane — see [Send A and Send B](#) for what lives in each. CHOKE is the same mutual-exclusion mechanism the Drum Machine uses: on a drumkit the GM closed and open hi-hats start in group 1, so a closed hat silences the open hat's ring.

These live in each channel strip — drumkit, melodic, and loop all share the same per-channel rack (there is no eight-pad limit and no separate knob row).

Drumkits

A drumkit is one sample per key, edited on the **drum-grid** (the same grid as the Drum Machine engine). A kit holds **any number of sounds** — not just eight — and you grow or shrink it with the + tile and each strip's X. Loom ships **68** ready-made sample kits — the **Drum Machines** group (classic-hardware one-shot sets), three hand-curated kits (**TR-808**, **Acoustic**, **Dirt**) and a 31-pad **GM Percussion**; pick one from the PRESET dropdown and the lane is ready to play. The kit is rebuilt fresh from its manifest on each session load, so the WAVs never need re-importing.

Loops



A loop is sliced into segments mapped to **consecutive notes**, so playing the keys in order replays the groove. Above the strips the **whole loop** is drawn as one continuous, colour-coded waveform with a cut line at every slice. Selecting a loop also **drops a note clip** onto the lane — one note per slice, each placed at its exact position so the notes form a continuous **staircase** — opened in the piano-roll; hit Play and it replays the loop, now as discrete notes you can move, mute, repitch or re-order. The clip keeps the loop's waveform as its header.

See [Editing Clips](#) for drawing patterns in the drum grid, and [Engines](#) for the Drum Machine engine (which lists every kit — synth and sample — in a unified preset table).

Audio channel

The **audio channel** is the first-class way to bring a finished loop into a Loom session: drop a WAV and it plays **tempo-locked to the project without changing pitch**, with its waveform shown as a header above the clip editor. It stays a pure audio loop; to chop a loop into individually editable note slices, load it through the Sampler's **Loop** family instead (see [Sampler](#)).



Creating an audio channel

Adding one takes two steps: make the lane, then give it a WAV.


1. **Make the lane.** Click the + at the top of the clip grid. It opens a menu listing every engine plus a final **Audio channel** entry — pick that and an empty audio lane appears.
2. **Give it a WAV.** Click any empty cell in that lane (or right-click it and choose *Import audio (WAV)*...) to open a file picker, or simply drag a WAV straight onto the cell. Loom decodes the file, stores it in IndexedDB (so it survives a reload), estimates its original tempo, and places the loop as an

audio clip in that cell. The clip opens in the inspector.

An audio lane holds **one WAV per clip**, so you can fill several of its cells with different loops and launch them like any other clip.

Each new audio lane gets a launch button on its scene row, so it is immediately playable alongside the rest of the session.

The audio-clip editor



An audio clip has no note grid. Clicking it opens the **audio-clip editor** — a **Warp ON / OFF** toggle (tempo-locking, see below) above a **waveform header**.

The waveform header shows a peak view of the buffer with a bar/beat ruler, any detected slice markers (orange), and a live playhead while the clip plays. This same header also appears **above** the normal piano-roll or drum-grid for any clip that references a buffer, so you always see the audio you are editing against.

When the transcription helper is configured, the audio-clip header also shows a **TRANSCRIBE** row with a **melodic/drums** toggle and a **Transcribe loop** button. This sends the clip's effective loop region to the local `/transcribe` pipeline and creates a fresh lane with an editable note/drum clip from that region.

Tempo-lock (Warp)

With **Warp ON** (the default), the audio channel plays in time with the project BPM using a pitch-preserving WSOLA time-stretch:

- **At the loop's native BPM** the stretch ratio is ≈ 1 , so playback is essentially identical to the source file.
- **At any other project BPM** the buffer is time-stretched to fit — faster or slower — **without changing pitch**. The stretched buffer is cached and re-rendered automatically whenever you change the project tempo, so the loop stays locked as you experiment.


With **Warp OFF** the loop plays at its natural speed with no tempo sync — useful when you want the audio exactly as recorded.

First-play note: on the very first loop iteration after import (before the stretch cache is warm) playback briefly falls back to a varispeed render — a slight pitch shift that self-heals from the next iteration. At the loop's native tempo the ratio is ≈ 1 , so even that first pass is near-identical.

Slicing a loop into notes

The audio channel itself is a pure WAV loop. To chop a loop into individually editable hits, load it through the Sampler's **Loop** family (see [Sampler](#)) rather than the audio channel. Loom detects slice points (from embedded **Acid** / **cue** / **AIFF** markers when present, or by onset detection plus a tempo estimate), stores one short sample per slice in IndexedDB, and creates a **note clip** that triggers the slices in order on a piano-roll — so the groove plays back identically, now as discrete, editable notes. Move, mute, repitch, or re-order the hits in the piano-roll, and tweak each slice's tune/cutoff/decay/level/pan in its channel strip (see [Per-sound parameters](#)); the clip keeps the original waveform as its header. (Earlier builds did this from a **& Slice** → **pads** button on the audio channel; that moved to the Sampler's Loop family.)

Editable warp markers



Amber warp markers appear as vertical lines in the waveform header; drag one to re-anchor that beat to the grid.

An audio clip can carry **warp markers** — sparse reference points that pin specific audio beats to exact positions on the session grid. With markers in place, the WSOLA time-stretcher locks each inter-marker region independently, so the groove tracks the grid precisely rather than relying on a single global stretch ratio.

Auto-seeding on stem import. When you separate a track into stems (see [Stem separation](#)), Loom analyses the drums stem (the most rhythmically reliable) and seeds a warp marker every **four bars** across all stems of that import — a sparse starting point you then refine by hand, not a marker on every beat. The drums stem becomes the **warp reference** for the group; its clip is the one with a fully editable marker overlay in the waveform header.

Editing markers. Open the drums-stem clip in the inspector. In the waveform header you will see amber vertical lines at each marker. You can:

- **Drag** a marker left or right to realign that beat to the correct grid position.
- **Double-click** an empty area of the header to add a new marker at that audio position.


- **Double-click** an existing marker to delete it.

Edits to the reference clip's markers propagate automatically to the other stems that share the same warp group — move one marker and all four audio clips update together.

Warp pill. An amber **Warp** pill in the clip-editor toolbar toggles warping on and off for the clip. Warping is pitch-preserving (WSOLA) and re-renders automatically whenever the project tempo changes, so stems stay locked to the grid as you experiment with BPM.

Stem separation (optional, local service)

Stem separation lets you drop a finished song into Loom and get it back as four separate **audio lanes** — **Vocals**, **Drums**, **Bass**, and **Other** — so you can mute, solo, and remix each part inside the existing session.



How it works

Click the  button in the session bar (tooltip "Separate a song into stems (local service)"), or use **File** ► **Separate into Stems....** A dialog titled "Separate into stems" opens and immediately checks whether the local helper service is reachable:

- **Service found** — the hint line reads "4 tracks (Vocals / Drums / Bass / Other) via the local service." and the **Separate** button becomes active once you pick a file.
- **Service not found** — the hint reads "Can't find the stems service at localhost:8765. Is it running?" and Separate stays disabled. Start the service (see below), then re-open the dialog.

There are two ways to feed it material:

- **Pick an audio file** with the file picker, then click **Separate**.
-  **Capture PC audio** — records whatever your computer is playing (you must tick "share audio" in the browser's capture prompt). The button becomes **Stop and separate**; press it when you have enough and the recording goes straight into the same separation flow, no file needed.

A **Replace the session with the stems** checkbox sits above the buttons and is **on** by default; see below for what each setting leaves you with.

Either way, the dialog then shows a progress bar:

1. **"Uploading..."** — the file is being uploaded to the local service.
2. **"Separating... m:ss"** — the service is running Demucs; the counter shows elapsed time. The bar may be indeterminate if the model does not report fine-grained progress.
3. **"Transcribing notes..."** — only when you asked for transcription (see below).
4. On success the dialog closes automatically and four new **audio lanes** appear in the session — one per stem (**Vocals, Drums, Bass, Other**). Each lane holds a full-length audio clip downbeat-anchored to bar 1 and tempo-warped to the session BPM. Hitting Play reconstructs the original mix; mute or solo any lane to isolate parts.

With **Replace the session with the stems** ticked you are left with a clean session holding just the four stems in a single scene named **Stems** (the project name, key, sends and master rack are kept). Untick it and the stem lanes are added to the session you already had.

With **Replace** ticked, the session BPM is conformed to the **detected floating-point tempo** of the imported track, so the warp ratio starts at exactly 1.0 and the grooves lock to the grid from the first bar. When you add stems to an existing session instead, the project tempo is authoritative and is left alone — the stems are warped to it.

The entire lane-creation is a **single undo step**, so you can undo all four lanes at once.

Transcribe to notes (opt-in). A **"Transcribe to notes (experimental)"** checkbox in the Separate dialog (off by default) tells the service to additionally analyse each stem and create a note clip alongside the audio:

- Pitched stems (Vocals, Bass, Other) → melodic note clips, with pitches estimated via pYIN onset detection.
- The Drums stem → a drum-machine pattern, with hits mapped to the Drum Machine engine.

Each transcription lands as **its own new lane**, named after the stem it came from (*Notes: Drums, Notes: Bass, ...*), in a **separate scene called "Transcription"** — kept apart from the audio stems so you can A/B the two against each other. The clips are immediately editable in the piano-roll or drum-grid. They are a convenience starting point, not a perfect transcription — pitch detection works best on monophonic or sparse material.

Cancel aborts a running job and frees the temporary files on the service. **Close** closes the dialog (only available when no job is running).

Opt-in nature

The feature is entirely opt-in. If you never start the service, nothing else in Loom changes — the ≡ button (and its File-menu twin) is the only touch point, and it degrades gracefully to a clear "service not found" message.

Stems land in IndexedDB as ordinary sample assets: they survive browser reloads just like any other sample you import.

Setting up the local service

The separation runs on your machine via a small Python service in `tools/stem-service/`. It requires **Python 3.10+** and **ffmpeg** on your PATH.

```
cd tools/stem-service
python -m venv .venv
# macOS / Linux:
. .venv/bin/activate
# Windows:
.venv\Scripts\activate

pip install -r requirements.txt
uvicorn app:app --port 8765
```

The first time you separate a track the service downloads the Demucs `htdemucs` model automatically (several hundred MB). Subsequent runs skip the download. Separation takes roughly **1–2 minutes per song** on CPU; a GPU-enabled PyTorch build is much faster.

Codespaces and custom service URL

If you want to run the service in a GitHub Codespace (a Linux VM with Python), start it there with the same commands above, forward port 8765, and paste the resulting HTTPS URL into the browser console:

```
localStorage.loomStemServiceUrl = 'https://your-codespace-url-8765.preview.app.github.dev'
```

The same override works for any non-default host. CORS for `localhost:5173` (dev), `localhost:4173` (preview), and the GitHub Pages origin is already configured in the service.

Note: Chrome's *Private Network Access* policy may add a preflight request when the Pages version of Loom calls `http://localhost:8765`. The lowest-friction setup is running Loom locally (`npm run dev`) alongside the service.

For full notes on CORS, the HTTP contract, and the Codespaces workflow, see [tools/stem-service/README.md](#).

Saving & Export

Loom keeps everything in the browser. There is no account, no cloud, and no server upload — your sessions live in your browser's `localStorage` and on your own filesystem when you export them. This chapter covers how to save and restore sessions, undo your work, and render a scene to a WAV file.

Sessions in the browser

Every parameter, every lane, every clip, and every scene is part of the current session. Changes take effect immediately; nothing is auto-committed to disk. When you close the tab, the browser preserves the last-saved state via an autosave entry in `localStorage`, so reopening Loom typically drops you back where you left off.


Three buttons on the session bar (the second header row) drive session management:

Button	What it does
 New	Discards the current session and starts a blank, empty one (tooltip "New empty session").
Save	Opens the Save Manager with the name field ready to type.
Load	Opens the Save Manager to browse and restore a saved session.

See [Transport](#) for the full two-row header layout (transport & tempo on top, session & I/O below).

New and Load both start from an empty desk. Whatever replaces the session — New, loading a save or the autosave, picking a demo, or a MIDI import in Replace mode — first releases everything the old session was holding: every lane's channel strip (level, pan, EQ, both sends, mute, compressor and sidechain), the master rack, the send buses, the master processors, mute/solo and the per-lane note-FX chains. Nothing rides through. It is worth stating because the opposite used to be true: a lane whose id existed in both sessions kept its live strip, so a demo you loaded could arrive already panned and EQ'd by the song before it, and pressing New did not put the mixer back to flat. The single deliberate survivor is the **master volume fader** — that is how loud your speakers are, not part of the song.

Save Manager



Clicking **Save** or **Load** opens the Save Manager modal. It is the single place where all session persistence happens.

Saving a session

Type a name in the text field at the top and click **Save current** (or press Enter). Loom writes the full session state — BPM, time signature, every lane's engine + inserts + clips + scenes, the mixer state, and the arrangement take if one exists — into `localStorage` under a unique key, and also updates the autosave slot. The entry appears in the list immediately.

The save format is versioned (`schemaVersion: 3`). Only version-3 saves load: anything else is refused outright with an **"Invalid save data"** alert rather than loading broken state. There is no conversion from older formats. What *is* handled is a version-3 save written before some field existed — those are backfilled on load, so a session saved months ago still opens with its insert racks and preset dropdowns intact.

The saved-session list

Each row in the list shows the session name, date/time, and size in KB. Per entry you can:

- **Load** — restores the session and closes the modal.
- ⬇ (download) — exports the entry as a `.json` file to your filesystem without closing the modal.
- ✎ (rename) — prompts for a new name in place.
- 🗑 (delete) — confirms and removes the entry.

The topmost row is **Auto-save (latest)**, which always reflects the state at the time of the last named save.

Load from file...

Imports a `.json` file you previously downloaded (or received from someone else). Loom validates the schema before applying it; an invalid file shows an alert.

Clear all saves

Removes every named entry from `localStorage`. The autosave slot is preserved. A confirmation dialog appears before anything is deleted.

Storage readout

The footer shows the total size of all named saves, so you can keep an eye on `localStorage` usage.

Undo / redo

Loom keeps a global undo history that covers **every** session mutation: adding or removing lanes, clips and scenes, editing notes in the piano roll or drum grid, moving or tempo-scaling clips, changing engine parameters via knobs and faders, and inline renames. You can step backwards and forwards from the header buttons or from the keyboard.

Header buttons. A pair of ↶ **Undo** / ↷ **Redo** buttons sits in the transport bar, just to the right of the Play / Stop controls. Each reflects the current state of the history — it is disabled (greyed out) when there is nothing to undo or redo — so you can see at a glance whether a step is available. Click ↶ to undo the last change, ↷ to redo it.

Keyboard shortcuts.

Action	Shortcut
Undo	Ctrl+Z / Cmd+Z
Redo	Ctrl+Shift+Z / Cmd+Shift+Z or Ctrl+Y

Capture is automatic — Loom snapshots the session after each interaction and coalesces a continuous gesture (a knob or fader drag, a note drag, a marquee move) into a **single** undo step, so one drag undoes in one click. The shortcuts are inactive while a text input has focus, so typing a save name — or an inline rename — never accidentally undoes your work. Loading a session, loading a demo, or starting a new session clears the undo history.

WAV export



WAV export is part of the unified **REC** group on the session bar (second header row), not a separate button. The ● **REC** button (tooltip "Record — pick the mode beside it") records using whichever mode is selected in the adjacent **mode selector** (`#rec-mode`):

- 📹 **take** (default) — records knob moves and clip launches into a performance take (not a WAV).

- 🕒 **live** — records the live master output in real time to a stereo 16-bit WAV file.
- ⚡ **offline** — renders the current scene to a WAV file offline (faster than real time).

To export audio, pick 🕒 **live** or ⚡ **offline** and press • **REC**. The two WAV backends behave as described below.

🕒 **live (real-time WAV)**

The real-time backend is the ground-truth render, and it works like a tape machine: **arm** → **play** → **stop**.

Pressing • **REC** in this mode does not start anything — it only **arms** the recorder, and the button reads • ARMED . Capture begins on the next downbeat, whichever way you start the transport (the ► button, a clip launch, a scene launch), and the button changes to • Recording... . It then records **open-ended** until you press Stop.

That is the important difference from the offline render: it is not one pass of one scene. It captures **the whole performance, including scene changes** — launch scene A, let it run, switch to B, drop a clip, then stop, and all of it is in the file. It taps the live master output after every insert, the master compression and the master FX, so what you hear is exactly what lands in the file, including any random variation from a voice that uses it.

A 2-second tail is appended after you stop so reverb and delay repeats are not cut off.

⚡ **offline (fast WAV render)**

The offline backend rebuilds the full audio graph — lanes, inserts, master bus — inside an `OfflineAudioContext` , applies every lane's current sound state, batch-schedules all note events, and renders faster than real time without touching the live session. It shares the same encoder and download step as the real-time path, so the output format is identical.

It loops seamlessly. The offline renderer deliberately renders **two cycles of the scene and gives you the second one**. The first cycle is thrown away because it starts from silence — no reverb tail from the previous bar, no delay repeats in flight. The second cycle inherits all of that, so the file loops without a seam. This is the reason to prefer the offline render for loop material.

For the same reason it appends **no** FX tail: the render is exactly the musical, bar-aligned length. A trailing tail would round up to an extra bar and the loop would drift.

Where it differs from the live sound. The divergence is structural, not random: per-pad FX sends and per-voice drum-strip sends/EQ are dropped offline, and sample-mode drum kits render through the sampler path. Both are approximations of the live per-voice mix, so a kit that leans on per-voice sends will not sound identical. Everything else — presets, clip automation, the worklet engines — is applied exactly as live. That includes automation drawn on the mixer column: a level fade or a pan sweep is scheduled point by point, so it renders as the move you drew rather than jumping to its end value.

What gets exported

- **One pass of the scene (offline only).** The offline export captures one full iteration of the longest clip across all sounding lanes; shorter clips loop to fill that window. The **live** take has no such window — it runs until you stop it.
- **A scene must be playing (offline only).** Press • **REC** in ⚡ offline mode with nothing playing and Loom shows a brief notice and does nothing; launch a scene first. In ⌚ live mode arming with nothing playing is normal — that is the point, since it starts capturing from the downbeat of whatever you launch next.
- **Switching mode disarms.** Changing between take / live / offline clears whatever the previous mode had armed.

Where the recording goes

When a render finishes — live or offline — Loom asks what to do with it:

- **Download a WAV**, named `loom-take-<timestamp>.wav`; or
- **Insert it back into the session** as a new audio channel, tempo-locked to the project BPM. This is how you bounce a busy scene down to a single audio lane and keep building on top of it.

Cancelling the dialog discards the take and writes nothing.


Live build and GitHub Pages

The public instance of Loom is deployed automatically to <https://ijol.github.io/Loom/> — every push to `main` triggers a GitHub Actions workflow that runs `npm run build:pages` and deploys the result to GitHub Pages. That script builds the plugins into `public/plugins/` before bundling, which matters more than it sounds: the engines and the insert effects are plugins loaded at runtime, so a build that skipped that step would deploy an app with no instruments in it. The standard `npm run build` does the same at base `/`, for local development or self-hosting on another path.

Performance & Arrangement

Loom has two main views: **Session** and **Performance**. The Session view is the clip grid you work in most of the time — lanes, clips, and scenes launched on demand. The Performance view is a linear **arrangement timeline**: a fixed song laid out from left to right with a playhead that moves forward through it, starts at the beginning, and stops at the end.

Switch between the two views using the **Session / Performance** toggle in the transport bar (`#mode-toggle`). Switching stops playback. See [Transport](#) for the full transport layout.



Five ways to fill the arrangement

The arrangement starts empty — and an empty arrangement invites its primary gesture: **drop audio loops on it**.

1. Drop audio loops (the primary gesture)

Drag audio files (WAV, MP3, FLAC, OGG, M4A, AIFF) from your file manager anywhere onto the Performance view — the big **"Drop audio loops to start"** invitation when the timeline is empty, or the **+ new lane** strip at the bottom once it has content. Each file becomes:

- a new **Audio lane**, named after the file;
- a clip **fitted to the session tempo**: the file is assumed to be a loop, its duration rounds to the nearest whole-bar count and the clip warps to fit exactly (pure arithmetic — no beat detection). The band's **bars-chip** shows the fit; if the rounding guessed wrong, the clip's own Warp editor corrects it;
- a **band** on the timeline at the bar you dropped it, showing the loop's **waveform** with tick marks where it repeats.

Dropping several files at once creates one lane per file, all starting at the drop bar. Nothing about this is a parallel world: the dropped loop is an **ordinary session clip on an ordinary Audio lane** — it appears in the Session grid too, launches from there like any clip, and saves in the session file with its sample in the project's sample store.

2. Record a loop (●)

The other way in is to **record** a loop instead of dropping one. The ● button — in the Performance toolbar, and on the big invitation when the timeline is empty — is a bar-quantized looper:

1. Press ●. Recording is armed for the **next bar boundary** — the button blinks and the toolbar shows *recording at bar N...* If the transport is stopped, it starts playing first.
2. When the bar arrives, recording begins. The button turns solid red with a live bar count, and a red ghost band grows on the + **new lane** strip where the take will land.
3. Press ● again. Recording continues to the **end of the current bar** and stops there — the cut is sample-exact, so the loop is always a whole number of bars with no tail.
4. The take lands exactly like a dropped file: a new **Audio lane**, a clip at the session tempo, and a band at the bar where recording **started**. Undo removes it like any drop.

What it records is chosen by the **source selector** next to ●:

- **Master** (the default) — resamples Loom's own output: everything you hear becomes the loop. No permissions needed.
- **System** — records the computer's audio via the browser's screen-share picker (tick **share audio** — a browser tab, a different app, anything).
- **Mic** — records the microphone, unprocessed (no echo cancellation or noise suppression — it is music, not a call).

For the external sources a 🎧 **monitor** toggle appears; it is **off by default** — turn it on to hear what you are recording through Loom (mind the feedback loop with a live microphone on speakers).

Escape hatches: **Esc** cancels a countdown or a recording outright; pressing ■ **Stop** while recording keeps the whole bars already played and drops the partial one (if no bar has completed, the take is discarded); switching views cancels a countdown but lets a running recording finish and land.

3. Copy to Performance

The fastest route from a working session to a playable song is the † **Copy-to-Performance** button (`#copy-to-performance`) in the session bar of the header — now an icon-only button with the tooltip "Copy the scenes to the Performance timeline". Clicking it calls `arrangementFromSession`, which walks your scenes in order and lays them out as a linear song:

- Each scene becomes one section.
- The section length equals the longest effective clip in that scene (measured in bars). If a clip has a loop sub-region enabled, its sub-region length is used instead of the full clip length.
- Every lane that has a clip assigned for that scene gets a timeline band covering the section.

After the layout is computed, Loom switches you to Performance automatically. **This one is not undoable** — it overwrites whatever arrangement was there, so copy before you start editing bands by hand, not after.

4. Record a take live

You can record the arrangement in real time while you play.

Recording is driven by the unified **REC** control in the session bar, which has three modes selectable beside the **REC** button:  **take** (the default — captures knob moves + clip launches into a performance take, described below),  **live** (records real-time audio to a WAV file), and  **offline** (renders the current scene to WAV offline, faster than real time). The steps below assume  **take** is selected.

1. Make sure the REC mode selector (`#rec-mode` , next to the REC button) is set to  **take** — its default. (The other two modes,  **live** and  **offline**, record audio to WAV instead; see [Saving & Export](#).) Then click **REC** (`#rec`) in the session bar to arm recording.
2. Stay in Session view and press Play. Recording begins.
3. Launch clips and scenes as you would for a performance. Move any knobs whose automation you want captured.
4. Press Stop. Loom finalises the take: any still-open clip events are clamped to the stop time, durations are computed, and the recorded content appears in the Performance view as timeline bands and automation curves.

If you arm REC and then switch to Performance mode before pressing Play, the arm is cleared automatically (a toast notification appears) because Performance mode drives playback from the arrangement directly rather than from the live session.

5. MIDI import

When you import a Standard MIDI File via **File** ► **Import MIDI...** (see [MIDI & Samples](#)), Loom calls the same `arrangementFromSession` logic after building the session. Because an imported MIDI file produces a single scene whose clips span the full song, the arrangement comes out as one long section per lane — the complete track laid out linearly from bar 1.

The timeline

Once the arrangement has content, the Performance view shows **one scrolling surface**: the ruler stays pinned to the top, the lane labels stay pinned to the left, and everything scrolls together (the rows used to scroll independently and desync on a long song). Zoom and horizontal scroll are remembered **per machine**, not in the save file.

- **Toolbar** — Length (bars), a Zoom slider (16–400 pixels per bar; Ctrl+wheel works anywhere on the timeline), the **Loop A–B** toggle with numeric **A** and **B** bar fields beside it, and a readout showing total bars and BPM.
- **Ruler** — a bar-numbered ruler across the top. **Click it to move the playhead** (playing or stopped); **drag on empty ruler space to set the A–B loop**. When the loop is active, its brace with two drag handles sits on the ruler.

- **Clip bands** — one row per lane. An audio band shows its **waveform** (with tick marks where the loop repeats, and a bars-chip naming the fit); a MIDI band shows a **mini note preview**. A **muted** band paints dimmed; a **selected** one carries a blue outline.
- **+ new lane strip** — the permanent drop target at the bottom: drop an audio loop, get a lane.
- **Automation curves** — below each clip band, any recorded or drawn automation curves appear. You can draw into curves directly using the Line or Flat brush.
- **Master automation** — curves routed to global (master) parameters, in a section at the bottom.
- **Playhead** — a vertical line that moves in real time while the arrangement plays, scrolling with the content.

Editing bands

Bands are first-class objects with a selection model — you do not have to re-record to change the layout:

- **Select** — click a band; **Shift-click** adds or removes it from the selection; drag on empty track space draws a **marquee**; clicking empty space clears.
- **Move** — drag the body. The default is **free movement**: the band lands where you drop it (snapped to the beat), gaps are allowed, and a collision clamps it against its neighbour. Hold **Shift** while dragging for the old **ripple** behaviour (later bands push forward); hold **Alt** to disable the beat snap. Dragging **vertically** onto another lane row moves the band to that lane. **Escape** cancels a drag in flight.
- **Resize** — drag the handle at either end. The **left edge is an honest trim**: it slides the band's content offset with it, so what you hear stays anchored to the bars — trimming reveals or covers material instead of shifting it. Resizing clamps against neighbours; it never moves another band.
- **Delete** — press **Delete** (or click the band's **x**).
- **Keyboard** — **Ctrl+D** duplicates the selection, **Ctrl+C / Ctrl+V** copy and paste it at the playhead (relative offsets preserved, on the bands' own lanes).
- **Right-click** a band for the context menu: **Mute** (the band stays but never fires — the gate is in the scheduler, so a muted band is genuinely silent), **Split at playhead** (two bands; the right half keeps playing the same material via its content offset), **Duplicate**, **Delete**.

All of it is undoable on Performance's own stack (see below).

A–B loop brace

The **Loop A–B** button in the Performance toolbar toggles the arrangement-wide loop brace. When active:

- Two drag handles on the ruler mark the loop start (A) and loop end (B). Drag either handle to set the window; handles snap to whole bars. You can also **type the bars** into the **A** and **B** fields next to the Loop A–B button, which is far easier on a long song — typing a value also switches the loop on.

- The playhead wraps within [A, B). When it reaches B, all lanes are stopped and playback re-anchors to A instantly. Any clip that was already active spanning A is relaunched at the wrap point so there is no gap.
- When Loop A–B is off, Loom plays in **song mode**: the arrangement runs from the beginning to `durationSec` and stops — every lane is halted and the transport stops.

This loop brace operates on the arrangement timeline as a whole, and it is **the same region** as the clip editor's **Global** loop: set A–B here and the active scene's shared loop follows, and when you switch back into Performance the brace picks the scene's shared loop up — but only if that loop is switched **on**. A scene whose shared loop is off leaves the A–B window you had here untouched rather than clearing it. It is still distinct from a clip's *own* (non-Global) loop brace, which repeats a sub-region inside that one clip. See [Editing Clips](#) for both.

Song playback

Press Play while in Performance mode to start playback from the beginning of the arrangement (or click the ruler to start anywhere). The arrangement's own play state is used — it is independent of the live Session runtime. The playhead advances, clips launch at their scheduled times, and automation curves are applied continuously. A band that was left-trimmed enters its clip **already started** — the music under the playhead is the music you placed there.

The timeline wins over WEAVE, with the same rule the Session grid keeps: the moment the arrangement drives a lane, that lane's weave (or follow) is suspended, exactly as if a scene had launched it. The WEAVE panel takes the lane back with its own ►.

At the end of the arrangement (when Loop A–B is off), all lanes stop and `onArrangementEnd` fires, which stops the transport. You can press Play again to restart from the top.

Knob automation written into the arrangement's curves is applied every lookahead tick alongside clip launches. Automation values are normalised (0–1) internally and mapped to each parameter's min–max range at playback time.

Automation lanes

Automation curves can be added by hand — you do not have to record a take first. Set a non-zero **Length** in the toolbar (or record/copy content so the arrangement has a duration), then use the automation header that appears just below the ruler.

Adding a lane

The header row contains a grouped parameter dropdown and a + **Automation** button. The dropdown lists every automatable parameter in the project, organised by prefix (lane ID or `master`). Each entry shows the parameter ID and its label — for example `lane-1.fx.reverb.wet – WET`. Each lane's **mixer**

column is in there too: `bus.level`, `bus.pan`, `bus.delaySend`, `bus.reverbSend`, the three `bus.eq.*` bands — and, under the MIXER heading, `mixer.mute` and `mixer.solo` as 0/1 stepped curves — so a fade-in, a pan sweep, a send that opens up over eight bars or a section that drops the drums is drawn here like any other curve. Select the parameter you want, then click **+ Automation**. A new lane appears below the clip band for that lane (or in the Master section for global parameters). The curve starts flat at the parameter's current value.

There is a quicker route for a knob you can see: **right-click it**. In Performance view the menu offers *Automate on the timeline* (or *Edit automation on the timeline* if a curve already exists) and jumps you straight to it. The same menu in Session view targets the open clip instead — see [Modulation & Note FX](#).

Drawing the shape

Two brush buttons sit in that same automation header, after the **+ Automation** button and under the label **Brush**. They only appear once at least one curve exists — per-lane or on master — so on a fresh arrangement there is no Brush label to find:

- **Line** — click and drag to draw a ramp between the start and end points of the gesture. Use this for smooth fades, filter sweeps, or any gradual change.
- **Flat** — paints a constant value across the drag range. Use this for step-style automation or to hold a value steady across a section.

The active brush is highlighted, and the choice applies to every curve. Each lane also exposes **On / Off** and **Smooth / Stepped** toggles in its header. **On / Off** mutes the curve without deleting it. **Stepped** switches interpolation from smooth linear to staircase, snapping the value at each sub-step boundary — useful for parameter jumps that should be instantaneous.

Removing a lane

Click the **×** button on the right side of the lane header to remove the curve entirely. The action is undoable with Ctrl+Z / Cmd+Z.

How automation curves play back

Curves added manually behave identically to curves captured by recording. During arrangement playback they are applied continuously alongside clip launches — knob values are updated every lookahead tick, mapped from the normalised 0–1 curve to the parameter's min–max range. Curves generated by a recorded take (see [Transport — REC](#)) and manually drawn curves live in the same list and can coexist on the same lane.

For modulator-driven per-lane automation (LFO / ADSR) that runs in Session view rather than the arrangement timeline, see [Modulation & Note FX](#).

Lane mute & solo — two kinds, side by side

The Performance lane header carries **two pairs** of buttons, and they mean different things:

- **S ▶ / M ▶ (launch-solo / launch-mute, accent-coloured)** — the **arrangement** stops driving lanes. Launch-solo a lane and the take launches clips and applies automation **ONLY** there; what was already sounding on the other lanes finishes **at the next bar** (a musical handover, not a cut). Un-solo and the freed lanes re-enter the band under the playhead immediately. This is "solo this loop" in the musical sense, and it is fully reversible mid-song.
- **m / s (the mixer's audio mute/solo)** — the classic desk pair: silences the lane's audio **bus** instantly, tails included. The lane keeps scheduling underneath. These are the **SAME** switches as the mixer's and the clip editor's — one pair, three places.
- **VU meter** — a vertical level meter shows the lane's live output RMS.

The audio pair is saved with the session; the launch pair is performance state.

The mixer's mute and solo are also automatable: every lane lists `mixer.mute` and `mixer.solo` (under the MIXER heading) in the + Automation picker, as 0/1 stepped curves — and pressing m/s **while a take records** writes the press into the take, so a mute performance replays like any knob move. See [Mixing & FX](#).

Persistence

The arrangement is saved as part of the session file (schema version 3). Recorded takes — clip bands, automation curves, loop brace position — survive Save / Load and browser restarts.

Performance has its own, separate undo stack. This surprises people, so it is worth stating plainly: while you are in Performance mode, `Ctrl+Z` / `Ctrl+Y` undo **timeline edits only** — moving, resizing or deleting bands, length and zoom changes, adding or removing curves, drawing automation. Those keystrokes never reach session undo, even when the timeline stack is empty and nothing happens. Conversely, session edits are invisible to Performance's undo. Switch back to Session view to undo session changes.

Two things are outside undo entirely: the **raw recording** (a take is finalised on Stop and does not enter the stack) and **Copy to Performance** (it overwrites the arrangement outright).

See [Saving & Export](#) for how sessions are saved and loaded.

WEAVE

WEAVE is a third view, beside Session and Performance. Where the Session grid launches clips you wrote, WEAVE **crossfades between loops** — and keeps crossfading, on its own, for as long as you let it. A scene made in WEAVE is not a fixed arrangement; it is one that keeps finding new material while you steer it.

Open it from the view toggle in the transport bar. WEAVE is a **panel plugin**: it ships in `plugins/weave/` and is loaded at runtime like any engine or insert, which means it can be absent from a build without the rest of the app noticing.



The idea in one paragraph

Every track in WEAVE plays a **blend of two or more loops** instead of its own clip. Where the blend sits is a position you can move by hand or leave to the master flow. The two loops never simply alternate: what they **SHARE** stays put — that shared skeleton is what holds the bar together — and what differs

hands over gradually. Melodic material is blended in **scale degrees**, so a crossfade cannot detune; it can only move between notes that are in the key.

You are never editing the loops. WEAVE is a tool that sits ON TOP of the session: switch a track's weave off and its clip is exactly as you left it.

The track row

Each track is one row of up to four lines. The top line is what a hand reaches for while the music runs; the second is what you set once and leave; the third appears only while the track is generating; the fourth is the list of loops it may draw from.

Top line, left to right:

- **VU meter** — what the track is putting out, the same meter the mixer column shows. It keeps reading with the transport stopped, because a held note or a reverb tail is still sound.
- **Loop ring** — where this track is in its bar, and how long is left of a queued change.
- **Name**, then ► ■ **M S** 🔒 — launch, stop, mute, solo, and **lock**. A locked track holds the loops it has while everything else travels. Mute and solo are the mixer's own, not a second pair.
- **Level** — the same gain as the mixer column, with its percentage beside it.
- **Topology** — how this track weaves. See below. The blank entry, **– off –**, returns the track to simply playing its clip.
- **The weaving control itself** — a dial, a queue or a square, depending on the topology.
- **🔊 Sound** — off by default. See *Playing one part on several instruments*.

Second line: the note bar this track is about to play, then which slot (on a rack), instrument, preset, part, **who it follows**, **GEN**, **how long it takes to repeat itself**, style shelf, half/double time, and octave. The three in bold have sections of their own below.

Third line: the generator's controls, and only while **GEN** is lit. On every other track the line is there and empty, so switching a track on does not shove the tracks below it down the screen.

Fourth line — PLAYS: the loops this track draws from. See *The list a track draws from*.

Topologies — how a track weaves

Pick one per track from the dropdown.

A→B

Two loops and a dial between them. When the dial reaches the far end, **B becomes the new A and a fresh loop is drawn for B** — so the journey never ends and never returns to where it began. Where that fresh loop comes from is decided by the track's **PLAYS** list, if it has one; see below.

The control is an **endless dial**, not a fader, and it is dragged **up and down**. A dial rather than a fader because a fader promises two ends, and with the scene evolving there are none.

The names sit either side as *from* → *to*.

Cloud

Four loops, one in each corner of a square, and a dot you drag. An edge behaves like a plain two-loop crossfade; the middle is the one place all four sound at once.

A cloud walks its lap along a **path**, because the master flow is one number and a square is two:

- **RIM** — the four sides. The dot spends its whole lap on an edge, so every corner arrives as a clean handover.
- **CROSS** — side, diagonal, side, diagonal. The same four corners, but it crosses the middle twice a lap, alternating clean handovers with everything piling up.

A cloud evolves too: each lap swaps out **the corner the dot is furthest from** — the one nobody is hearing — so four laps renew the whole square, one corner at a time.

Queue

An ordered list with a cursor, always crossfading between the two entries it sits between. Finite where A→B is endless, but walkable backwards. It is no longer offered in the dropdown; a saved track already on it keeps working.

The master flow

The header controls move every unlocked track at once.

- **Flow** — the master dial, and a readout of where the scene is.
- **Drift** — **Together**, **Offset** or **Free**: where the tracks stand relative to one another. This is a **layout**, **applied once, the moment you pick it** — *Together* puts every track on the same number, *Offset* fans them out evenly, and *Free* places nothing at all. After that the dial moves every track by the same amount, so **the distances between them are carried, not replaced**. Move one track's own dial and its new distance from the others is the distance from then on. A **locked** track is left where it is, and so is a track with no weave.
- **Speed** — how long one full journey takes, in bars. **Off** by default: a panel that started travelling the moment you opened it would change a session you had not touched.
- **EVOLVE / STATIC** — what happens at the END of a lap. **STATIC** is a scene you place by hand and it stays placed. **EVOLVE** draws fresh material on arrival, which is what makes a scene keep moving.
- **Journey** — two controls, because how far a journey goes and whether it comes home are two decisions. The button reads → **One way** or ⇌ **There and back**; beside it a length picks **2, 4 or 8 laps**. Going out it draws new loops; coming back it walks the ones it already played, so the way

home is the same journey in reverse rather than a different one. On a one-way journey the length decides nothing, so it greys out and says so — and it keeps the value you gave it, so turning the round trip back on resumes at the same length.

-  **HOLD** — freeze the arrangement. The chord walk keeps going and the step rack keeps writing; only the loops stop moving.
- **SURGE** — held, not toggled: a momentary push that restores itself when you let go.
-  **Reshuffle** — **tap** to re-deal only the loop you are NOT hearing, which is why it is safe mid-phrase; **hold** it to re-deal every loop at once. A locked track ignores both.

What each track PLAYS — parts and shelves

A track's **part** decides which material it draws and in which register: **Bass**, **Comp**, **Pad**, **Arp** or **Melody**. Leave it unset and the track behaves as it always did. An engine can declare its own default — a TB-303 lane comes out as bass whether or not you mark it.


- **Bass** and **Melody** read the pattern library, at their own octave.
- **Comp**, **Pad** and **Arp** read **no** library at all. Their material is GENERATED as chords, because a chord written as fixed semitones cannot stay in key across the eight scales a session may be in. What you pick for them is a **shape** — the rhythm and the voicing — and the notes follow the session's key, scale and chord progression.
- Marking a track **Arp** also gives it an arpeggiator. Unmarking it does not take the arpeggiator away: by then it is a card with your settings on it.

Style picks which shelf of the library this track draws from, independent of the session's own style. A style change lands on the **next** loop drawn, not the one playing — changing it mid-phrase would be a splice, not an evolution.

Half and double time ($\times 2$ / $\div 2$) and the **octave** (- +) both live on the weave, never on the clip. The phrase is always delivered whole: half time stretches it and takes the room it needs.

The list a track draws from

Part and style say what a track **may** draw. The **PLAYS** list, on the fourth line of every track row, says what it **does** draw, and in what order.



Leave it empty and it reads **everything it may**: the track draws from the whole shelf its part and style allow, at random, which is what WEAVE always did. Add a loop and the behaviour changes completely:

- **It is walked in the order you wrote it, and it wraps.** An A→B hand-over goes to the *next* entry rather than drawing at random, so the same four loops come round in the same order, every time.
- **A list of one freezes the track.** There is nowhere to hand over to, so it holds where it is. Falling back to the shelf would play exactly the material you excluded.
- **It outranks the track's own clips and the shelf** — for the hand-over and for the generator both.
- **The picker offers the whole library**, not just this track's shelf. That is deliberate: the list is how you hand a track material from outside its part and its style.
- **A cloud ignores it.** A cloud renews itself by corners, not by hand-overs; the list is A→B's.

Editing: ↑ and ↓ move an entry, ✕ takes it out, and the picker at the end adds one. Buttons rather than drag-and-drop, because the row repaints under your finger. If you delete the loop a travelling track is currently leaving, it rejoins the list at the **top** — the front is what you wrote next.

The list is saved with the session, in the save file.

Following — a track that plays what another track implies

A track can accompany another instead of playing material of its own. Set **Follow** <track> on its second line and it stops reading loops and clips entirely: every bar it works out what to play from what the track it follows is playing.

What it plays is decided by its **part** — the same Bass / Comp / Pad / Arp choice as above. Following says *whose* harmony; the part says *which* player. A bass takes roots and leans on the change; a comp reads a rhythm from the style; a pad takes the nearest inversion; an arp walks the chord's own tones and is deliberately not voice-led, because an arpeggio is heard as a line.


The harmony is **taken, not deduced**, in that order of preference: a progression you wrote by hand, then the session's own, and only then an analysis of the leader's notes. The analysis is the last resort because it is the weakest — a one-bar library loop transposed per bar reads back as a single chord, and most of the shelf does exactly that.

Worth knowing:

- **It reads the leader AHEAD**, not behind. That is what lets a cadence be placed before it arrives, and what makes an offline render come out the same as what you heard.
 - **Its harmony changes only where a BAR does**, and only while it is playing. A chord arriving mid-bar would re-schedule music already sounding.
 - **Following and weaving are exclusive, both ways.** Choosing a leader shelves the track's weave; switching back to — *plays its own* — hands it back exactly as it was. Choosing a topology clears the follow.
 - A **percussion** track cannot follow: it has no part to play. The control is simply not built rather than offered empty.
-

GEN — a read head over the track's own material

GEN is the amber switch beside the follow picker. It is the third thing that can decide what a track plays, and the strongest of the three after following: with GEN on, the track stops playing its loops and starts **reading** them.



What it actually does. The track's loops are folded into a single bar exactly as the weave folds them. That bar's **pitches** become an ordered pool — the rhythm is thrown away — and a read head walks a pattern you set, taking one pitch per step. It invents nothing: everything you hear was already in the material. What you are dialling is *how it is read*.

Where the material comes from, in this order: the track's **PLAYS** list, then its own clips that have notes in them, then the style shelf. With none of the three there is nothing to read, so the switch simply does not come on. Switching GEN on shelves the track's weave selection and switching it off hands it back untouched, so nothing is lost either way.

The controls are the third line of the track row, and they come in five groups.

The grid — how the bar is cut up

Control	Range	Default	What it does
DIV	1–16	4	How many steps cut one bar. 4 is quarters in 4/4; the meter still says how long a bar is, so 6/8 wants 6
BARS	1–16	1	How many bars before the pattern comes round again. This is also the phrase length PHRASE reads
×2^	0–3	0	A power-of-two multiplier on that length — ×1, ×2, ×4, ×8. Separate from BARS because "my phrase is four bars" and "now take four times as long to get through it" are different gestures

Which steps fire

Control	Range	Default	What it does
CADENCE	0–1	1	A floor on metric weight , not a pattern. At 1 every step of the division fires; turning it down drops the weak positions first and the strong ones last; at 0 the track is silent
CAD MOD	0–1	0	How much that floor varies from step to step. At 0 every bar of the pattern has an identical rhythm. It swings either side of the knob, so the average stays where you put it
PHRASE	0–1	0	How much a bar's position in the phrase floors it: the opening bar states, the middle bars thin out, the turnaround comes back

Which note it lands on

Control	Range	Default	What it does
IN KEY	OFF / SCALE / CHORD	OFF	SCALE snaps to the track's key — a passing note the chord does not contain still gets through. CHORD locks to the chord tones of the progression at that bar: it cannot sound wrong, and it cannot sound like a melody. OFF is honest rather than lazy — the blend already walks scale degrees
VOICING	−7...+7	0	Moves the note by whole tones of the set — chord tones on CHORD, scale degrees on SCALE. Not semitones: that would be a transposition, which is what the octave buttons are for
VOICE MOD	0–1	0	Per-step variation on VOICING, up to three tones either way at full depth

VOICING and VOICE MOD do nothing while IN KEY is OFF. Off means off in both halves.

Where the hit lands, and how long it holds

Control	Range	Default	What it does
NUDGE	−1... +1	0	How far off its step a hit sits, in fractions of a step; negative is early. Fractions, so the feel survives a change of DIV
GROOVE	0–1	0	Per-step variation on NUDGE. This is what makes it a groove rather than a shift — at 0 every hit moves the same way, which just reads as the track being late
HOLD	0.05–4	1	Note length as a multiple of its step. 1 fills the step, below it detaches — and above 1 the notes overlap, which is how the generator slides : on a 303 (or any engine that infers portamento from overlap) this is the glide control
HOLD MOD	0–1	0	Per-step variation on HOLD, up to a step either way

The two wheels — six controls that move the read head itself, three per wheel: **x** how far one turn moves it, **CYC** how many turns the wheel takes to come round (1 means it is not turning), and **%** how much of it is applied. The **%** is the one that switches the wheel on; the other two describe a wheel that is not yet moving anything.

- **BAR x/CYC/%** turns once per bar of the pattern, so what it does repeats when the pattern does — a figure you can learn.
- **LOOP x/CYC/%** turns once per lap, so it never repeats inside a pass. This is the one that reaches material a short pattern would otherwise never get to.

Two readings that explain most surprises:

- **The pool is read on every step, whether or not that step fires.** Turning CADENCE down gives you *the same line with holes in it*, not a different line.

- **CADENCE reads the pattern's own bar; IN KEY reads the song's.** The rhythm repeats with the pattern, the harmony walks the progression.

Worth knowing:

- **BAR MOD does nothing out of the box.** The default pattern is one bar long, and a wheel that turns once per bar has one turn to make. Raise BARS (or $\times 2^{\wedge}$) first.
- **A pool longer than the pattern has a tail nobody hears** until LOOP % is raised.
- **A scene launch does not stop the generator.** The grid takes the *weave* off a track; generating and following are properties of the song, saved on the track, and they carry on.
- **Switch GEN off before you give a track a leader.** Turning GEN on clears a follow, but picking a leader does not turn GEN off — and following wins, so the switch stays lit over a track that is not generating.

Everything here is saved with the session and undoable, and it lives on the track rather than in the panel: close WEAVE and the track keeps generating.

How long a track takes to repeat itself

Beside the follow picker is a five-rung ladder: **Loop · Turns · Travels · Wanders · Never repeats.**

It does not add notes and it does not make anything busier. It buys **length** — how long the track plays before it plays the same way twice. At *Loop* the accompaniment is a loop in the ordinary sense. Each rung up brings in another wheel that turns at its own rate: which figure it uses, which colour the chord takes, which register it sits in, how dense it leans. Because those rates share no common factor, the combination takes far longer to come round than any of them does alone.

Set it by ear rather than by the label. *Turns* is the default and is enough for most tracks. The top of the ladder is a long piece of music, and on a short one it can read as drifting rather than developing.

Playing one part on several instruments

The **●** button is a second axis: the weave decides which NOTES play, this decides **what they are played on**.

Turning it on builds what it needs — the track becomes a **rack** of instruments, its own sound in the first slot and others alongside, each on its own preset. The control then wears the shape of that track's loop control: a **fader** on A→B, a **square** on a cloud. Off, it collapses back to just the button.

Without it, each note plays on the instrument of the loop it came from. With it, every note reaches every instrument and the control balances them. Either the loop chooses the instrument or you do — never both.

The 1 2 buttons on the second line pick which slot the instrument and preset dropdowns are pointing at.

The chord progression

Every chordal track follows one progression, walked across the session's bars and shown as a bar with the current chord in Roman numerals. Pick a catalogue entry, or edit the strip and it becomes yours — the catalogue is never written to. "Stay home" is a single chord and no journey.

The progression travels too, in the same legs the loops do. The first chord never moves — home is what a departure is measured against — an interior chord may be swapped for its diatonic relative, and the last alternates with V as a turnaround. Only ONE chord moves per leg: harmony is far less forgiving than timbre, and two changing at once stops reading as a variation and starts reading as a wrong chord.

If the progression is longer than a track's clip, the track hears a **window** of it at a time rather than the first two chords for ever; a short progression tiles instead.

The progression is the session's, not the panel's. It is stored with the key and the scale — the same question asked at three different rates — which is why the Chord note-FX's *chord* setting and the generator's IN KEY = CHORD can both read it, on any lane, whether or not WEAVE is open.

The macros

Four knobs colour everything at once, and each of them changes the MUSIC rather than the treatment of it.

- **Density** and **Energy** rewrite the notes themselves.
- **Mood** chooses the scale the blends walk their degrees in, from brightest to darkest. At its neutral it does not mean "no scale": the session's own scale is home, and a travelling scene still wanders a step either side of home and holds there a couple of legs. Mood only says where home is.
- **Style mix** decides how far a track may stray from the session's style, and how often. A stray steps one to three places along a catalogue ordered by family, so a downtempo scene drifts to its neighbours instead of landing in jungle and back.

At their neutral positions they say nothing at all.

There used to be six. **Space** and **Motion** are gone: they were the only two that moved PARAMETERS rather than notes, and a knob that sweeps a reverb is spent after its first pass while one that rewrites the music keeps giving.

The step rack

Below the tracks is a row of bars pointing at a **destination** — any parameter a knob can move. It moves that parameter **in time with the loop** rather than on a curve of its own, which is what a step row on an old sequencer was for. + **ROW** adds another; each row has its own length and its own mode, so a sixteen-step row cannot silence a four-step one.

Its value goes to the audio object and never to the track's saved sound: the row owns the value, and stamping its momentary position into a preset is the mistake that door exists to avoid.

Switching the whole thing off

🔌 **WEAVE ON / OFF** unplugs the panel. Off, it stops the transport and every track goes back to playing exactly what it played before **WEAVE** existed — nothing is undone, nothing is lost, and plugging it back in does not start anything on its own. It is the honest way to hear the session underneath.

+ **WEAVING TRACK** adds a track that arrives already weaving: a new lane, an empty carrier clip, and its first two library loops picked for you. A track that arrived empty would leave the panel as useless as it was.

PRINT — keeping what you heard

PRINT freezes what the weave is playing into a new scene: a full lap of the chord progression, laid end to end, as ordinary clips you can edit. It is an output, never the goal — the weave carries on folding either way.

Launching a printed scene plays what was printed, even on a track that is still weaving — the grid wins, as below. (This used to be a known limitation, with a note telling you to switch the track's topology off first. Do not: choosing a topology is now exactly the gesture that takes the track *back* from the grid.)

Two things worth knowing

The grid wins. Launch a scene and every weaving track hands itself back to it: each one plays the clip that scene gives it, or is silent where the scene leaves it empty. Launching a single clip from the grid does the same for that track alone. Nothing is destroyed — the topology, the loops, the position and the list are all still there — and the panel takes the track back the moment you use it: its own ►, a topology, or choosing or moving a loop. While the grid holds a track, **Play skips it**, and the panel's own ► is how you hand it back.

What the grid takes is the **crossfade**, not the whole track. Following and generating are properties of the song rather than of an open panel, so a track that follows or generates goes on doing it through a scene launch.

(This used to work the other way round, and it is worth knowing which build you are on: a weaving track once read its own selection and nothing else, so every scene sounded the same on it — including a scene of empty clips.)

Nothing here is undoable. A weave move is deliberately not an undo entry; it is a performance, not an edit. It IS saved with the session, and it nudges the autosave when it moves.

Developer Guide

This chapter is for contributors who want to extend Loom or understand how its internals fit together. Read it alongside `CLAUDE.md` at the repo root, which is the shortest architecture summary the project keeps — this chapter expands on it in prose.

Both drift. Where a document and the code disagree, **the code is right**: every claim below carries the file it was checked against, so you can check it again.

The spine

Three structures hold everything together:

1. **A plugin registry** — engines, FX and modulators are all plugins, and most of them are no longer in this repository's source at all. A plugin is a directory under `plugins/`, compiled into `public/plugins/<id>/`, and loaded by the browser **at runtime** off a JSON manifest; it compiles against `@loom/plugin-sdk` and talks to the host through one small runtime ABI, `globalThis.Loom`. Six engines, fifteen inserts, two modulators and one whole VIEW (the WEAVE panel) ship that way today. A shrinking in-tree SPI, discovered by a Vite `import.meta.glob` scan, still carries the LFO/ADSR modulators and the note-FX.
2. **SessionState** — the pure data model: lanes contain clips, scenes reference which clip each lane plays. No audio side-effects live here.
3. **LaneResourceMap** — owns the live Web Audio nodes for each lane. One entry per lane, holding a `ChannelStrip`, a `SynthEngine` instance, and an `InsertChain`. The lane allocator in `src/app/lane-allocator.ts` is the sole path for creating and swapping these resources; nothing else should construct them directly.

Boot

`src/main.ts` is a boot script, not a feature file. It builds the handful of objects everything else needs, then hands them to one wiring module per concern in `src/app/`:

- `bootstrapPlugins()` first, before anything reads the registry (`main.ts:92`), then `loadPlugins()` — the external plugins, fetched and validated (`main.ts:98`). That second one is a **promise**, and the difference matters: everything the in-tree registry knows is available synchronously, while every engine and insert that arrives from `public/plugins/` lands later. Code that enumerates components at boot sees only the built-ins unless it awaits it, which is written down at `main.ts:647`.
- `createAudioGraph()` (`main.ts:120`) and the worklet `addModule` calls, which every lane allocation waits on (`main.ts:136 – 154`). The order inside that block is load-bearing too: `loom-processor` installs `globalThis.Loom` and must be registered **before** any plugin's `dsp.js` is added, or the

plugin has nothing to register its renderer with.

- The `Sequencer` (`main.ts:162`), the `DestinationRegistry` (`main.ts:170`), the lane allocator (`main.ts:192`), and the `SessionHost` (`main.ts:377`).
- Then roughly a dozen `wireX(...)` / `createXFeature(...)` calls, each of which owns one concern end to end and lives in its own file under `src/app/`.

The order is load-bearing, and the comments in `main.ts` say why at each step. `wireMenuBar` is deliberately the last statement of `boot` (`main.ts:810`): its action table names a handle from nearly every feature above, so the bar appearing is the proof that boot ran start to finish.

If you are looking for the code behind a control, it is almost never in `main.ts` — find the `wireX` call that mentions it and open that module.

How a plugin gets in

The external ABI — the normal case

The browser cannot list a directory, so `public/plugins/index.json` is the **discovery mechanism**: a flat array of plugin ids, rewritten by `npm run build:plugins`. For each id, `loadPlugins` (`src/plugin-host/plugin-host.ts`) fetches `plugin.json`, validates it as **data** before a single line of the plugin runs (`manifest-validate.ts`), adopts its declared components into the engine / modulator / fx registries, seeds any `presets.json`, and collects its `dsp.js` URL for the worklet.

Two properties of that loop are worth relying on:

- **Failure is contained and complete.** A plugin that throws — even one whose components were already adopted before its own module blew up — is rolled back in full and recorded in the load report. One bad plugin never takes the app down, and never leaves a half-registered engine behind: visible in the selector, silent at note time, is exactly the failure this design refuses to allow.
- **A plugin's JS is loaded through a blob: URL**, not a bare `import(url)` (`module-loader.ts`). A plugin is a static asset in `public/`, and Vite's dev middleware refuses to serve one as a module — so the obvious import works in the production build and breaks under `npm run dev`. Fetching the source and evaluating it from a blob behaves identically in both, and is the same mechanism a user-installed plugin will need when its bytes live in IndexedDB instead of at a URL.

The worklet half is symmetrical: `loom-processor.ts` installs `globalThis.Loom` (`registerRenderer`, `registerModulatorKernel`) and the host guarantees that module is `addModule` 'd **first**, so every plugin `dsp.js` added afterwards shares the realm and can reach the registry. The ABI deliberately carries no DSP — everything a plugin uses from `@loom/plugin-sdk` is compiled into the plugin's own bundle, which is what lets the ABI stay stable across versions.

What is left of the in-tree SPI

`src/app/plugin-bootstrap.ts` still calls `import.meta.glob` at build time over `src/engines/*.ts` and `src/plugins/**/*.ts` (`*.test.ts` excluded), and any exported value matching the `PluginFactory` shape (`{ kind, manifest, create }`) is registered. **No insert is found this way any more** — all fifteen are external. What the glob does today is mostly to *evaluate* modules so their module-scope registrations run:

- **Modulators** call `registerModulator` at module scope into `src/modulation/modulator-registry.ts` — their own door, the same one a plugin modulator arrives through. The shape check does not match them.
- **Note-FX** declare `defaultParams()` instead of `create()`, so the shape check skips them too; those files call `registerPlugin` themselves (`src/plugins/notefx/arp.ts:11`).

`bootstrapPlugins` then bridges every registered engine descriptor into the plugin registry, so `listPlugins('engine')` keeps seeing all nine.

The engine registry (`src/engines/registry.ts`) supports both a **singleton** pattern (`registerEngine`) for shared instances and a **factory** pattern (`registerEngineFactory` / `createEngineInstance`) for per-lane instances that need independent state.

Never ask an engine its name

Anything the core needs to know about a component goes through **one door**,

`src/plugins/capabilities.ts` — `clipContent`, `slide`, `outputTrim`, `acceptsNoteFx`, `defaultNoteView`, and so on. A built-in answers from code, a plugin from its manifest, and the caller cannot tell which; that is the whole point, because it is what let six engines move out of `src/` without the core noticing. Every accessor there returns a safe default for an unknown id and never `undefined` — an unregistered engine that blanked its lane's UI would fail silently. **An `engineId === '...'` comparison anywhere outside that file is a bug.**

`listEngines()` reads from the singleton map and is the source of metadata (name, type, polyphony, parameter specs) used to populate the lane engine selector.

What "declaring a param" actually buys you

Engines declare their parameters as `EngineParamSpec[]`, and for the catalogue that is the whole story:

`listAutomationTargets` walks `engine.params` and pushes ``${lane.id}.${spec.id}`` for every continuous one (`src/automation/automation-targets.ts:120`). Automation and the modulation dropdown both read that catalogue through `DestinationRegistry.list()`. Neither of them calls `getAudioParams()`.

`getAudioParams()` / `getSharedAudioParams()` are a narrower thing — the **Web Audio binding surface**, used only where a modulator has to reach a real `AudioParam` through the depth-gain bridge in `connection-binder.ts`. That means FX inserts, channel strips, and the two engines that still expose shared params, Drums and Sampler. For the six melodic worklet engines `getAudioParams()` returns an empty `Map` on purpose (`src/engines/worklet-lane-engine.ts:150`); their modulation is applied per sample inside the worklet by `ModulationRuntime`.

So: declare the param in the spec and it is automatable and modulatable. Return it from `getAudioParams()` only if it is a genuine `AudioParam` on the main thread.

That rule is why the lane **mixer** is automatable at all. The seven strip controls — level, pan, sends A and B, and the three EQ bands — were once declared inline by `drums-machine` alone, which is why a drum lane's volume could be automated and a Subtractive lane's could not: the capability belonged to one engine's source file rather than to the mixer. They now live in `src/core/channel-strip-params.ts` and every engine with a `ChannelStrip` spreads `STRIP_PARAM_SPECS` into its own params, keeping the original `bus.*` ids so existing drum-lane envelopes and connections still resolve. One caveat if you touch the binder side: a **modulator** binds to the strip's multiplicative trims, not to the real gains, because a bipolar modulator summed onto a gain can drive it below zero — which inverts phase instead of quietening, and the lane cancels against the rest of the mix. Automation writes the real gains through `setStripParam`.

SessionState data model

`src/session/session-types.ts` defines three levels (re-exported from `session.ts`, which holds the factories):

- **SessionLane** — has an `engineId`, a list of `SessionClip` | `null` slots, and an `engineState` bag that persists knob values, modulator configs, note-FX, sampler keymap, pad params, and kit mode.
- **SessionClip** — holds `notes: NoteEvent[]` (the unified note list for both melodic and drum clips), optional `clipEnvelope[]` for per-clip automation, and an optional `sample` field for loop/song audio clips. Clips also carry `loopEnabled` / `loopStartTick` / `loopEndTick` for sub-region looping, and a `gridResolution` hint for the drum editor.
- **SessionScene** — a `clipPerLane` map from lane id to clip slot index (or null for a stopped lane).

Notes carry a `velocity` field (0–127). The `velToColor` function in `src/core/velocity-color.ts` maps velocity to a blue-to-yellow ramp used by both the piano roll and the drum grid.

Saves are written as `schemaVersion: 3` (`SavedStateV3` in `src/save/`). Older saves are normalised by `session-migration.ts` at load time before anything else touches the data.

Replacing the session is New plus a load, in that order. Every route that swaps the whole session — New, the boot demo, the demo picker, a save or autosave load, an import that replaces — calls `resetAllResources` (`src/session/session-host-reset.ts`) *first*, and only then applies the incoming state. The reason is that applying a session only ever pushes the fields the incoming JSON happens to carry, and `ensureLaneResource` is idempotent: a lane id present in both the old and the new state used to keep its live `ChannelStrip`, so level, pan, EQ, sends, mute, compressor and sidechain rode straight through New and through every demo switch — as did the master rack, the send buses, mute/solo and the note-FX chains. The fix is not a longer list of restores at the apply site (that list goes stale the next

time something joins the desk); it is that a load starts from a released desk, so anything the new state does not set sits at its constructed default. The one deliberate exception is the master **volume** fader: that is the listening level of the room, not a property of the song.

LaneResourceMap and the audio graph

The master audio path assembled in `src/app/audio-graph.ts` runs:

```
master (sum GainNode)
  → MasterBusStrip (EQ / pan / mute)
  → InsertChain (the master rack)
  → MasterShaper (air / glue / width)
  → MasterCompressor (the safety limiter)
  → soft-clip WaveShaper (4x oversampled)
  → AnalyserNode → ctx.destination
```

Two details of that chain are deliberate. The shaper sits **before** the limiter, because air/glue/width are mix decisions and the limiter must be the last thing that sees the signal. The soft-clip after it is the absolute ceiling: identity below ± 0.8 , then a tanh knee that maps everything above — including overs beyond ± 1 — to about ± 0.95 , so the master output cannot digitally clip.

A second analyser, `masterMeterAnalyser`, taps off the soft-clip and is **not** connected to the destination. It feeds the master VU meter and the PERF peak/clip readout, so both read the true, clip-free output.

`SidechainBus` is not a node in that chain at all. It is a lane-id → tap registry (`src/core/sidechain-bus.ts`): each `ChannelStrip` registers a `GainNode` fed off its post-mute output, and a ducker subgraph reads `getTap(sourceLaneId)` to drive its envelope follower. The allocator hands it to every lane strip it builds.

That follower is an **AudioWorklet** (`duck-processor` / `duck-node`, wrapping the pure `DuckDetector` in `src/audio-dsp/duck-detector.ts`), not a filter chain. It was two `BiquadFilterNode` lowpasses until 2026-07-27: at a 0.25 s release the cutoff is 0.64 Hz, the pole pair sits a hair from $z=1$, and in the float32 node graph the rounding error accumulated like an integrator — the "envelope" grew without bound with the input at exactly zero, so the duck multiplier drifted past 0 and came back negative, i.e. the ducked lane went silent and then returned phase-inverted and louder. A one-pole follower is a convex combination of the input and itself, so it is bounded by construction at any time constant and returns to zero when the source stops. Keep the multiplier's `[0, 1]` invariant if you touch it; `src/audio-dsp/duck-detector.test.ts` and `src/core/strip-ducker.dsp.test.ts` pin it.

Each lane's `LaneResources` consists of a `ChannelStrip` (level, EQ, send levels), a `SynthEngine`, and an `InsertChain` of per-lane FX. `LaneResourceMap.replaceEngine` hot-swaps only the engine while keeping the strip and inserts in place — the channel-level resources survive an engine swap.

The lane allocator (`src/app/lane-allocator.ts`) is the only module that constructs a `LaneResources` , in exactly one place: `ensureLaneResource` (line 228). The second entry point, `ensureExtraPoly` , went with the `PolySynth` class it existed to feed. Call `ensureLaneResource` once per lane before accessing anything in the map. Test code that needs a lane wired up must call it explicitly as setup.

The scheduler

The `Sequencer` class (`src/core/sequencer.ts`) fires every 25 ms (the poll interval) and looks **200 ms** ahead. On each tick it calls `sessionTick(now, lookaheadSec)` with `lookaheadSec = 0.2` , and the session host fans that out to `tickLane` for each playing lane.

`tickLane` (`src/core/lane-scheduler.ts`) implements the Chris Wilson two-clocks pattern: for every `NoteEvent` whose absolute schedule time falls in the window `[now, now + lookaheadSec)` , it calls `ctx.onTrigger` . Schedule times are derived by converting clip-tick positions to seconds using the current BPM and projecting onto the absolute timeline from the loop-start anchor. Step duration for a 16th note is `60 / bpm / 4` seconds.

Two important consequences for contributors:

- `bpm` and `length` are mutable at runtime; the next scheduled step picks up the new values immediately.
- **Continuous** engine params reach the note already sounding. Each lane keeps a smoothed live bag (`ParamSmoother` , driven by `VoiceManager`) that every voice re-reads per sample, so turning a knob bends the held note instead of waiting for the next trigger. **Structural** params are the exception and still apply to the next trigger only: waveform, filter model, unison size, and every envelope *time* — the envelopes are closed-form over elapsed time, so re-reading an attack mid-note would step the amplitude. Drums is deliberately outside this: its params are read at trigger time.

The scheduler asks `laneLoopRegion` (`src/core/clip-loop.ts:40`) how long one iteration of a clip is, and there are **two** ways the answer comes back shorter than the whole clip. The active scene's global loop wins first: when `GlobalLoopOverride.enabled` is set, `[startBar, endBar)` becomes the region for every lane in the scene, whatever the clips say. Absent that, `effectiveClipLoop` (line 19) applies the clip's own `loopEnabled / [loopStartTick, loopEndTick)` . The brace UI in `src/core/clip-loop-brace.ts` is the editing surface for the clip's own region.

That precedence has a consequence worth knowing before you touch clip automation: a clip's envelope array spans the clip's `lengthBars` and is blind to **both** shortenings, so inside a shorter loop the curve slides against the notes. It is written down as known debt at the top of `src/core/clip-envelope-length.ts` and pinned by tests — meet it as a decision, not a mystery.

How-to recipes

Add a synth engine

Write a plugin. Do not add a file to `src/`. All six melodic engines are external plugins, and yours should be too — it needs no access to this repo's source, and it is a directory no future PR review has to carry. Full walkthrough in [docs/plugin-development.md](#); the shape:

1. **Scaffold** — `npm run plugin -- new plugins/<id>`.
2. **Declare it in `plugin.json`**. The manifest is the whole declaration: `id`, `loomApi: 1`, and a `components` array. An engine component carries `kind: "engine"`, `polyphony`, its `params`, a `groups` table (which sections exist, their title, order, colour, and which share a row — omit it and the grid falls back to first-appearance order, one row per group, no colour), an optional `modulators` array shipped with the sound, and a `capabilities` block. Add `"dsp": "dsp.js"` and `"presets": "presets.json"` alongside.
3. **Write `dsp.ts`** — the pure per-sample voice renderer, registered with `Loom.registerRenderer(id, ctor)` at module scope. It is plain TypeScript with no `AudioContext`: unit-test it directly, next to the code, in `plugins/<id>/dsp.test.ts`.
4. **Implement the live-params hook** — read your **continuous** params out of the live bag every sample so a knob moves the note already sounding; copy the **structural** ones — waveform, filter model, unison size, envelope times — into your own fields once, at construction, from the trigger-time snapshot. The hook is optional, so a renderer that skips it compiles clean and passes the whole suite; it is just the one engine whose knobs go dead mid-note. The registry-driven test in `src/audio-dsp/live-params.dsp.test.ts` catches it.
5. **Build it** — `npm run build:plugins`. **Nothing you did is visible until this runs:** the app loads `public/plugins/`, and the dev server does not compile `plugins/` for you. A stale build looks exactly like a broken engine.

There is no `WORKLET_ENGINE_IDS` list to add yourself to any more — it is a live view over "is this id worklet-hosted", and any plugin id answers yes.

Adding an engine in-tree is the exception and needs a reason. The three that still live in `src/engines/` — `sampler`, `audio`, `drums-machine` — are there because each owns browser resources an external plugin has no way to hold: decoded buffers, IndexedDB, per-voice `AudioParams`. "I have a checkout, so it is easier" is not that reason.

If you do write one, two things about the registered descriptor look like bugs and are not:

- **Its synthesis surface is inert, not throwing.** `createDescriptorEngine` gives you a `createVoice()` that returns a no-op `Voice` with an empty `getAudioParams()` (`src/engines/descriptor-engine.ts:78` → `:48-50`). Nothing on the live or offline path calls it — the registered singleton is purely metadata, and modulation for these engines runs sample-accurately inside the worklet (`src/audio-dsp/modulation-runtime.ts`).
- **The bridged plugin does throw.** `bootstrapPlugins` wraps each engine descriptor in an engine `PluginFactory` so `listPlugins('engine')` keeps seeing every engine, and that wrapper's `create()` throws on purpose (`src/app/plugin-bootstrap.ts:90`). It is a tripwire: if you see it, something called `createInstance('engine', ...)` instead of going through the lane allocator.

See [Engines](#) for the full engine catalogue.

Commit an engine param edit through one seam

Any control you build for an engine param must write it with `commitParam` from `src/engines/engine-param-commit.ts`, never `engine.setBaseValue` alone:

```
commitParam(engine, ctx, paramId, value); // engine + the engineState mirror
```

`setBaseValue` moves the sound. It does **not** persist it. The mirror into `lane.engineState.params` is the only vehicle by which a knob value reaches a save, and builders that forgot it threw the edit away silently — that was the knob-loss bug on FM, Wavetable, Karplus, Westcoast and TB-303, fixed by routing every builder through this one seam.

Two siblings exist for the cases a UI context cannot cover:

- `commitParamForLane(engine, sessionState, laneId, id, v)` — same seam for a caller that holds the session directly, e.g. a MIDI control surface writing a lane whose editor is closed.
- `commitEngineBaseValues(engine, sessionState, laneId)` — the bulk sibling for the programmatic applies that move a whole sound at once (recall a preset, load a user preset, Randomize). Those push values straight into the engine, so no `onChange` fires and `commitParam` never runs.

`withoutParamMirror(...)` suppresses the mirror. The load path uses it to apply a lane's preset without clobbering the saved params it is about to replay — a saved tweak beats its lane preset.

Add an FX insert

An insert is a plugin like an engine, with one structural difference: **it is not a worklet**. Only synthesis runs in the worklet; an insert builds ordinary Web Audio nodes on the main thread.

1. `npm run plugin -- new plugins/<name>`.
2. In `plugin.json`, declare a component with `kind: "fx"` and its `params`.
3. Write `main.ts` — the factory that builds the node graph. The SDK's main-thread builders are there for you (`modulated-delay`, `signal-max`, `envelope-follower`); they sit **below the dividing line** in `packages/loom-plugin-sdk/src/index.ts`, and importing one into a `dsp.ts` renderer will not work, because that half runs inside the worklet with no `AudioContext`.
4. `npm run build:plugins`.

The "+ Add insert" picker is an unfiltered `listPlugins('fx')` (`src/session/lane-insert-ui.ts`), and the same builder serves lanes, the master rack and both send racks, so a new insert appears in all four at once.

Separately from the picker, the `FxBus` seeds send A with `delay` and send B with `reverb` when it is constructed (`src/core/fx.ts:30`). That is a default, not a restriction — both are offered as ordinary inserts too. **Seeding needs the plugin registry to be up**: the `FxBus` cannot seed a send with an insert that has not loaded yet.

Add a modulator

A modulator can be a plugin too — `plugins/sh/` (sample & hold) is the shipped proof. Declare `kind: "modulator"` with a `modulator` block naming its `driver` (`time` or `gate`), the `scopes` it supports (`shared`, `per-voice`) and an `idPrefix`; a `time`-driven one also ships a per-sample kernel registered with `Loom.registerModulatorKernel` inside the worklet. In-tree modulators come through the same door, `registerModulator` in `src/modulation/modulator-registry.ts` — the LFO and the ADSR just call it directly.

Binding is unchanged: `ConnectionBinder.apply` builds `modulator.output → GainNode(depth × (max - min)) → targetAudioParam` (`src/modulation/connection-binder.ts:44`) for the native-node targets, while a worklet engine's params are modulated per sample by `ModulationRuntime`.

See [Modulation and Note FX](#) for the user-facing side.

Add a note-FX

`kind: 'notefx'` is the fourth plugin kind — a transform applied to notes before they reach the engine, per lane, persisted in `lane.engineState.noteFx`. A `NoteFxFactory` declares `defaultParams()` and has no `create()`, so the bootstrap's shape check ignores it: the file **must** call `registerPlugin` itself at module scope (`src/plugins/notefx/arp.ts:11`). The processor that does the work lives beside it in `src/notefx/`.

Add a preset

For one of the six melodic engines, the presets ship **with the plugin**: append an entry to `plugins/<id>/presets.json` and run `npm run build:plugins`. `plugin-host` seeds them at load. For Sampler or Drums, the file is still `public/presets/<engine>.json`.

Either way the `gm` field is optional (an integer GM program number for MIDI-import matching), JSON is the source of truth, `preset-loader.ts` validates and `preset-apply.ts` applies it at runtime by calling `engine.applyPreset`. Each engine's JSON keys are its own vocabulary — do not use a generic `setBaseValue` loop.

Add a synth drum kit

Append an object to the `KITS` array in `src/core/drums.ts`. Kits are parameter bags over shared DSP primitives. To add a new drum *voice* (not just a new kit): extend the `DrumVoice` union, add it to `DRUM_LANES`, add an entry to every kit, implement a `play<Voice>()` method, and add a `trigger()` case.

Add a sampled drum kit

1. Create a subdirectory `public/drumkits/<id>/` containing WAV files for each voice (e.g. `kick.wav`, `snare.wav`, `closedHat.wav`).
2. Add a manifest file `public/drumkits/<id>.json` with `id`, `name`, and a `samples` array. Each entry needs `voice`, `note` (GM MIDI note number), and `file` (path relative to `public/drumkits/`).

3. Register the kit in `public/drumkits/index.json` by appending `{ "id": "<id>", "name": "<display name>" }`.

68 sampled kits ship under `public/drumkits/`; `tr808`, `acoustic` and `dirt` are the hand-curated three and the clearest reference for this layout.

Conventions

File size: 300 lines of code as a target, 500 as a hard cap. Lines of *code* — comment lines and blank lines do not count towards either number. The distinction is not pedantry: `src/main.ts` is 816 raw lines and 473 lines of code, which is inside the cap by the rule that applies and over it by the one that does not. A file that is long because it explains itself is fine; a file that is long because it does too much is the thing the cap exists to catch. When one crosses the line, split it by concern — `src/app/` is what that looks like in practice.

Assertions are relative. See Testing, below.

Source layout tour

src/

- core/
 - DSP primitives + pure logic (drums, sequencer, lane-scheduler, lane-resources, fx, meter, notes, history, knob, pianoroll, ...). The reference TB303 class that used to live here (synth.ts) is gone with the last node-per-note path; the renderer is the reference now
 - velocity-color.ts / velocity-gain.ts / velocity-lane-editing.ts
 - note-velocity colour ramp, gain curve, lane editing helpers
 - clip-loop.ts / clip-loop-brace.ts
 - clip sub-region resolver + drag-brace UI primitive
 - channel-strip-params.ts
 - the seven strip params (level, pan, sends A/B, EQ x3) declared ONCE: engine params, automation targets, modulation targets and mixer knob ids all read this
- engines/
 - SynthEngine abstraction, registry, and the HOST side of an instrument. Only FOUR engines still have a file here – sampler, audio (the dedicated audio channel), drums-engine and layers – because each reaches for something a plugin cannot: the first three own browser resources, and layers builds OTHER engines through the renderer registry, which the plugin ABI deliberately forbids. The other six arrive from plugins/ and register the same descriptor shape from their manifest. Ten in total.
 - layers-engine.ts + layers-rack-ui.ts – the host half of the rack; it carries itself on three declared hooks (dynamicParamsFor, structuralFor, extraUI) rather than an engineId === ... anywhere in the core
 - Also: the lane engine wrappers (worklet-lane-engine, sampler-/drums-/audio-worklet-engine), engine-selector UI, engine-param-commit (the one write path for a param edit) and engine-randomize (the 🎲 dice, derived from each engine's declared EngineParamSpec rather than per-engine knowledge)
- plugin-host/
 - How an external plugin gets in: plugin-host (index.json is the discovery mechanism; validate-as-data, then roll back in full if it throws), manifest-validate, module-loader (blob: URLs, because Vite dev refuses to serve public/ as a module) and loom-api (the main-thread half of globalThis.Loom)
- session/
 - SessionState model + all session UI (session-host, session-ui, session-inspector, clip-editors/, session-migration)
 - session-host-reset.ts – resetAllResources(): the ONE teardown every route that replaces the session runs FIRST, so a load starts from a released desk
- modulation/
 - LFO/ADSR voices, ModulationHost, ModulatorScope, connection binder
- plugins/
 - What is left of the in-tree SPI. There is NO fx/ directory any more – all fifteen inserts are external plugins
 - capabilities.ts – the ONE door through which the core asks what a component can do. An engineId === '...' anywhere else

is a bug

modulators/ – lfo, adsr (they call registerModulator at module scope, the same door a plugin modulator arrives through)

notefx/ – arp, chord, random (the processors themselves live in src/notefx/)

presets/ Preset loader + apply logic. A melodic engine's presets ship with its plugin; public/presets/ now holds only what the in-tree engines need

midi/ SMF parser, MIDI-to-session transform, GM lookup, import UI

samples/ Sample types, IndexedDB store, buffer cache, keymap, import metadata

stems/ Stem-separation client + config + lane-plan builder (talks to the local Python service in tools/stem-service/)

performance/ Arrangement / record model: arrangement-from-session, arrangement-ops, arrangement-runtime (records clip-launches + knob automation; surfaced via the REC group's take mode – see performance-feature)

audio-dsp/ THE SYNTHESIS KERNEL, minus the engines. voice-manager, scheduler-queue, modulation-runtime + modulator-kernels, and renderer-registry – which is now filled AT RUNTIME by plugin dsp.js modules, not by in-tree imports; there is no <id>-renderer.ts per melodic engine any more. The DSP that stayed in-house is drums/ and sample/. The shared primitives (osc, filter, ladder, sync-osc, unison, fold, adsr) moved to packages/loom-plugin-sdk, where a plugin can reach them. No AudioContext – unit-test directly

param-index.ts – params are addressed BY INDEX in the audio loop, so a per-sample render does zero name lookups

param-smoother.ts – the lane's live param bag: slews only the params still in flight, and is what makes a knob move the note already sounding

duck-detector.ts – the sidechain envelope follower (one-pole, asymmetric attack/release), run by the duck worklet

layers/ – LAYERS: a voice made of other voices. layer-spec (the rack, and the TWO ways a note picks a slot – by keyboard ZONE when a person plays, by an INDEX the note carries when something upstream already knows) and layers-renderer, which builds sub-renderers through the registry and sums them

audio-worklet/ The processors + typed node wrappers: loom-processor/loom-node (melodic), drums-*, sampler-*, duck-* (the sidechain follower). A processor is referenced ONLY via ?worker&url and its registered name – never imported on the main thread (see processor-name.ts)

export/ Offline scene/WAV render + the live take recorder

patterns/ The pattern library (styles x patterns) + its picker UI

perf/ Performance diagnostics (the PERF HUD)

instrument-presets/ The instrument page's USER-PRESET SURFACE, and nothing else:

the preset dropdown + Randomize, apply, store, templates and the param id list (poly-params). Was src/polysynth/ before the poly → instrument rename; the PolySynth class it was named after went with the worklet cutover. User presets are keyed BY ENGINE under 'loom-user-presets-v1'; the old 'tb303-poly-presets-v1' key is gone, and nothing reads it, because no released build ever wrote one. There are no migrations in this project. The `poly` names that ARE load-bearing live elsewhere: polyBlep in the SDK DSP, the manifest's polyphony field, and the poly group / poly.voices param id carried by five plugin manifests

app/ Boot wiring, one module per concern (37 files). main.ts calls into these; it does not contain them – see "Boot" above

audio spine – audio-graph, lane-allocator, engine-swap, trigger-dispatch, live-voice-registry

boot wiring – transport-controls, engine-selector-wiring, midi-control-wiring, midi-import-wiring, import-lane-prep, recording-feature, stems-feature, session-lifecycle, save-history-wiring, menu-wiring, xy-panel-wiring, knob-menu-wiring, automation-writes, lane-host-wiring

UI plumbing – knob-mounting, knob-registry-prune, mute-solo, bpm-broadcast, track-ids, toolbar-status-chips, about-dialog, modal-dialog, and the four menu files (menu-spec / menu-actions / menu-bar / menu-shortcuts) that menu-wiring mounts

features – performance-feature, arrangement-playback, automation-recording, stretch-resync, warp-resync

plugin-bootstrap – the build-time glob scan, now mostly a way to RUN module-scope registrations (see above)

save/ SaveManager (schemaVersion: 3), auto-history (AutoHistory: snapshot-diff undo/redo + gesture coalescing, wired to the transport-bar ⏮/⏭ buttons), history-wiring (withUndo / attachKnobUndo + the undo keyboard – LIVE and load-bearing: withUndo wraps mutation sites across the app)

notefx/ Note-FX processors (arpeggiator, chord spread, random) – per-lane, applied to notes before they reach the engine. The arp's `steps` pattern is written in POSITIONS rather than notes (arp-steps.ts), so it survives a transpose or a change of key, scale or octave

weave/ WEAVE's pure core – no DOM, no AudioContext, no module state. The blend (blend-rhythm / blend-melody / blend-clip), the three topologies, the master flow, the macros, the loop ids, loop-pool (the list a lane draws from, walked with wrap) and weave-runtime, which answers the scheduler's "what does this lane play" through SchedulerContext.notes

harmony/ The DERIVED accompaniment: what a lane plays when it FOLLOWS another rather than reading clips. follow-source is the door,

generator/	shaped as the same WeaveSource the weave hands the scheduler; render-part picks between parts/ (bass, comp, pad, arp); cycle is the arrangement ladder that buys LENGTH, not density GEN – a read head over a lane's own material rather than a producer of new notes. grid (how a bar is cut), pool (the blended bar's pitches, ordered), cadence (a floor on metric weight), chord (conform + voicing), note-timing (nudge and hold), displace (the two wheels) and generate, which walks the head. The panel row is data – src/app/panel-context-generator.ts declares the params and the plugin draws them
arranger/	The chord progression as session data: progression.ts and chord-track (chordDegreeAtTime – what chord is sounding at a given moment, read by the note-FX and by the generator)
automation/	Clip envelope recording + read-back, the automation painter and its LFO, the knob right-click menu – and the DestinationRegistry, the ONE catalogue every parameter picker must read (see docs/automation-destinations.md)
control/	Live MIDI controller subsystem: APC Key 25 profile, live keyboard, LED mediator, profile registry
demo/	Baked MIDI demos + demo picker
styles/	SCSS
plugins/	PLUGIN SOURCE – 24 shipped directories, each a plugin.json plus its code: six engines (tb303, subtractive, fm, wavetable, karplus, westcoast), fifteen inserts (autowah, bitcrusher, chorus, compressor, delay, distortion, flanger, gate, limiter, multifilter, phaser, reverb, ringmod, tremolo, width), two modulators (sh, pernote) and one PANEL (weave – a whole view, handed a DOM root and a PanelContext). A component that synthesises ships dsp.ts; an insert ships main.ts and builds native nodes on the main thread. Its tests live beside it. audio-probe is private:true – a fixture for the host's own tests, never shipped
packages/	
loom-plugin-sdk/	@loom/plugin-sdk – the surface a plugin author compiles against: manifest + param types, and the shared DSP primitives. One line in src/index.ts divides two species: above it runs per-sample in the worklet, below it builds native Web Audio nodes on the main thread
public/	
plugins/	PLUGIN BUILD OUTPUT – written by `npm run build:plugins`, discovered by the browser through index.json. Versioned, because the deploy publishes it. Editing plugins/ and not rebuilding ships the OLD plugin, silently
presets/	What the in-tree engines need: sampler, drums-machine and drum-kits.json. A melodic engine's presets ship with its plugin instead
drumkits/	Sampled drum kits: index.json + <id>.json manifests + WAVs


```
tools/
  loom-plugin/    The plugin CLI: `new` scaffolds a directory, `build` compiles
                  one (or plugins/*) into public/plugins/ and rewrites the index
  stem-service/   Local Python service (FastAPI + audio-separator / Demucs)
                  exposing an HTTP job queue for stem separation.
                  Run: uvicorn app:app --port 8765
                  Tests: python -m pytest test_app.py (not part of npm test)
  manual/         The pipeline that builds the manual you are reading:
                  assemble.mjs (chapters → one HTML document), shots.mjs
                  (screenshots), pdf.mjs, web.mjs, shot-list.mjs (the
                  hand-maintained list of screenshots + their selectors)
                  and manual.css. Driven by build-manual.mjs
```

Building the manual

The chapters in `docs/manual/*.md` are the only hand-written source. `index.html`, `Loom-Manual.pdf` and everything in `images/` are generated and committed — never hand-edit them.

Command	What it regenerates
<code>npm run build:manual</code>	Everything: builds the app, then screenshots + PDF + <code>index.html</code>
<code>npm run manual:shots</code>	Screenshots only
<code>npm run manual:pdf</code>	The PDF and <code>index.html</code> (no app build, no server)

`manual:shots` photographs whatever is already in `dist/` — it does not build. Run `npm run build` immediately before it, or you will capture a stale bundle (the same trap as `test:e2e`, below). A new chapter file is invisible to both outputs until it is added to the `CHAPTERS` array in `tools/manual/assemble.mjs`.

Testing

Loom has four test layers, one per risk class.

Pure logic (`src/**/*.test.ts`, excluding `.dsp` and `.wiring` suffixes) — schemas, scales, migrations, session/arrangement logic, modulation math. These run fast and have no audio dependencies.

Scheduling with a fake clock — `src/core/lane-scheduler.test.ts` and `src/session/session-runtime.test.ts` drive the look-ahead scheduler through a mock `AudioContext` clock. The fake clock advances in controlled steps so timing edge-cases are deterministic.

Real DSP (`*.dsp.test.ts`) — audio actually rendered and measured, in two techniques. The **pure kernel** is driven sample by sample with no `AudioContext` at all: `src/audio-dsp/drums/new-voices.dsp.test.ts` calls `renderSample()` in a loop and asserts each voice sounds like what it claims to be. `src/audio-dsp/modulation-scope.dsp.test.ts` is the same technique, and so is `src/audio-dsp/live-params.dsp.test.ts` — the registry-driven one that renders every engine in `WORKLET_ENGINE_IDS` twice,

moving a knob mid-note in the second render, and fails the engine whose sound does not change. The **Web Audio nodes that stayed native** render through `OfflineAudioContext` via `node-web-audio-api`, globalised in `test/setup.ts`: `comp-block`, `master-comp`, `master-shaper`, `strip-ducker`, `multifilter`, the sample/warp helpers and the offline export. (`src/performance/arrangement.dsp.test.ts` carries the suffix but does neither — it is arrangement maths. Glob for `*.dsp.test.ts` rather than trusting this list to stay complete.)

A plugin's DSP is tested inside the plugin, beside the code: `plugins/<id>/dsp.test.ts`, plus `plugins/tb303/tb303-parity.dsp.test.ts`, which pins the 303 against a committed `reference-render.json` so a refactor that changes the sound has to say so out loud. `packages/loom-plugin-sdk/src/sdk-parity.test.ts` guards the primitives those plugins share. They run in the same `npm run test:unit` sweep as everything else.

The golden-WAV loop runs on the worklet path. `test/preset-render.dsp.test.ts` renders a fixed phrase for the first three presets of each melodic engine through `renderKernelLane` — the same pure kernel the offline export uses — and writes `test/output/`; `npm run test:wav-diff` compares that against the 18 WAVs committed in `test/golden/` and prints peak / RMS / L2 deltas (it never fails CI — it is a human inspection tool), and `npm run test:wav-bless` overwrites the golden set on purpose, so commit the result. Its sibling `test/preset-stability.dsp.test.ts` is the divergence detector over **every** shipped preset: finite, bounded, and the voice ends. (The pre-worklet batteries and their 90 orphan WAVs are gone; `test/dsp-battery.ts` no longer exists.)

Two listening tools live in `tools/` and are gated on an environment variable, so the regular sweep collects them and runs nothing. `tools/render-preset-auditions.test.ts` (`AUDITION=1`, `AUDITION_ENGINES=fm,westcoast`) renders every factory preset of the named engines four ways — plain, accented, a four-note chord, a long low note — raw and through the live master tail, and writes `test/output/auditions/` with an `index.html` listening page whose ratings export as JSON. It answers "does this **factory preset** render clean". `tools/render-session-lanes.test.ts` (`REPRO=1`, `REPRO_SAVE=/path/to/session.json`, required) answers "does **my lane**": for each melodic lane in a saved session it renders that lane's exact saved params with its own busiest clip at several polyphony caps (`REPRO_VOICES`, default `1,8,39`), raw and mastered, into `test/output/session-lanes/`. That is the reproduction shape for "sounds fine at minimum polyphony and turns to noise when Voices rises" — if the offline kernel reproduces the noise it is DSP, if it does not it is worklet-only. The session tool multiplies the manifest `outputTrim` by the synth category gain, as the host does; the audition tool applies the manifest trim alone, as the golden loop does.

Modulation, objective and end-to-end — `src/audio-dsp/modulation-pipeline.test.ts` drives the real in-engine path (`ModulationRuntime` → `VoiceManager` → `renderer`) for each of the six melodic engines, with an LFO at full depth on a continuous param, and asserts the rendered RMS envelope differs measurably from the unmodulated render, plus a negative control. It exists because the worklet rewrite dropped the per-engine coverage the old `.wiring.test.ts` files had. One `.wiring.test.ts` remains — `src/core/ducker-subgraph.wiring.test.ts` — and it covers Web Audio subgraph wiring like the sidechain, which is the only place that pattern still applies.

Assertion rule: always write relative assertions (`a > b` , `a > b * 2`). Never hard-code absolute magnitudes — they are a brittleness smell. If you must write one, justify it in a comment.

Colour-free output: every `npm test` script runs under `cross-env NO_COLOR=1` . When invoking Vitest directly, prefix with `NO_COLOR=1` . Do not add `--reporter=...` — the scripts already configure the right reporter.

Key commands:

Command	What it runs
<code>npm run dev</code>	Vite dev server with hot reload at http://localhost:5173
<code>npm run build</code>	Every plugin, then <code>tsc</code> typecheck + Vite bundle to <code>dist/</code>
<code>npm run build:plugins</code>	<code>plugins/*</code> → <code>public/plugins/</code> + its <code>index.json</code>
<code>npm run plugin -- new plugins/<id></code>	Scaffold a new plugin directory
<code>npm test</code>	Full suite: unit + e2e (always build first)
<code>npm run test:unit</code>	Vitest only, no browser
<code>npm run test:fast</code>	Unit tests excluding DSP renders (inner-loop TDD)
<code>npm run test:dsp</code>	DSP renders only (slow; needs <code>node-web-audio-api</code>)
<code>npm run test:e2e</code>	Playwright against <code>vite preview</code> on port 4173

e2e gotcha: `test:e2e` and `npm test` serve `dist/` with no build step. Playwright boots `vite preview` over the last production bundle. If you changed `src/` without rebuilding, the newest features are absent from the bundle and tests fail with "element not found" — which looks like a regression. Always run `npm run build` before `npm run test:e2e` .

The same trap has a plugin-shaped twin. `public/plugins/` is build output, and nothing recompiles it for you — not the dev server, not `vite build` on its own. Change a plugin, reload, and you hear the previous version with no warning; add a new one and its engine is simply missing from the selector. `npm run build:plugins` after every plugin edit, or `npm run build` , which does it first.

Vitest runs test files serially (`fileParallelism: false`) because `node-web-audio-api` 's `OfflineAudioContext` is not safe under parallel forks. The teardown occasionally exits non-zero with `ERR_IPC_CHANNEL_CLOSED` after all tests pass — that is a tinypool shutdown race, not a test failure; re-run to confirm green.

`CLAUDE.md` at the repo root is the short-form architecture summary; this chapter is the long form. The convention is that a design doc is pruned from the tree once its work ships (specs drift faster than anything else), so recover rationale from git history when you need it: `git log --diff-filter=D --name-only -- docs/superpowers/` .

What is still open — code debts, and the specs that have come back into `docs/superpowers/` since the last prune — is inventoried in `docs/superpowers/REMAINING-WORK.md` .

(Links to files outside `docs/manual/` are ordinary links: the single-page build turns a link to a sibling *chapter* into an in-page anchor and leaves everything else alone, promoting `../../*` paths to absolute GitHub URLs so they work from the shipped page too. See `rewriteChapterLinks` and `rewriteRepoLinks` in `tools/manual/assemble.mjs` .)